

A WORKING PACK

AI for Finance Teams

The Working Pack



Tang Shiuan Jenn

Chartered Accountant

No invented statistics. No fabricated case studies. No promises about your income.
Everything in this pack is either an artefact you can use or a claim you can check.

Contents

1. How to Use This Pack
2. Choosing What to Automate First
3. The Prompt Library
4. Verification, Red Lines and Choosing a Tool
5. Your Policy Documents
6. The 90-Day Plan and the Training Curriculum

How to Use This Pack

What this is

This is a working pack, not a book. Nothing in it needs reading cover to cover. It is a set of artefacts you can put in front of a team on Monday: prompts to paste, a policy to adapt, a plan to follow, a scoring sheet to argue over, and a list of the places where the honest answer is to close the laptop.

It was written by a chartered accountant, for people whose name goes on the work.

What it deliberately does not contain

No statistics about AI adoption. No productivity percentages. No case studies about firms that saved a fortune. No promises about your income or your job.

Everything in that list is standard in this market and almost none of it is checkable. The moment a pack like this quotes a figure nobody can verify, every other claim in it becomes a guess too. So there are no numbers here except the ones you generate yourself, in your own team, measuring your own work.

What that costs you: this pack cannot tell you what results to expect. What it buys you: everything it does tell you can be checked.

How to use it

If you have not started at all. Read *Choosing What to Automate First*, run the 45-minute meeting at the end of it, and pick two use cases. Two, not ten. Then take the policy from *Your Policy Documents*, fill it in, and circulate it before anyone else starts experimenting. That order matters — a policy written after the tools have spread is a document about a problem you already have.

If your team is already using AI without rules. Start with *Your Policy Documents* today. Then *Verification*, *Red Lines* and *Choosing a Tool*. The prompts can wait a fortnight; the red lines cannot.

If you are the one who has to teach it. Go to the four-level curriculum in *The 90-Day Adoption Plan and the Training Curriculum*. Each level is written so you can teach it without outside help, with the live exercise already specified.

If you just want better output tomorrow morning. *The Prompt Library*. Find the workflow that matches what is on your desk, adapt the prompt rather than pasting it blind, and read the check line underneath it before you use anything it gives you.

The one rule underneath all of it

Whoever signs the work owns the work.

A tool is never the reason something went out wrong. Every prompt in this pack has a line under it telling you what to check, because a fluent answer and a correct answer look identical until you look. That is the whole discipline. The rest is technique.

A note on what is general and what is yours

Everything here that touches law, tax, professional standards or regulatory duty is general information. Rules differ between countries and professions, and they change. No document in this pack is legal advice, and adopting any of it is not a statement that you comply with anything. Have the policy read by whoever carries risk in your organisation, and check it against your own regulator's guidance and your client engagement terms before you use it.

Choosing What to Automate First

Ask a finance team where they want to start with AI and you will usually hear one of three answers: variance commentary, board paper summarising, or "the audit file". These are the three most interesting problems in the room. They are also, on any honest assessment, three of the worst places to begin.

The reason is not that the work is too hard. It is that in all three cases you cannot cheaply tell whether the output is right. A draft that sounds like a competent colleague wrote it, and is wrong in one place you did not look, costs you more than doing the job yourself. You pay twice: once for the draft, once for the review that has to be as careful as the original work.

What follows is a way of choosing that does not rely on anyone's enthusiasm: five factors, scored one to five, applied to the tasks your team actually does. It takes forty-five minutes with the right people in a room, and it produces two starter use cases you can defend to a sceptic, plus a written reason for everything you left out.

Part One — The five-factor score

Score every candidate task on five factors, one to five. Total out of twenty-five. Two of the factors are inverted, which matters — read the definitions before you score anything.

Factor 1 — Frequency

How often does this task recur?

Score	What it looks like
1	Once a year or less. A one-off.
2	A few times a year. Quarterly.
3	Monthly.
4	Weekly, or several times a month.
5	Daily, or many times a day.

Frequency is where the money is. A task that takes an hour and happens once a year is worth an hour a year. The same task done weekly is worth fifty. Teams routinely pick the annual monster over the weekly irritation, because the monster is memorable and the irritation is invisible.

One exception, and it is a real one. Some low-frequency tasks produce an artefact that then gets used constantly — a close checklist, a set of KPI definitions, a process note. Score the *task* honestly at 1 and note the durable output separately. A 1 with a permanent artefact behind it is a legitimate reason to promote a task up the list. A 1 with a disposable output is not.

Factor 2 — Time

How long does one instance take now, done by the person who normally does it?

Score	What it looks like
1	Under ten minutes.
2	Ten to thirty minutes.
3	Half an hour to two hours.
4	Half a day to a full day.
5	More than a day.

Score the time it takes *now*, not the time you imagine it should take. If nobody knows, time three instances before the meeting. Guessing here corrupts everything downstream, because time and frequency together are the only two factors that describe the prize. The other three describe the cost of collecting it.

Do not score the time you think the tool will save. You do not know that yet, and the point of the exercise is to stop deciding on things you do not know.

Factor 3 — Verifiability

How easily can you tell whether the output is right?

Score	What it looks like
1	You cannot tell without redoing the work from scratch, or correctness is a judgement no one available can make quickly.
2	Checking takes nearly as long as doing. Errors are plausible and quiet.
3	Checkable, but only by someone who knows the subject, reading the whole thing.
4	A knowledgeable person spots an error quickly — minutes, not hours.
5	Correctness is mechanical and obvious. It ties, it foots, it matches the ledger, or the sheet returns the answer you already knew.

This is the factor teams overscore, every time, because verification feels like an afterthought rather than part of the work. It is not an afterthought. It is the entire economics of the thing.

Write the verifiability score last, and write it as a sentence first: "I would know this was wrong because ____." If you cannot finish that sentence in one breath, the score is 1 or 2. If the sentence is "because I would read it carefully", that is a 2 or a 3, not a 5 — careful reading is exactly the expensive thing you were trying to avoid.

Factor 4 — Consequence of error — inverted

If a mistake gets through, what happens? **High score means cheap mistake.** This factor runs backwards on purpose, so that every factor points the same way and the total means something.

Score	What it looks like
1	The error reaches a client, a regulator, a lender, a court or a board, or a professional duty attaches to the work. Expensive and slow to unwind.
2	The error reaches someone outside your team and has to be corrected formally. Embarrassing.
3	The error flows into something else — a return, a report, a system — and gets found later, out of context.
4	The error is caught in normal internal review. Costs minutes and some grumbling.
5	The error is caught instantly by the person who asked, and nobody outside the room ever knows.

Score the consequence of the error *escaping*, not of it being made. Everything gets reviewed in theory. Ask what happens on a Friday afternoon in a busy week.

Factor 5 — Data sensitivity — inverted

What has to go into the tool for the task to work? **High score means nothing confidential.** Also inverted.

Score	What it looks like
1	Impossible without named personal data, client confidential material, unpublished price-sensitive information, or credentials.
2	Real internal records — ledger detail, supplier and customer names, transaction data.
3	Internal material that identifies your organisation and its people, but no client data and nothing price-sensitive.
4	Internal but generic. Procedures, structures, definitions.
5	Nothing confidential at all. Public, generic, or fully anonymised without effort.

Two practical notes. First, "we would anonymise it" only counts if anonymising is quick and reliable. Twenty minutes of careful redaction per instance belongs in the Time score, and the task is less attractive than it looked.

Second, this is general information and not legal advice. What you may put into which tool depends on your contracts with clients, your own confidentiality obligations, the rules of your professional body and the data protection law where you and your clients sit. Those differ by jurisdiction and they change. Get your own position checked rather than relying on a score in a table.

Reading the total

Total	Reading
20–25	Start here. Rare. If you have one, do it this month.
16–19	The working middle. Most of a sensible programme lives in this band.
11–15	Conditional. Worth doing only if you can redesign the task to lift the weak factor.
10 or below	Leave it. Revisit in a year if something changes.

Then apply three vetoes, which override the total completely:

1. **Verifiability of 1 or 2 — vetoed.** Whatever the total.
2. **Consequence of 1 — vetoed.** Whatever the total.

3. Data sensitivity of 1 — vetoed. Whatever the total.

A data sensitivity of 2 is not a veto but a constraint: that task is only available to you inside a tool where the material genuinely stays under your control, on terms someone has actually read. If you do not have such a tool, every 2 behaves like a 1 and drops out.

The vetoes exist because a total is an average, and averages hide the one thing that will hurt you. A task scoring 4, 5, 2, 1, 1 — frequent, slow, unverifiable, catastrophic if wrong, full of client data — totals 13 and sits comfortably mid-range. Thirteen out of twenty-five is not a comfortable task. It is a disaster with good attendance.

Why the "time" trap catches everyone

Here is the failure mode this whole method exists to prevent.

Someone proposes a task that scores 4 or 5 on time. It is the task that eats their Thursday. It is genuinely painful and everybody agrees it is painful, which makes it feel like an obvious candidate. But it scores 2 on verifiability, because the only way to know the output is right is to do the underlying work.

The arithmetic is brutal. The task takes four hours; the tool drafts it in ten minutes. You have not saved three hours and fifty minutes. You have saved nothing until the draft is checked, and checking an unverifiable four-hour output takes close to four hours, because the check *is* the work. You have added ten minutes and a false sense of progress.

Worse, the draft is persuasive. It is fluent, well-structured and confident. A checker reading a fluent draft looks for errors; a person doing the work from scratch encounters the facts. Those are different mental activities and they catch different things. A fluent wrong draft is harder to correct than a blank page, because it tells you where to look and the place it tells you to look is not where the error is.

So: time tells you the size of the prize; verifiability tells you whether you can collect it. A high time score with a low verifiability score is not a good use case with a snag. It is a bad use case that looks good, which is worse than one that looks bad, because someone will champion it.

The corollary is cheering. Tasks that score modestly on time — twenty minutes, done daily, checked in ten seconds — are the ones that quietly pay. They are unglamorous, nobody will present them at a conference, and they will save your team more hours this year than the Thursday monster ever will.

Part Two — Thirty tasks, scored

Below are thirty tasks common to finance functions and professional firms, scored against the five factors. The scores are mine and they are arguable — deliberately so. The point of a published scale is that you can look at a row, disagree with one number, and see exactly which definition we read differently. Change the score, change the total, carry on. What you must not do is change a score because you want a different answer.

Assume a general-purpose assistant of the kind anyone can reach from a browser. Where a row is constrained by data sensitivity, the constraint means an enterprise-tier arrangement your organisation controls, on terms someone has read.

Key: **F** frequency, **T** time, **V** verifiability, **C** consequence (inverted), **D** data sensitivity (inverted).

#	Task	F	T	V	C	D	Total	Verdict
1	Explaining an inherited spreadsheet formula	4	2	5	5	5	21	Best possible starter — paste the formula, not the data; the sheet marks the answer.
2	Writing a formula or query from a plain description	4	2	5	5	5	21	Same family. Test on a sample where you know the result before trusting it.
3	Drafting an internal training exercise for the team	2	3	4	5	5	19	Low risk, genuinely useful. Check the model answers yourself.
4	Reconciliation prep — formatting and first-pass matching	5	3	5	4	2	19	High frequency and the ledger marks the work. Constrained only by where the data may go.
5	Meeting notes into circulated minutes	5	2	4	4	3	18	Strong starter. You were in the room, so you can mark the output.
6	Routine client emails — updates, chasers, acknowledgements	5	1	5	4	3	18	Small saving each time, large in aggregate. Never send unread.
7	Process documentation for a routine task	2	4	4	4	4	18	High value; the person who does the task is the reviewer. Backlog makes it better than it scores.
8	Month-end close checklist — build or refresh	1	3	5	4	5	18	One afternoon, permanent artefact. An excellent pilot.
9	Internal finance policy — first draft	1	4	4	4	5	18	Good use. The draft is a skeleton; the control detail must be yours.
10	Job description for a finance vacancy	2	3	4	4	5	18	Fine filler task. Do not let it invent qualifications, or say anything about pay you have not decided.
11	Audit request list / information list for a client	2	2	4	4	5	17	Quick win before every audit. Your auditor will amend it anyway.
12	Interview questions and scoring sheet for a finance hire	1	3	4	4	5	17	Solid and low-stakes. Keep the scoring criteria yours.
13	User acceptance test cases for a system or control	1	4	4	4	4	17	Good. The tool supplies breadth; you supply the nasty cases.
14	KPI definition dictionary	1	3	5	4	4	17	Unglamorous, durable, week-one material. Your team recognises a wrong definition instantly.
15	Plain-language rewrite of a technical memo	3	2	4	4	3	16	Common and useful. The check is fast because you own the source.
16	Forecast scenario — building the structure	2	4	4	3	3	16	Use it for scaffolding, never for assumptions. Test the mechanics with known inputs.
17	Cleaning and standardising a messy data extract	4	3	3	3	2	15	Only with row-count and total controls in place. Silent row loss is the risk, not visible errors.
18	Invoice coding suggestions against the chart of accounts	5	1	4	3	2	15	Volume is the attraction. Treat every suggestion as a first pass a human approves.
19	Fixed asset register tidy-up — descriptions and categories	1	4	3	3	3	14	Middle of the pack. Check categories yourself; they drive depreciation and tax.

#	Task	F	T	V	C	D	Total	Verdict
20	Long email chain into a decision log	4	2	3	3	2	14	Useful for your own recall. Weak as a record anyone else relies on.
21	Variance commentary — first draft from figures you supply	3	3	3	3	2	14	The classic trap. Quick to draft, slow to trust — the tool cannot know <i>why</i> the number moved.
22	Contract summarising into a key-terms schedule	3	4	3	2	1	13	Vetoed on data. The check is re-reading the contract, which was the work.
23	Board paper summarising	2	3	3	2	1	11	Vetoed on data. The omission you cannot see is the entire risk.
24	Audit working paper — documenting work performed	4	3	2	1	1	11	Vetoed three ways. The record must describe what you actually did.
25	Tax computation review	2	4	2	1	1	10	Vetoed. Use it for the covering memo if anything, never the computation.
26	Payroll queries involving named individuals	3	2	3	1	1	10	Vetoed. Named pay data, and the error lands on a person.
27	Credit assessment of a named customer	3	3	2	1	1	10	Vetoed. A decision about a named party on data you cannot check.
28	Researching the treatment of a novel transaction	2	4	1	1	2	10	Vetoed. Where confident-sounding wrong answers do the most damage.
29	A difficult client email — bad news, fee dispute, complaint	2	2	3	1	1	9	Vetoed. In a live dispute your words become evidence. Write it yourself.
30	Drafting a response to a regulator or professional body	1	3	2	1	1	8	Vetoed. The lowest score here, and rightly.

What the table is telling you

Read the shape of it rather than the individual rows.

The top is boring. The two highest scores are explaining and writing spreadsheet formulas — tasks so small that nobody would propose them in a strategy meeting. They win because the sheet tells you immediately whether the answer works, and because you paste a formula rather than data. Frequency and verifiability, not ambition.

The bottom is everything anyone was excited about. Board papers, tax computations, working paper documentation, regulator correspondence. These are the tasks that hurt most, which is precisely why they involve the most confidential material and the heaviest consequences. Pain and suitability are not correlated. If anything they run against each other.

The middle is where the argument is. Rows 17 to 21 all score 14 or 15 — data cleaning, invoice coding, variance commentary. These are the rows worth an hour of your team's disagreement, because the score depends on facts about your organisation that I do not have. If your ledger data can sit in a tool you control, invoice coding moves from 15 to 18 and becomes a serious candidate. If it cannot, it drops out entirely. Same task, different organisation, different answer.

Two rows deserve a word each.

Variance commentary (row 21) is the single most requested use case in finance, and it scores 14. The figures are verifiable; the explanation is not. A tool can see that distribution costs moved. It cannot know that you changed courier in week three, and it will not say "I don't know" — it will produce a plausible reason. Plausible reasons in management commentary are worse than blank space, because someone will act on one. If you use a tool here at all, use it to draft the *structure* and to check that every movement above your threshold has been addressed. Supply the causes yourself.

Audit working paper documentation (row 24) scores 11 and is vetoed three times over. It deserves the emphasis. A generated narrative describing work performed is indistinguishable, on the page, from a narrative describing work performed — whether or not it was. That is not a technology problem you can solve with a better prompt. It is the reason the record exists. Where professional duties attach to your file, the documentation is the evidence of the work, and it has to be written by the person who did it. Check what your own professional body and your firm's methodology require; this varies and it changes.

Disagreeing properly

Someone in your team will say a score is wrong. Good. Handle it like this:

1. Name the factor. Not "this should be higher" — "verifiability should be 4, not 3".
 1. Quote the definition. Read the row from the scale out loud.
 2. Say what you would do differently. "I would know it was wrong because the control total does not match" is an argument for 5. "I would read it properly" is not.
1. Change the number and let the total fall where it falls.

The one rule: nobody argues about a total. Totals are arithmetic. If you find yourself defending a total, you have decided the answer and are reverse-engineering the scores, and the method has stopped working.

Part Three — Running this in your own team in forty-five minutes

This is a single meeting. It does not need a workshop, a facilitator or a wall of sticky notes, and it produces a decision rather than a discussion.

Who is in the room

Four to eight people. You need three roles covered, and one person can cover two:

- **People who do the work.** They score frequency and time. Nobody else is entitled to.
- **People who review the work.** They score verifiability and consequence. The reviewer's verifiability score wins any tie, because they are the one who will pay for a bad one.
- **Whoever owns confidentiality and data protection.** They score data sensitivity, and they hold a veto that needs no explanation.

Do not invite anyone whose only contribution is enthusiasm.

Before the meeting

Circulate the five factor scales — just the five tables from Part One, one page. Ask everyone to bring **three tasks they personally do**, written on separate lines, described concretely. "Reporting" is not a task. "Writing the commentary on the departmental cost report" is.

If nobody knows how long a candidate task takes, ask the person who does it to time three instances that week. Ten minutes of measurement beats an hour of estimating, and the number is defensible afterwards.

The agenda

Time	Item	Output
0:00–0:05	Purpose and rules. Read out the three vetoes.	Everyone knows a high total can still lose.
0:05–0:12	Silent listing. Each person writes their three tasks. No discussion.	12–24 raw candidates.
0:12–0:20	Read out and merge. Combine duplicates, discard anything nobody actually does.	A clean list of 15–25.
0:20–0:35	Score. Thirty seconds a task, by role, in order F-T-V-C-D.	Every task has five numbers.
0:35–0:42	Apply vetoes. Strike the failures. Sort by total.	A ranked, filtered list.
0:42–0:45	Pick two. Name an owner each. Set the review date.	A decision.

Fifteen minutes for scoring sounds tight and is not. Thirty seconds a task is enough when the scales are on the table in front of people, and the clock is doing useful work: it stops the room relitigating a task everyone already understands.

Rules for the room

- **Silent listing first.** Open with discussion and the most senior or most talkative person sets the agenda; you will score their three tasks.
- **Score by role, not by consensus.** The doer's time score stands, the reviewer's verifiability score stands. Consensus drifts towards whatever the room already wanted.
- **No scoring your own task alone.** If you do the task and review it, someone else scores verifiability.
- ****Write the "I would know it was wrong because _____" sentence out loud**** for anything scoring 4 or 5 on verifiability. Half of them collapse to a 3 the moment they are spoken.
- **Vetoes are silent.** No debate. The veto is not a judgement about the

task or the person who proposed it.

What you leave with

Two use cases. Not five, and not one.

One gives you no comparison: if it goes badly you cannot tell whether the method or the task was at fault. Five gives you five half-started experiments and no evidence, because no small team learns five new habits at once.

For each of the two, write down before you leave the room:

1. The owner. One named person, not a team.
2. The baseline: how long the task takes now, measured over three

instances, and how often it happens.

1. What "working" would look like in four weeks, in time or rework — not in enthusiasm.

1. The verification step, in one sentence, written before anyone starts.
2. The date you will look at it again. Four weeks is about right.

After four weeks

Come back for fifteen minutes. Three possible answers:

- **It worked.** Write down what verification actually cost, because that

is the number you will need when someone proposes the next task. Then score the next two candidates from your ranked list.

- **It half-worked.** Usually this means the verification step was invented afterwards rather than designed first. Fix that and run another four weeks. Do not widen.

- **It did not work.** Say so plainly, record why against the five factors, and move to the next candidate. A use case that fails against a written scale is information. A use case that fails against nothing is just a bad mood.

If nothing scores well

This happens, particularly in small practices where nearly everything touches client confidential material. If your list has no task above 15 and every high scorer is vetoed on data, that is a genuine finding, and it is worth more than a forced pilot.

It usually points at one of two things. Either your tooling is the constraint — the work would qualify if the material stayed under your control, which is a procurement question rather than a use case question. Or your suitable tasks are all small and you have been looking for something big. In the second case, go back to the top of the table: formula explanation, checklists, definitions, internal drafting. Small, frequent, instantly checkable, nothing confidential. That is not a consolation prize. On the arithmetic, it is the answer.

Re-run it

Score again every quarter, or whenever one of three things changes: the tools available to you, your policy on what data may go where, or the shape of the work. Keep the old scoresheets. Being able to show why a task was rejected in March and what changed by September is worth as much as the decision itself — particularly to whoever has to answer for it.

The Prompt Library

A prompt is a briefing, not a spell. Paste one, take whatever comes back and put your name on it, and this will go badly. Every prompt below assumes you know things the model does not: what your entity does, which numbers matter this period, and what happens if the answer is wrong. The prompt gets that out of your head in an orderly way. You do everything after.

Change three things in almost every prompt here. The role line, because "assisting a finance team in a group with overseas subsidiaries and a manual consolidation" produces different work from "assisting a finance team". The constraints, so they carry your house rules — anything a new joiner would get wrong. And the format, so the output feeds your template rather than being retyped into it.

Leave two things alone: the instruction to mark uncertainty and refuse to guess, and the ban on inventing figures. Remove either and you get confident filler, the expensive failure in this work. First, though, settle a separate question — whether you may put this information into this tool at all. Client data, staff data and anything confidential are decisions for whoever owns that policy, not for you at your desk. If nobody has decided, assume not.

Change this	Keep this
The role line, to match your entity	The uncertainty rule
The constraints, to carry your house rules	The ban on inventing figures
The format, to fit your template	The instruction to list what it could not determine

1. Month-end close and reconciliations

P1.1 — Reconciliation break triage

Use when: you have unexplained differences and little time.

Paste this:

You are assisting a finance team at month-end on [ACCOUNT], [PERIOD].

Task: group the reconciling items below by likely cause and say what evidence would confirm or rule out each.

Constraints: use only what I paste. Invent no amount, date or counterparty.

Format: a table – Item, Amount, Likely cause, Confirming evidence, Who to ask. Then a list headed "Cannot classify".

Anything you are unsure about goes there. Never guess.

Items:
[PASTE]

Check before you use the output: the groups are inference, not findings. Confirm each against the source document.

P1.2 — Break explanation, written up

Use when: you must write up a difference you have explained.

Paste this:

You are helping an accountant write a file note on a difference on
[ACCOUNT], [PERIOD].

Task: turn my notes below into a note another accountant could follow in
a year without asking me anything.

Constraints: use only my facts. Add no causes or conclusions, and
neither soften nor strengthen mine.

Format: four short paragraphs – what the difference was, why it arose,
what was done, what remains open. Under 300 words.

Fill no gaps; list them under "Questions I could not answer".

Notes:
[PASTE]

Check before you use the output: the "remains open" paragraph. Models tidy loose ends away.

P1.3 — Close checklist builder

Use when: your close runs on memory and one person's spreadsheet.

Paste this:

You are documenting the month-end close at [ENTITY].

Task: turn the tasks below into a checklist someone who has never run
our close could follow.

Constraints: include only tasks I describe. Add no standard tasks you
assume we perform, and no deadlines I have not given.

Format: a table – Working day, Task, Owner, Depends on, Evidence. Then
any task I did not mention, as a question.

Where sequence or ownership is ambiguous, write UNCLEAR.

Our tasks, in my own words:
[PASTE]

Check before you use the output: the dependency column. A task placed before its dependency is worse than no checklist.

P1.4 — Intercompany mismatch plan

Use when: two entities disagree and you need a plan before the calls.

Paste this:

[ENTITY A] and [ENTITY B] disagree on the intercompany balance for [PERIOD].

Task: an investigation plan – causes to test in order, cheapest and likeliest first.

Constraints: work only from what I tell you. Assume nothing about our systems, currencies or cut-off conventions.

Format: numbered steps, each stating what to check, what result confirms the cause, and what to do if it does not.

End with "Assumptions I had to make". If you need more, ask.

What each entity reports:
[PASTE]

Check before you use the output: the assumptions list, then every step resting on it.

P1.5 — Accrual and prepayment review questions

Use when: you are reviewing someone else's schedule.

Paste this:

You are reviewing a schedule of accruals and prepayments for [ENTITY], [PERIOD].

Task: the questions a careful reviewer would ask – aimed at what would be wrong if the preparer had been rushing.

Constraints: recalculate nothing and call no figure wrong; you cannot see the evidence. Name no accounting standard.

Format: questions under Completeness, Cut-off, Basis of estimate and Reversal. Maximum 20, ordered by money at stake.

Where the schedule gives too little for a real question, say so.

Schedule:
[PASTE]

Check before you use the output: whether the questions are specific to your schedule. Generic ones mean thin context.

2. Variance analysis and management commentary

P2.1 — First-pass variance questions

Use when: the variance report has landed and nobody has spoken to the business.

Paste this:

You are assisting a controller investigating variances for [ENTITY], [PERIOD].

Task: for each line I flag, the questions to put to the budget holder – those separating timing, volume, price, misposting and real change.

Constraints: do not explain the variances yourself; you do not know our business. Let no presumed cause appear in a question's wording.

Format: for each flagged line, the account name then three to six questions.

If a line is too small to be worth a conversation, say so.

Table:
[PASTE]

Check before you use the output: leading questions. A question naming a cause returns the answer you suggested.

P2.2 — Commentary from your explanations

Use when: you have the reasons and need them written up.

Paste this:

You are drafting management commentary for [ENTITY], [PERIOD], for [AUDIENCE]. My notes give each account, its figure and the business's explanation – the only permitted source of cause.

Task: write the commentary.

Constraints: add no context, market conditions, forward-looking statements or figures. Do not tell the reader how to feel.

Format: one short paragraph per variance, largest first, opening with the account name in bold. Maximum [WORD LIMIT] words.

Flag any explanation that does not account for its movement.

Notes:
[PASTE]

Check before you use the output: every causal phrase. Each must trace to a named person.

P2.3 — The walk from last period to this

Use when: you need a bridge a reader can follow without a chart.

Paste this:

You are explaining a movement in [MEASURE] from [PRIOR PERIOD] to [CURRENT PERIOD] at [ENTITY].

Task: write the bridge as prose, in the order that reads clearest.

Constraints: the components must sum to the movement as given. If they do not, stop and tell me the difference. Insert no balancing item.

Format: a lead sentence stating the movement, then one sentence per component.

Say plainly if you cannot reconcile my components to my movement.

Inputs:
[PASTE]

Check before you use the output: whether it reported a reconciliation failure rather than writing around it.

P2.4 — De-waffling a draft

Use when: the commentary reads like a press release.

Paste this:

You are an unsentimental editor of finance writing.

Task: rewrite the draft below so every sentence carries information. Cut hedging and adjectives that assert rather than describe.

Constraints: change no figure, cause or conclusion, and add no facts. If cutting would lose a fact, keep the fact and cut the decoration.

Format: the rewritten text, then the sentences you could not tighten because they made a claim the draft does not support.

Never invent a stronger version of a weak claim. Flag it.

Draft:
[PASTE]

Check before you use the output: the figures, line by line. Editors round numbers to make sentences sit better.

P2.5 — Attack my explanation

Use when: you are about to present a reason and want it tested.

Paste this:

You are a sceptical audit committee member.

Task: attack the explanation below. Find the questions that would expose it if it were wrong, incomplete, or an artefact of the numbers.

Constraints: do not present alternatives as fact, and invent no figures. Frame each challenge as a question I can answer.

Format: three groups – size, timing, simpler explanation. Maximum ten, hardest first.

If you cannot fault it, say so rather than manufacture objections.

My explanation:
[PASTE]

Check before you use the output: whether each challenge is answerable with evidence you can obtain. Park the rest.

3. Budgeting and forecasting

P3.1 — Assumption map for a budget line

Use when: a budget line has become one number nobody can defend.

Paste this:

You are making assumptions explicit for [BUDGET LINE] at [ENTITY] for [PERIOD].

Task: break the line into the drivers needed to build it from scratch, and say what source holds each answer.

Constraints: give no values, ranges, percentages or benchmarks – not one number. Where a driver depends on a decision, say whose.

Format: a table – Driver, Unit, Type (fact / estimate / decision), Where the answer lives, Question to ask.

If the line could be built two ways, show both.

Check before you use the output: that the drivers actually rebuild your line.

P3.2 — Budget holder briefing

Use when: the budget pack goes out to non-finance managers.

Paste this:

You are briefing non-financial managers who must submit figures for [PERIOD]. Their cost lines are [LIST]; the deadline is [DATE].

Task: write the note that accompanies the template.

Constraints: no jargon without a plain gloss. No threats and no cheerleading. State no figure, limit or target.

Format: a short opening on what we need and when; a numbered list of what to submit; another of what not to do.

Leave a marked [] wherever a figure or name is needed.

Check before you use the output: that every placeholder is still empty.

P3.3 — Forecast assumption challenge

Use when: the forecast is built and you want the weak joints found.

Paste this:

You are reviewing the assumptions behind a forecast for [ENTITY] over [PERIOD].

Task: identify which assumptions the outcome is most sensitive to, which are mutually inconsistent, and which cannot be tested.

Constraints: never tell me an assumption is too high or too low – you have no basis. Compute no revised outcomes.

Format: three sections – Most sensitive, Mutually inconsistent, Not testable as written. Name each and state the problem.

Where facts I have withheld might reconcile something, ask.

Assumptions and values:
[PASTE]

Check before you use the output: the sensitivity claims. Without your model they are a list to test, not a finding.

P3.4 — Scenario definitions, without numbers

Use when: scenarios were asked for and nobody has defined them.

Paste this:

You are defining scenarios for [ENTITY] over [PERIOD].

Task: define [NUMBER] scenarios. Each is a coherent story about the world, not a percentage applied to a base case.

Constraints: no numbers, no probabilities, no likelihoods. Each must stay internally consistent and use only my uncertainties.

Format: for each – a name, a paragraph, bullets showing which assumptions move and which way, and the earliest observable sign.

If two uncertainties cannot vary independently, say so.

Business and uncertainties:
[PASTE]

Check before you use the output: the early-warning signs — the part that gets used, and the likeliest platitude.

P3.5 — Budget submission review

Use when: submissions are in and you must review many quickly.

Paste this:

```
You are reviewing a budget submission from [DEPARTMENT] for [PERIOD]
against our rules.

Task: list every breach, and every required explanation missing or too
vague.

Constraints: do not assess whether the figures are reasonable – that is
mine. Quote the submission's exact wording when you flag it.

Format: a table – Rule, what the submission says, Problem, what to
request. Then: ready for review, or return to sender.

If a rule is ambiguous, flag the rule rather than choose.

Rules, then submission:
[PASTE]
```

Check before you use the output: the quotations. If one is not verbatim, discard the review and run it again.

4. Audit and assurance support

P4.1 — Making sense of a request list

Use when: the auditor's list has arrived and you must allocate work.

Paste this:

```
You are planning a response to an auditor's request list for [ENTITY],
[PERIOD].

Task: for each item, say what is likely being established, what document
satisfies it, and which role holds it.

Constraints: this is general information about how such requests are
framed, not advice on what any standard requires – say so at the top.
Never call an item unnecessary.

Format: a table – Request, Likely purpose, Document, Owner, Effort.

Where you cannot tell what is asked, write "ask them to clarify".

List:
[PASTE]
```

Check before you use the output: the "likely purpose" column. It is inference, and a wrong one doubles the work.

P4.2 — Response to an audit query

Use when: you have the facts and need the reply written cleanly.

Paste this:

You are helping an accountant respond to an audit query on [TOPIC] for [ENTITY], [PERIOD].

Task: draft the response.

Constraints: use only my facts. Never assert a treatment is correct or compliant – state what was done and why, and let the auditor judge. Name no standard.

Format: our position in short; the facts in order; the documents attached; any point that remains a matter of judgement.

If my facts do not fully answer the query, say what is missing.

Query and facts:
[PASTE]

Check before you use the output: any sentence asserting compliance. Replace it with what you actually did.

P4.3 — Control description write-up

Use when: a control exists in practice but was never written down.

Paste this:

You are documenting an internal control at [ENTITY].

Task: write the control description.

Constraints: describe only what I say happens, weaknesses included. Never write the control as it ought to be, and add no threshold or approver I have not stated.

Format: What it is; Who performs it; How often; What evidence it leaves; What it would catch; What it would not catch.

If you cannot complete the last section honestly, ask me questions.

What actually happens, informal parts included:
[PASTE]

Check before you use the output: the "would not catch" section. If it is short or reassuring, it is wrong.

P4.4 — Walkthrough interview questions

Use when: you are about to trace a transaction with its process owner.

Paste this:

You are preparing a walkthrough of the [PROCESS] process at [ENTITY].

Task: write the questions for the person who performs it, designed to reveal where the documented process and the real one differ.

Constraints: open questions only, nothing answerable with yes. Do not presume the process is followed. Assess nothing.

Format: questions in the order the transaction flows, each with a note on what a concerning answer sounds like.

End with the questions that surface undocumented workarounds. Where my description is too vague to build a question on, say so.

Believed process:
[PASTE]

Check before you use the output: that it does not read as an interrogation. People stop describing reality when tested.

P4.5 — Evidence gaps in a testing plan

Use when: you have drafted a testing approach and want its holes found.

Paste this:

You are reviewing a draft testing approach for [AREA] at [ENTITY].

Task: find where the evidence would not support the conclusion drawn.

Constraints: this is a logical review, not a professional or regulatory one; say so at the top. Name no standard.

Format: for each gap – what the test does, what conclusion is drawn, why it does not follow, what would close it.

If it is sound as far as you can tell, say so.

What we test, how, and over what population:
[PASTE]

Check before you use the output: whether it stayed inside logic. Telling you what is required is outside its competence.

5. Tax and compliance preparation

Everything here is preparation, not advice. No output should reach a filing without a qualified person in your jurisdiction reviewing it: tax rules differ between countries, differ within them, and change. Use these to get organised, not to reach conclusions.

P5.1 — Information gathering checklist

Use when: a filing is coming and the documents are scattered.

Paste this:

You are organising a finance team before a [FILING TYPE] for [ENTITY], [PERIOD].

Task: a checklist of documents to gather for whoever prepares the filing.

Constraints: this is a gathering aid and general information, not tax advice; rules differ by jurisdiction and change. Never state what is deductible, taxable or claimable, and reference no law or form.

Format: a table – Activity, Records to gather, Where they live, Owner, Status.

Where treatment depends on local rules, write "for your adviser".

Activities carried out:
[PASTE]

Check before you use the output: that nothing drifted from "gather this" into "this is treated as". Strike any treatment.

P5.2 — Categorisation queries

Use when: you are cleaning a transaction listing before it goes to an adviser.

Paste this:

You are preparing a transaction listing for [ENTITY], [PERIOD], for our tax adviser.

Task: identify transactions described too vaguely for someone else to categorise, and write the question for whoever incurred them.

Constraints: categorise nothing and state no tax treatment; this is preparation, not advice. Never infer purpose from a supplier name.

Format: a table – Reference, Description as given, Why it is unclear, Question to ask, Who to ask.

Leave out anything adequately described. A short list is a good outcome.

Listing:
[PASTE]

Check before you use the output: whether purpose was inferred from vendor names. That assumption unwinds badly later.

P5.3 — Understanding a letter you have received

Use when: correspondence arrives and you need to know what it asks.

Paste this:

You are helping a finance professional understand a letter received by [ENTITY].

Task: say plainly what it appears to ask for, what it requires by when, and what it omits.

Constraints: this is a reading aid and general information, not legal or tax advice; procedures differ by jurisdiction and change. Never tell me how to respond or what my obligations are.

Format: What it asks for; Dates it states; What is unclear.

Quote ambiguous passages rather than resolving them.

Letter:
[PASTE]

Check before you use the output: the deadlines, word for word. A misread date is the one irreversible error here.

P5.4 — Questions for your adviser

Use when: you have adviser time booked and want to use it well.

Paste this:

You are preparing a finance team to meet its tax adviser about [TOPIC] at [ENTITY].

Task: the questions to ask, ordered so the answers build on each other, with the documents to have to hand for each.

Constraints: answer none of them. Never hint at the likely answer or state a rule. This is meeting preparation only.

Format: numbered questions, each with a sub-line "Bring:".

End with "Things we should tell them without being asked". Where my description leaves a gap, ask rather than assume.

Our situation:
[PASTE]

Check before you use the output: the final section — the most valuable part, resting on facts you may have omitted.

P5.5 — Obligations calendar

Use when: deadlines live in several people's heads.

Paste this:

You are building an internal compliance calendar for [ENTITY].

Task: turn the obligations below into a calendar, with the internal milestones before each external deadline.

Constraints: use only my obligations and dates, and correct none of them — you do not know my jurisdiction. List anything you think I omitted as questions.

Format: a table — External deadline, Obligation, Internal milestone, Days before, Owner. Sorted by date.

Open by saying the dates come from me and need confirming.

Obligations, with deadlines as we understand them:
[PASTE]

Check before you use the output: that no deadline was silently adjusted and no obligation added.

6. Contracts, agreements and policy documents

Reading a contract for its finance consequences is not the same as advising on it. These prompts extract and organise; they do not interpret, and none replaces your legal adviser.

P6.1 — Finance summary of an agreement

Use when: a signed agreement lands and you need its effect on the numbers.

Paste this:

You are explaining the financial mechanics of an agreement involving [ENTITY].

Task: summarise only the provisions affecting money, timing, or what must be tracked.

Constraints: this is a reading aid, not legal advice; law differs by jurisdiction. Quote the clause reference and wording for every point, never interpret ambiguous drafting, and state no accounting treatment.

Format: Amounts payable or receivable; Timing and triggers; Variation; Termination and its financial consequences; Obligations to track.

End with "Clauses I could not interpret with confidence", quoting each.

Agreement:
[PASTE]

Check before you use the output: every clause reference. If the references drift, nothing else is safe.

P6.2 — Payment terms extraction

Use when: you are loading terms into a system or comparing a portfolio.

Paste this:

```
You are extracting data from a contract for [ENTITY].

Task: extract the payment terms into a table for system entry.

Constraints: extract only; never normalise into your own categories.
Every value must appear with its clause reference. If a field is not
addressed, write "not stated" – never a default.

Format: a table – Field, Value, Clause. Fields: currency, amount basis,
invoicing frequency, payment period, late payment, price adjustment,
discounts, expenses, disputes.

Then list any payment provision that did not fit.

Contract:
[PASTE]
```

Check before you use the output: every "not stated", and any odd clause reference.

P6.3 — Obligations and dates register

Use when: the contract creates things somebody must remember to do.

Paste this:

```
You are building an obligations register from an agreement involving
[ENTITY].

Task: list every obligation falling on us, with every deadline, notice
period, renewal date and triggering condition.

Constraints: quote the clause reference and wording for each entry.
Calculate a date only where the contract gives both the reference point
and the period. Never interpret conditional wording – quote it.

Format: a table – Obligation, Owner function, Trigger or date, Notice
period, Clause, Wording quoted.

Flag any obligation whose trigger you could not determine.

Agreement:
[PASTE]
```

Check before you use the output: every calculated date. Business-day and month-end definitions break the arithmetic.

P6.4 — Questions for legal review

Use when: legal budget is limited and you want it spent well.

Paste this:

You are preparing for legal review of a draft agreement involving [ENTITY].

Task: find where the draft and the deal I describe diverge, and write the questions for counsel.

Constraints: give no legal advice and never say whether a term is standard or reasonable. Confine yourself to: what I said the deal is, what the words say, the question.

Format: a table – Commercial expectation, Clause and wording, The mismatch, Question for counsel.

Where drafting matches my description, say so rather than invent concern.

The deal, then the draft:
[PASTE]

Check before you use the output: that the expectation column reflects what you said, not a tidier version of it.

P6.5 — Internal policy drafting

Use when: you need a policy people will actually follow.

Paste this:

You are drafting an internal [POLICY TYPE] policy for [ENTITY].

Task: write the policy.

Constraints: include only my rules. Add no thresholds, approval levels or exceptions – if one seems missing, list it as a question. Reference no law. No sentence over 25 words.

Format: Purpose; Who it applies to; The rules as numbered clauses; What to do if they do not cover your situation; Who owns it.

End with "Decisions you have not made yet".

Rules and limits we have decided:
[PASTE]

Check before you use the output: the gaps list, before circulating anything.

7. Board papers, reports and written communication

P7.1 — Board paper structure

Use when: you know what to say and not yet how to arrange it.

Paste this:

You are structuring a board paper on [TOPIC] for [ENTITY].

Task: propose the structure – sections, what belongs in each, roughly how long each should be.

Constraints: structure only. Draft no content, supply no facts, suggest no recommendation.

Format: a numbered outline. For each section: purpose in one line, what goes in it, target length, and whether it belongs in an appendix.

End with "Facts you will need that you have not mentioned".

Background:
[PASTE]

Check before you use the output: the gaps list against your own knowledge — the closest thing here to a rehearsal.

P7.2 — One page from twenty

Use when: the paper is written and the summary is due.

Paste this:

You are summarising the paper below for [AUDIENCE]. The decision required is [DECISION].

Task: a summary of no more than one page.

Constraints: every statement must appear in the paper. Where it qualifies a key point, the qualification must survive – never harden a hedged conclusion. Change no figure.

Format: the decision in one sentence; three to five bullets of what the reader must know; the main uncertainty; next steps.

List any point where the paper is unclear about its conclusion.

Paper:
[PASTE]

Check before you use the output: the hedges. Summarising turns qualified statements into confident ones.

P7.3 — The difficult email

Use when: you must tell someone something they will not like.

Paste this:

You are helping a finance professional write a difficult message.

Task: draft the message.

Constraints: state the substance in the first three sentences – no warm-up. Never apologise for facts, soften a deadline I have given, or add one I have not. No flattery and no exclamation marks.

Format: subject line, then the message, under 200 words.

Then one alternative opening with a different balance of directness. If my description leaves the outcome unclear, ask.

Situation, recipient, outcome I want:
[PASTE]

Check before you use the output: whether a commitment crept in that you did not intend to make.

P7.4 — Anticipating the questions

Use when: you present tomorrow.

Paste this:

You are preparing a presenter for questions on [TOPIC] at [ENTITY].

Task: predict the questions, hardest first – from a summary-only reader and from one who read the appendix.

Constraints: never answer them or draft my answers; I need to know what I do not know. Invent no facts about our business.

Format: three groups – what the paper answers, what it raises but leaves unanswered, what it does not anticipate.

Say plainly where a likely question has no answer available.

Paper and audience:
[PASTE]

Check before you use the output: the middle group — the questions you will be asked and are least ready for.

P7.5 — Plain-English pass

Use when: a document must be read by people outside finance.

Paste this:

You are rewriting a finance document for readers who are competent professionals but not accountants.

Task: rewrite it to be read once and understood.

Constraints: keep every figure and qualification intact. Keep unavoidable technical terms and gloss them in brackets on first use — never a loose synonym, and never make a conditional statement unconditional.

Format: the rewritten document, then the terms you glossed and the gloss used.

Where a passage cannot be plainer without changing meaning, leave it under "Kept as written".

Document:
[PASTE]

Check before you use the output: every gloss. A wrong definition is harder to correct later than the jargon was.

8. Process documentation and team training

P8.1 — Procedure from your own narration

Use when: the process lives only in the head of the person who does it.

Paste this:

You are turning a spoken description of [PROCESS] at [ENTITY] into a written procedure.

Task: write the procedure as numbered steps a competent newcomer could follow.

Constraints: include only steps in the source; add no system names or approvals not mentioned. Where the speaker says "usually" or "it depends", preserve the conditionality. Flag any reordering.

Format: numbered steps starting with a verb; sub-steps where the source gives detail; a "Before you start" list.

End with "Points the description left open".

Notes of the description:
[PASTE]

Check before you use the output: the conditionals. Turning "usually we check" into "check" changes the process.

P8.2 — Desk instructions for a new joiner

Use when: somebody starts on Monday.

Paste this:

You are writing desk instructions for a new joiner at [ENTITY], covering [TASK].

Task: rewrite the procedure below for a first-timer, without our internal shorthand.

Constraints: change no step and no sequence. Expand an abbreviation only where I have said what it means – otherwise mark it [EXPAND THIS].

Format: the instructions, each step with a "what this is for" line, then "When to stop and ask".

List anything in our procedure you could not make sense of.

Procedure, and what the joiner knows:
[PASTE]

Check before you use the output: the "when to stop and ask" section, with whoever does the job now.

P8.3 — Training exercises

Use when: you are teaching a technique rather than a task.

Paste this:

You are designing exercises for a team learning [TECHNIQUE] at [ENTITY].

Task: design [NUMBER] exercises surfacing exactly the mistakes below, in increasing difficulty.

Constraints: use fictional data, clearly labelled. Never present invented figures as real market or benchmark data, and name no real organisation. Each needs a definite right answer or marking list.

Format: for each, the scenario, what the learner produces, the mistake it surfaces, marking notes.

If a mistake cannot be surfaced by an exercise, say so.

Their level and current mistakes:
[PASTE]

Check before you use the output: the marking notes. Vague ones produce arguments, not learning.

P8.4 — Handover note

Use when: somebody is leaving, or going on long leave.

Paste this:

You are building a handover note for [ROLE] at [ENTITY].

Task: organise the notes below into a handover document and identify what is missing.

Constraints: use only what I give you. Never assume responsibilities from the job title or name contacts I have not given. Where I was vague, mark "needs detail".

Format: Recurring tasks by frequency; One-off items in progress; Contacts; Access required; Known workarounds; Things only this person knows.

Prompt me with questions to fill the last section.

What the role does:
[PASTE]

Check before you use the output: the questions it puts to you, and answer them before the person leaves.

P8.5 — Explaining a spreadsheet to its future maintainer

Use when: a model has grown and one person understands it.

Paste this:

```
You are documenting a spreadsheet formula for its future maintainer.

Task: explain what the formula does, step by step, and whether that
matches my intent.

Constraints: work only from the formula text and the layout described.
Assume no ranges or sheet contents I have not given.

Format: one sentence on what it does; a numbered breakdown from the
inside out; then Edge cases – blanks, errors, wrong types, growth.

Say so where an edge case cannot be determined from what I gave.

Formula, layout and intended result:
[PASTE]
```

Check before you use the output: the edge cases, by testing them in a copy of the file.

Building your own

Every prompt above is the same six parts in the same order. Once you see the frame, your own take about two minutes.

Part	What it does	Typical failure if you omit it
Role	Sets who is speaking, and to whom	Generic writing pitched at nobody
Task	States the single job	Several half-done jobs
Context	Supplies the facts only you have	Plausible invention filling the gap
Constraints	Rules out what you do not want	Added figures, added causes, added cheer
Format	Names the shape of the output	Retyping it by hand afterwards
Uncertainty rule	Says what to do when it does not know	Confident, fluent, wrong

The last is the part people leave out, and the part that separates a tool from a liability. "If you are unsure, say so" is the weak version. The strong ones name the behaviour: put unclassifiable items in a separate list; quote the clause you could not interpret; say plainly that you cannot reconcile the components to the total. Give it somewhere to put its ignorance and it will use it. Give it nowhere and it will smooth the gap over with something that reads well.

Two habits are worth forming. Write the check before you write the prompt — decide what would be wrong in the output and how you would spot it. If you cannot answer that, you are not ready to use the output for anything that

matters. And keep your prompts where the team can improve them, with a line under each recording what went wrong last time. A prompt library nobody edits is a library of first attempts.

Verification, Red Lines and Choosing a Tool

Everything in this section is general information, not legal, regulatory, tax or professional advice. Rules differ by jurisdiction, firms differ in what they permit, and both change. Check your own position with your regulator, your professional body and, where it matters, your own advisers.

Part One — The Three-Check Rule

You already have a review habit, tuned over years to the mistakes humans make: transposed digits, a stale opening balance, a paragraph copied from last year and not updated. AI output fails differently. It fails smoothly. The wrong number is formatted correctly, sits in the right column, and is surrounded by three that are right. The invented source has a plausible title. The reasoning reads like reasoning.

So the habit needs a small extension, not a replacement. Three checks, in order, every time output leaves your hands: **the figures, the sources, the reasoning.**

The order matters. Figures first: a wrong figure makes everything downstream irrelevant. Sources second: that is where fabrication hides. Reasoning last, because it takes the most thought, and you should not spend that on a document that has already failed.

It should take a few minutes on short work and under half an hour on something substantial. If it would take longer, the work was the wrong shape to give to a tool.

Check one — the figures

What you do. Take every number and trace it back to where it came from. Not "does it look right" — trace it. Open the source file. Find the figure. Match it. Then check the arithmetic between figures: components add to the total, the variance equals the difference, percentages are of the base you think, periods compared are the same length.

Numbers you did not supply get harder treatment. A figure that was not in your input came from somewhere you cannot see. Delete it or verify it independently. There is no third option.

What it catches. Numbers altered in transit — rounded when they should not have been, pulled from the wrong row, carried from one entity to another, restated to a different currency without saying so. Totals that do not foot. Percentages calculated the wrong way round. Prior-period figures that quietly became current-period. Numbers invented to fill a sentence that needed one.

What it misses. A number faithfully copied from a source that is itself wrong: the check verifies transmission, not truth. It also misses the number that should be there and is not — an omitted provision, a missing intercompany line, a segment left out of a total you did not think to foot. Absence is harder to see than error.

How long. A minute or two per page of numbers. If you cannot trace a figure inside about thirty seconds, that is not a slow check, it is a finding — write it down and go back to whoever produced it.

A prompt that makes this check faster, because it forces the output to show its own working:

Before the answer, produce a table of every number you are about to use. Columns: the number, the exact label or cell reference it came from in my material, and one word – "supplied" if it came from my material, "derived" if you calculated it. For anything derived, show the calculation. Do not include any number you cannot place in one of those two categories. If a figure is missing from my material, say so and leave it blank rather than estimating it.

Material: [PASTE OR ATTACH]

Task: [WHAT YOU WANT PRODUCED]

You still trace the figures, but against a list the tool has committed to, which is faster than hunting through prose. The "derived" rows tell you which calculations to re-perform.

Check two — the sources

What you do. For every factual claim that is not a figure — a rule, a treatment, a deadline, a definition, a contractual term, "the standard requires", "market practice is" — find the source and read it yourself. The primary document, not a summary of it. If the output names a source, confirm it exists, confirm it says what the output says it says, and confirm it is current.

If the output does not name a source, the claim has no source. Treat it as a hypothesis to prove or drop.

What it catches. Confident references to things that do not exist, or that exist but say something else. Rules described as they were years ago. Guidance from one jurisdiction presented as though it applies in yours — easy to miss, because the language of finance is shared across borders while the rules are not. Thresholds that have moved. Requirements that apply to a class of entity yours is not in.

What it misses. A source that is real, current, correctly quoted and beside the point — relevance is a judgement no check automates. Also the source that is right for your jurisdiction and wrong for your client's, and the well-known position with an exception your client sits inside.

How long. Budget two to five minutes per claim you actually rely on. That sounds heavy until you notice how few claims in a document carry weight. Mark the ones that do — the ones a reader would act on, or that would appear in a complaint — and check those properly. Skimming twelve claims is worth less than checking three.

A useful move here is to separate the drafting from the sourcing:

List every factual assertion in the text below that is not a number, one per line. For each, mark it "verifiable" if a reader could check it against a document, or "judgement" if it is an opinion or an inference. For each verifiable assertion, state what kind of document would settle it. Do not tell me the answer, do not cite anything, and do not add assertions that are not already in the text.

Text: [PASTE THE DRAFT]

That gives you a checking list rather than a citation list — more honest, because it does not invite the tool to produce references, which is where fabrication lives.

Check three — the reasoning

What you do. Close the output. In your own words, state the argument: the conclusion, what it depends on, and what would have to be true for it to hold. Then reopen it and compare. You are looking for three things — a step asserted rather than argued, an assumption never stated, and a conclusion wider than the evidence under it.

Then invert it: what is the strongest argument against the conclusion? If you cannot construct one, you do not yet understand the position well enough to send it.

What it catches. Confident conclusions built on one unrepresentative period. Correlation dressed as cause — "margin fell because of the pricing change", when three things changed that quarter. Recommendations that quietly assume the reader's circumstances. Analysis that answers a slightly different question from the one asked — the most common failure, and the hardest to see, because the answer is good.

What it misses. Anything that depends on knowing the client, the history, the personalities and the thing the last auditor flagged. Reasoning that is sound in general and wrong here. And your own blind spots — you are checking with the head that approved the approach.

How long. Five to ten minutes on anything carrying a recommendation; less on descriptive work. It is the check people skip when busy, and the one that catches the errors that embarrass people: figure errors get found by the next reviewer, reasoning errors get found by the client.

The prompt that helps here is the one that argues with you:

Argue against the conclusion in the text below. Give the three strongest objections a well-informed sceptic would raise, the assumption each attacks, and what evidence would settle it. Do not soften them, do not offer a balanced view, and do not rewrite the text.

Text: [PASTE THE DRAFT]

Read the objections as prompts for your own thinking, not as findings. Some will be weak. The point is to make the shape of the argument visible so you can test it.

The verification card

Print this and pin it where work gets reviewed. It is deliberately short: a checklist nobody completes is worse than none, because it produces the paperwork of assurance without the assurance.

THE THREE-CHECK CARD

Before this leaves my hands.

1. FIGURES — trace, do not glance

- Every number traced to a named source file, cell or document
- Totals footed; variances re-performed; percentages checked against the right base
- Periods, entities and currencies confirmed to be the ones stated
- Any number I did not supply: verified independently, or deleted

2. SOURCES — read the original, not the summary

- Every claim I am relying on has a source I have opened myself
- The source is current, and applies in this jurisdiction and to this entity
- Unsourced claims: proved or removed
- Nothing quoted that I have not read in full

3. REASONING — *state it without looking*

- I can state the conclusion and what it depends on, in my own words
- Every assumption is on the page, not in my head
- The conclusion is no wider than the evidence under it
- I can make the strongest case against it — and it still stands

SIGN-OFF

- This answers the question that was actually asked
- I can defend every line of it, unaided, to someone who challenges it
- Where I am uncertain, the document says so
- If this is wrong, I am the one who is wrong. That is correct.

Checked by: [NAME] Date: [DATE] Time taken: [MINUTES]

Log the time taken — for you, not anyone else. After a fortnight you will know which work is genuinely faster with a tool and which is only faster until the checking starts. It is the only productivity figure in this pack, and the only one worth having, because you measured it.

What the three checks do not catch

They do not catch everything, and a pack claiming otherwise would be the sort of thing this pack exists to be an alternative to. What remains after all three checks are done properly:

The plausible omission. The checks look at what is on the page, not at what should be there and is not. A tool that summarises fifty documents and silently drops the one that matters passes every check.

The correct-but-inapplicable. Right rule, right numbers, wrong client. Nothing in the checks tests fit.

Compounding. Verified output becomes input to the next piece of work, and the check does not repeat. Three steps on, the assumption is load-bearing and nobody remembers it was one.

Drift. The tenth document of the day gets a thinner check than the first, and reviewing fluent output is not the same mental act as writing it. Read enough finished-looking work and the reading gets lighter without your deciding it should.

Three things to do about the residual risk, none of which eliminates it:

1. Record your assumptions, and keep the input. An assumptions block

at the top of anything you pass on, with the material you supplied stored beside the output. If you cannot reconstruct what the tool was given, you cannot work out how it went wrong.

1. Vary who checks. Whoever framed the task is the worst person to

check it; the second-worst is anyone who agrees with them.

1. **Set a ceiling.** Decide in advance, in writing, the value or

sensitivity of work above which the tool is not used at all. Deciding under deadline is deciding badly. Part Two is a starting draft.

Part Two — Ten Places Not to Use It

There is a way of writing about this that treats every limitation as a temporary inconvenience soon to be engineered away. This is not that. Some of what follows may change; most will not, because the reason is not capability. It is that a duty sits on a person, and cannot be delegated to something that cannot be held to it. For each: the situation, why the tool fails there, and what to do instead.

1. Signing, certifying or expressing an opinion

The situation. Anything where your name, signature or firm goes on a document as the person vouching for it — accounts, certificates, confirmations, formal opinions, representations a third party will act on.

Why it fails. Your signature is not a statement that the words are good. It says *you* did the work and *you* stand behind it. The signature is about accountability, not text, and a tool cannot supply accountability. Sign something you did not verify and the signature is false, whether or not the content turns out to be right.

Instead. Use a tool, if at all, on drafting and checking well before signature — never on the substance of what is certified. If a tool touched the draft, the signatory must have re-performed the underlying work, not reviewed the prose.

2. Anything you cannot verify

The situation. The output is in a subject, language, jurisdiction or technical area where you cannot check whether it is right.

Why it fails. Every safeguard in Part One assumes a checker who can tell good from bad. Remove that and the tool is not an assistant, it is an oracle. And fluency scales with unfamiliarity: output in an area you do not know reads *more* convincing, because you cannot see the joins.

Instead. Either learn enough to check it, find someone who can, or do not send it. A tool used to orient yourself in an unfamiliar area gives you things to look up, never the answer. The tell: if you could not answer a follow-up question without going back to the tool, you do not own the work.

3. Live disputes, investigations and privileged matters

The situation. Litigation, arbitration, a regulatory investigation, a disciplinary process, a dispute with a former employee — anything where a lawyer is involved or about to be.

Why it fails. Two reasons, both serious. In a contested matter the working papers may become visible to the other side, and a tool-produced draft — including discarded versions and the framing of your questions — records how you thought about the case. And protections attaching to communications with your lawyers can be fragile; routing the same material through a third-party system may affect them, depending on your jurisdiction.

Instead. Take instructions from the lawyers on the matter, in writing, before anything goes near a tool. Assume until told otherwise that nothing about a live dispute goes into any external system. This is general information, not legal advice; the rules here are specific to your jurisdiction, and consequential.

4. Decisions about a named individual

The situation. Pay, promotion, redundancy, hiring, performance ratings, discipline, creditworthiness, a reference, a decision to end someone's contract.

Why it fails. The decision has to be explicable to the person affected, and possibly to a tribunal or a regulator; "the system suggested it" is not an explanation. The tool cannot see the context that matters — the year they had, the thing their manager did, the reason the numbers look like that. And several jurisdictions place obligations around automated processing of personal data and decisions affecting people; those vary, and you need to know your own.

Instead. Make the decision yourself, on evidence you can point to, and write your reasons before drafting any communication. A tool may make an already-made decision clearer or kinder in the telling. It must not help you reach it, and it must not receive the person's personal data.

5. Anything where the client's data cannot leave your control

The situation. Data covered by a confidentiality undertaking, a data processing agreement naming permitted sub-processors, a client policy prohibiting external processing, or a requirement to keep information in-jurisdiction. In practice also: personal data, bank details, board material before it is public, price-sensitive information.

Why it fails. Not because the tool is careless. Because pasting the data into it is a disclosure to a third party, and you either had permission for that or you did not. The tool's own security is not the question; whether you were allowed to send it is, and that is settled by contracts you signed before the tool existed.

Instead. Check the engagement terms and data agreements before the work starts, not during it. Where the answer is no, either work without the tool or use one deployed inside your own environment under terms you have read. When in doubt, redact: much analytical work runs on structure and anonymised figures, without names, account numbers or identifiers leaving your systems.

6. Matters where a professional duty attaches to you personally

The situation. Independence and conflict assessments. Ethical judgements. Whether to accept or resign from an engagement. Whether something must be reported, and to whom. Anything your professional body treats as a matter for your own judgement.

Why it fails. These duties sit on you as a person precisely because they exist to be exercised when it is uncomfortable — when the client is large, the fee is significant and the easy answer is available. A tool produces the conventional answer, usually the comfortable one, and has nothing at stake. That is the problem: the duty exists because you do.

Instead. Decide it yourself, document the reasoning at the time in your own words, and consult a human — a partner, an ethics helpline, your professional body — where you are unsure. The written record of your reasoning is the point, and it has to be yours.

7. Suspicions of fraud, error or dishonesty

The situation. Something looks wrong and a person may have done it deliberately.

Why it fails. These situations carry reporting obligations that vary by jurisdiction and can be time-critical, and some carry prohibitions on telling anyone what you suspect. A drafting tool cannot know your obligations or who you are permitted to speak to. There is also a practical trap: writing up a suspicion in a system you do not control creates a record you may not be able to retrieve, delete or explain.

Instead. Follow your firm's escalation procedure immediately — and if you do not know it, find out today rather than the day you need it. Take advice on your reporting obligations from someone qualified to give it. Keep your own note, written at the time and held locally. Nothing external, nothing informal.

8. Advice that turns on the specific rules of a place

The situation. Tax positions, filing obligations, statutory deadlines, entity-specific reporting requirements, licensing — anything where the answer changes when you cross a border or a threshold.

Why it fails. These rules are numerous, local and revised often. Material describing them ages badly and is uneven across jurisdictions, so the answer reflects the places most written about rather than the place you are in. That is dangerous, because the *structure* of the answer is usually right — the right considerations, in the right order — while the threshold, rate or date inside it is wrong or foreign.

Instead. Use the primary source published by the relevant authority, or a subscription service that states its currency date. A tool may help you frame the question or draft the covering letter. The rule itself comes from the source every time, and you note the date you checked.

9. The record of what happened

The situation. Minutes, file notes, attendance notes, records of what was said and decided.

Why it fails. A record is not a summary. Its value is that it is a first-hand account by someone present — exactly what a generated version is not. A tool asked to tidy up rough notes smooths them, fills gaps with what usually happens, and turns the awkward half-sentence someone said into a clean sentence they did not. In a dispute that difference is everything, and the smoothing is invisible once done.

Instead. Write records yourself, close to the event, in plain language, including the parts that are unresolved or unflattering. If you use a transcription tool, keep the transcript as the record and mark it as such — and check what your policy and local law say about recording people, which is a real question, not a formality.

10. Anything you would not be able to defend line by line

The situation. The general case. You are about to send work you could not walk through, unaided, if the recipient rang and asked why.

Why it fails. Not a limitation of the tool. A limitation of the arrangement. Work you cannot defend is work you have not done, and whether a tool, a junior or a template produced it changes nothing about your position.

Instead. Either do the work until you can defend it, or say plainly that you have not. "I have not verified this — it is a starting point" is a perfectly professional sentence. Most of the trouble people get into with these tools starts with an unwillingness to say it.

Part Three — Assessing a Tool Before It Touches Client Data

Marketing pages are written to be reassuring. Contracts are written to be enforceable. They routinely say different things. This section is about getting answers in the second kind.

Work through the table below with the vendor, or with your own IT team if the tool is deployed internally. Two rules make it worth doing. Get the answers in writing, with the clause reference: "yes, of course" is not an answer, "clause 8.3 of the data processing addendum" is. And if a question comes back vague after you have asked twice, record the vagueness as the answer. It is information.

Fill in the third column with what they said and the fourth with your own decision — a judgement your firm makes, not one a checklist makes for you.

What happens to our input

#	Question to ask	Their answer	Acceptable?
1	Is anything we submit used to train or improve your models, or anyone else's?		
2	If so, how do we switch it off, is the switch on our plan, and does it cover every feature?		
3	Which third parties process our input — model provider, host, subcontractors? Name them all.		
4	Do you keep our input for safety, abuse or quality review? Who sees it, and for how long?		
5	Does any of this differ for uploaded files or connected data sources, versus typed text?		

Where it is stored, and for how long

#	Question to ask	Their answer	Acceptable?
6	Which countries is our data stored, processed and backed up in? List them all.		
7	Can we require a named region, and is that contractual or a setting you can change?		
8	How long is our input retained by default, and can we shorten it?		
9	How long does it persist in backups, logs and caches after deletion?		
10	Is it encrypted at rest and in transit, and who holds the keys?		

Who at your end can see it

#	Question to ask	Their answer	Acceptable?
11	Which of your staff can access our content, in what circumstances, on whose approval?		
12	Is that access logged, and can we see the log?		
13	Do your staff sign confidentiality undertakings covering our clients' information?		
14	If a government or law-enforcement request reaches our data, will you tell us where permitted?		
15	Have you had an incident affecting customer content, and what is your notification period?		

What happens when we leave

#	Question to ask	Their answer	Acceptable?
16	On termination, what is deleted, when, and how is deletion evidenced?		
17	Can we export content, history and configuration in a usable format, as a contractual right?		
18	What survives termination in your favour: licences to our content, aggregated data, derived material?		
19	If you suspend or terminate our account, do we still get our data out, and for how long?		
20	If the service is discontinued, what notice do we get?		

The contract versus the marketing page

#	Question to ask	Their answer	Acceptable?
21	Which document governs — terms, data processing addendum, security page, sales deck? Show the order of precedence.		
22	Which of the claims on your website appear as obligations in the contract? Name the clauses.		
23	Can you change the terms unilaterally, with what notice, and can we exit without penalty if you do?		
24	What is your liability cap, and what is excluded from it?		
25	Do you accept the role of processor, on our instructions, in writing?		
26	Do you hold current independent security certification, and may we see the report, not the badge?		

What changes if you are acquired

#	Question to ask	Their answer	Acceptable?
27	On a change of control, do our terms transfer unchanged, and for how long?		
28	Is our data an asset that could be sold or transferred in insolvency or acquisition?		
29	Do we get notice and a right to terminate on a change of control?		
30	On a funding or ownership change, what happens to the commitments in questions 6 to 9?		

Two notes on using this. Question 22 produces the most silence, and the silence is the finding. And a sensible way to handle the paperwork is to have a tool extract, not interpret:

Read the attached terms and produce a table with three columns: the question below, the exact wording from the document that addresses it, and the clause number. If the document does not address a question, write "not addressed" – do not infer or fill the gap. Quote verbatim. Do not tell me whether the terms are good.

Questions:

[PASTE THE QUESTIONS FROM THE TABLE ABOVE]

Document: [ATTACH THE TERMS]

That gives you a map of where to read. You still read the clauses yourself, and anything that matters commercially or legally goes to someone qualified to advise on it. Extraction is a task a tool does well. Judging a liability cap is not.

The assessment has a shelf life

An assessment describes a supplier on the day you did it, and nothing more. Terms are revised. Sub-processors are added. Features arrive that handle data differently from the feature you assessed, and they arrive switched on. Free and paid tiers can differ on exactly the points you cared about, so an upgrade changes your position. Companies are bought. Retention settings get reset by a migration.

None of this is bad faith; it is what happens to software companies. The consequence is that a one-off assessment filed away is a record of a decision, not evidence of a current position.

So repeat it. Once a year for anything touching client data, and immediately on any of these triggers: the terms change, the ownership changes, you start using a materially new feature, your plan changes, a client asks where their data goes, or you take on work under a stricter confidentiality regime than usual. Keep the completed tables, dated and signed by whoever filled them in, so the next review is a comparison rather than a fresh start.

And keep one thing in your own hands whatever the assessment says: know which of your clients' information must never be submitted to an external system at all, write that list down, and make sure everyone who could paste something has read it. Vendor assurances describe what someone else promises to do. That list describes what you control.

Your Policy Documents

What is in this section

Three documents, in ascending order of length:

1. **The one-page team note** — for a team of two to fifteen who need a rule by Friday.
1. **The firm policy** — for a practice with clients, engagement terms and a regulator.
1. **The client-facing paragraph** — the wording you put in an engagement letter or a proposal when a client asks how you use these tools.

Take whichever fits, put your own name in the brackets, delete what does not apply, and have someone who carries risk read it before you circulate it. All three are starting drafts and general information, not legal advice. Rules differ by jurisdiction and profession and they change.

There is also a free web version of the one-page note that assembles it from nine questions. It is on the same site this pack came from, under Free tools. It runs entirely in your browser and collects nothing.

1. The one-page team note

USING AI TOOLS — [TEAM OR FIRM NAME]

Applies to: everyone in the team, including contractors. **Owner:** [name and role]. **Next review:** [date].

Why this exists. We use these tools because they save real time. They also produce confident, well-written output that is sometimes wrong. This note exists so the time saved is not paid back later with interest, and so nobody has to guess what is allowed.

Approved tools. [List them.] Anything not on this list is not approved yet — that is not a refusal, ask and the list can grow. Personal accounts and free consumer tiers are not approved for any work involving client material.

What may go in. [Choose one: no client or personal data at all / client data only in the approved tools whose terms state our input is not used for training / anonymised material only.] Identification numbers, bank details and payroll figures for named individuals stay out regardless of the tool. If you are unsure whether something counts, it counts. Ask first.

What you check before it goes out. Every figure, every quotation, every reference, every claim about a rule — against its source. Recalculate anything numerical independently. Never cite a standard, a case or a section you have not opened yourself.

What we tell clients. [Choose one: internal aid, not routinely disclosed / we answer honestly and specifically when asked / we state it in engagement terms.] We never deny using these tools.

Where they are not to be used. Signing, approving or certifying anything. Any final decision about a person's pay, employment or standing. Anything under privilege, live dispute or investigation without written approval. Anything you would not be able to check.

Records. Where a tool materially helped produce a deliverable, note it in the working papers: which tool, what for, who checked it. One line is enough.

If something goes wrong. Tell [owner] the same day. Say plainly what happened. Nobody is disciplined for reporting quickly. Concealment is the only thing treated seriously.

2. The firm policy

Longer, for a practice that has clients and a regulator. The clauses below are drafting blocks — keep the ones that apply.

2.1 Purpose and scope

This policy governs the use of artificial intelligence tools in the work of [FIRM]. It applies to all partners, employees, contractors and temporary staff, on any device, whether the tool is provided by the firm or not. It covers general-purpose assistants, AI features embedded in software the firm already licenses, transcription and note-taking tools, and any tool that sends firm or client material to a third party for processing.

It does not change any existing duty. Where this policy and a professional obligation appear to conflict, the professional obligation governs and you raise it with [ROLE] immediately.

2.2 Responsibility

The person who signs, reviews or issues a piece of work is fully responsible for it. That responsibility is not reduced, shared or transferred by the use of any tool. "The system produced it" is not an explanation the firm will offer to a client, a regulator or a court, and it is not one you may offer internally.

2.3 Approved tools and how the list changes

Only tools on the approved register may be used for client work. The register is maintained by [ROLE] and records, for each tool: the provider, the licence tier, whether firm input may be used to train the provider's models, where data is processed, the retention period, and the date the terms were last reviewed.

Requesting an addition is normal and encouraged. Using an unapproved tool on client material is a disciplinary matter, not because the tool is suspect but because the firm cannot answer for what it cannot see.

2.4 Client and personal information

[Choose the stance that matches the firm's risk appetite and delete the others.]

Option A — prohibition. No client information, personal data or confidential material may be entered into any AI tool. This includes names, initials, entity names, account numbers, identification numbers, addresses and figures attributable to a named person or client, and it includes combinations of details that would identify someone together even if each is innocuous alone.

Option B — approved tools only. Client information may be entered only into tools on the approved register whose terms state that firm input is not used to train the provider's models and is not retained beyond the period recorded on the register. It may never be entered into a personal account or a consumer tier.

Option C — anonymised only. Only material from which identifying details have been removed may be entered. Where anonymising would take longer than performing the task unaided, perform the task unaided.

In every case: identification numbers, bank account details, payroll figures for named individuals, health information and anything received under an explicit confidentiality undertaking stay out.

2.5 Verification

Before AI-assisted work leaves the person who produced it:

(a) every figure is recalculated or agreed to source; (b) every quotation, reference, standard and authority is opened and read in the original; (c) the reasoning is read as though a reviewer were trying to break it, with particular attention to anything stated confidently and without qualification; (d) [for higher-risk work] a second qualified person reviews the deliverable and records that they have done so.

Fluent output is not verified output. These tools are at their most persuasive on the things they have invented.

2.6 Disclosure

[Choose one.] The firm does not deny the use of these tools in any circumstances. Where a client asks, the answer given is honest and specific: what the tool was used for, what it was not used for, and who checked the result.

2.7 Prohibited uses

AI tools are not to be used to: sign, certify, approve or give an opinion; make a final decision about any individual's employment, remuneration, credit or professional standing; process any matter under legal privilege, live dispute or regulatory investigation without the written approval of [ROLE] for that specific matter; produce anything the user would be unable to verify; or generate material presented as the independent work of a person who did not review it.

2.8 Records

Where an AI tool materially contributed to a deliverable, the working papers record the tool, the purpose, and the reviewer. This is not bureaucracy. It is what allows the firm to answer a question about the work two years after the engagement closed.

2.9 Incidents

Any suspected error in issued work, or any client material entered somewhere it should not have been, is reported to [ROLE] on the day it is noticed. Prompt reporting attracts no sanction. Concealment does.

2.10 Review

This policy is reviewed every [3/6/12] months by [ROLE], and sooner if a tool on the register materially changes its terms, if the firm's professional body or regulator issues guidance, or if an incident suggests a clause is not working. Superseded versions are retained.

3. The client-facing paragraph

For engagement letters and proposals. Short on purpose — clients want to know you have thought about it, not to read your internal policy.

Use of technology. We use software tools, including tools that use artificial intelligence, as an aid in performing our work. They are used under a written internal policy which restricts what information may be processed, requires that outputs be verified against source before use, and prohibits their use for judgements and approvals. The individuals named in this engagement remain fully responsible for the work and for every conclusion in it. We are happy to describe our approach in more detail on request.

If your engagement terms already contain a confidentiality or subcontracting clause, check that this paragraph sits consistently with it rather than beside it.

Before you adopt any of this

These are drafts assembled from general practice, not legal advice, and not a statement of compliance with any law, standard or professional requirement. Requirements differ between jurisdictions and professions and they change. Have whoever carries risk in your organisation read the version you intend to use, and check it against your regulator's current guidance and your own client engagement terms.

The 90-Day Plan and the Training Curriculum

Ninety days is long enough to change how a finance team works and short enough that one person can hold the whole thing in their head. What follows is a plan for that period, then a curriculum you can teach in your own meeting room with no outside help and no budget beyond people's time.

It assumes a finance function of between five and fifty people that already has access to some tool. Choosing the tool is a separate exercise; this is what happens after the licence exists.

None of the numbers here are benchmarks, because there are none worth quoting — the only numbers that will matter are the ones your team records. This is general guidance only: anything touching confidentiality, client consent, professional standards, data protection or your engagement terms differs by jurisdiction and changes over time. Check your own position with your professional body, your regulator and whoever handles your data obligations.

Part one: the ninety-day plan

What you settle before day one

Four decisions, in writing, in about an hour. Skipping them causes most of the mess later.

Who owns this. One named person, not a committee, and not your busiest month-end resource. They need enough standing to say "not yet" to a partner.

What you will measure. The time a task takes end to end, and whether the output needed rework. Not logins, not prompt counts. Enthusiasm is a mood, not a measure, and it evaporates around the twentieth of the month.

The one rule that applies immediately, before any policy exists: nothing produced with an assistant leaves the team — not to a client, a board, a regulator or a colleague outside finance — without a named human who has checked it and will put their name on it. That can go out in an email this afternoon.

What data is not going in until the policy exists: client-identifiable information, personal data, anything under an NDA, anything price-sensitive. Whether that starting position is right for you depends on your tool's terms, your contracts and your local rules. Read the terms.

The people involved: a sponsor who approves the pilot and protects the time (an hour a month); the owner (half a day a week); three to five pilot users (two or three hours a week each); a reviewer, within their normal review time; and whoever handles systems and contracts, for a couple of hours in month one.

Phase one — Days 1 to 30: contain it, then choose

The goal. A policy people have read, and two chosen use cases with a baseline. You are not trying to produce results this month. You are trying to stop the uncontrolled version of this happening while you think.

Week 1 — find out what is already happening. The owner asks every person, individually and without consequence, what they have used and for what. In person, not in a form. Say plainly that nobody is in trouble and that you need the truth to write a sensible policy. Respond to what you hear with a fix, not a disciplinary — this is the

only week you get the truth cheaply. *Who*: owner. *Done when*: you have a written list of current use, and the whole team has the one rule above.

Week 2 — draft the policy. Two pages at most; nobody reads four. It covers: approved and unapproved tools; what data may never be entered; that a named human reviews everything before it leaves; how that review is recorded; and who to ask when the policy does not cover the situation. Have your data/IT contact check it against the tool's terms. If you owe clients confidentiality, get the consent question answered rather than assumed. *Who*: owner drafts, sponsor approves, data/IT checks. *Done when*: signed and sent, with a request to reply "read".

Week 3 — choose two use cases. Two, not ten. Starting ten produces ten half-finished experiments and evidence about none of them. Four tests: the task happens at least monthly; it eats time you can name; a wrong answer is caught before it does damage; and you can tell within a week whether the output was good. Good early candidates are drafting and summarising work that already has a review step around it. Poor ones are anything where an error goes straight out. *Who*: owner and pilot group; sponsor breaks ties. *Done when*: two use cases written in one sentence each, with a named reviewer for each.

Week 4 — measure the "before". Record how long each use case takes now and how much rework it needs, using the log below. The baseline is an estimate: people recall duration poorly, and measuring changes the number. Label it as an estimate wherever you quote it — a soft baseline you admit is soft beats a precise-sounding figure you invented. *Who*: pilot group. *Done when*: three baseline observations per use case, and the pilot group has done Level 1 of the curriculum.

Date	Task	Who	Minutes, unassisted (est.)	Minutes, assisted	Review minutes	Rework needed?	What went wrong
------	------	-----	----------------------------	-------------------	----------------	----------------	-----------------

How this phase fails. Someone senior sees a demonstration and the scope grows in a fortnight from two use cases to "a strategy". You end up with a deck, no baseline and no policy, while the team carries on as before. The defence is the written pair of use cases: everything else goes on a list marked "after day 90", and the list is shown, not argued with. A policy written to be safe rather than followed fails the same way: if it turns the honest answers you got in week 1 into sackable offences, people will comply by not telling you things.

Phase two — Days 31 to 60: run a real pilot

The goal. Evidence: thirty or so logged observations, and an unexcited sense of whether either use case is better, worse or the same.

Week 5 — train and start. Pilot group completes Levels 1 and 2, then uses the assistant on the two chosen use cases only, logging every run including the abandoned ones. Abandoned runs are data; dropping them is how a pilot lies to you. *Who*: owner trains, pilot group runs. *Done when*: two logged runs each.

Week 6 — the first honest look. Thirty minutes with the pilot group. Not a status update: ask each person for the worst output they got this week and what they had to fix, and the reviewer for what they corrected before sign-off. Write the failure types down. You are building a list of what this tool gets wrong *on your work*, which is worth more than any general warning. *Who*: owner chairs, reviewer contributes. *Done when*: a written failure list has been circulated.

Week 7 — tighten, do not widen. Improve the prompts, not the scope. The pilot group does Level 3 this week. If a use case is clearly not working, say so now and either fix the prompt once or drop it. Dropping one of two at week 7

is normal, and far cheaper than at week 12. *Who*: pilot group. *Done when*: one agreed prompt per use case, saved where everyone can reach it, placeholders written out.

Week 8 — the honest read. Compare the log against the baseline and answer three questions in writing. Did the task get faster *once review time is included*? A draft produced in two minutes that takes forty to check has not saved you thirty-eight. Did rework go up, down or stay flat? Would the reviewer be comfortable with this going out under their name every month? Good answers mean widen; mixed answers mean extend the pilot a month; poor answers mean stop that use case and say why, in a paragraph, to everyone. *Who*: owner writes, sponsor reads. *Done when*: a one-page finding exists with the log attached.

How this phase fails. The pilot group logs the successes and forgets the failures, usually without meaning to, and by week 8 the data says the pilot went beautifully while the reviewer is quietly rewriting everything. Two defences: ask for the worst output every week, and treat the reviewer's correction time as the real measure, because it is the one number nobody can flatter. Watch also for the review step being skipped under deadline pressure once output is "usually fine". Make sign-off a recorded step, not a habit.

Phase three — Days 61 to 90: widen only what held

The goal. Move from a pilot that worked to a routine that survives the owner going on leave.

Week 9 — write the run sheets. One page per surviving use case: the prompt with placeholders, what to attach, what to check before accepting the output, who signs off, where the result is saved. If it does not fit on a page, the workflow is too complicated to hand over. *Who*: pilot group writes, owner edits. *Done when*: one page each, stored with the policy.

Week 10 — teach the next group. The pilot group teaches Levels 1 and 2 to the next cohort, with the owner present. People who did it teaching people who have not is more convincing than any external session, and it exposes gaps in your run sheets fast. *Who*: pilot group. *Done when*: the second cohort has run the workflow once each, supervised.

Week 11 — a deliberate stress test. Run one period-end or one job with the new workflow and the owner not intervening. Log it. What matters is the list of questions the run sheet did not answer. *Who*: the new cohort. *Done when*: the run sheet has been amended.

Week 12 — review the policy against reality. Update the approved tools list, the data rules that proved unworkable or too loose, and the failure types you found. The two-page limit still applies. *Who*: owner and data/IT contact. *Done when*: version two is signed and re-circulated.

Week 13 — decide what happens next. One page to the sponsor: what you tried, what the log says, what you have stopped doing, the next two use cases, and what you are deliberately not doing yet. Add who keeps the log, and when it and the policy are next reviewed. *Who*: owner. *Done when*: the sponsor has the page and the next review date is in a calendar.

How this phase fails. Everything widens at once. The pilot held for two use cases with five trained, watched people; it gets treated as general permission, quality drops, and the whole thing is banned in month five after one bad incident. Widen the use cases you tested, with the training you tested. Watch also for the quieter failure: if nobody but the owner can run the workflow from the run sheet, you have a hobby, not a process.

The senior who will not engage

You will have one, often your best technical reviewer, which is not a coincidence — they got good by not trusting things.

Do not put them in the pilot group and do not sit them in a demonstration. Both read as pressure and produce a public position they then defend for a year. Give them the role they are already good at: reviewer. Ask them to check the pilot output and tell you what is wrong with it. That uses their scepticism as an asset and puts them in contact with the quality of the output rather than the claims about it. Some people who start out opposed adopt parts of it once they know where it fails. Some do not, which is fine if they are not blocking others.

Hold firm on two things. They do not get to veto other people's use inside the policy. And they do get to say "not on my client work" for a period if they carry the professional responsibility for it — that is judgement, not obstruction. Where work is regulated, whose name is on the file matters more than your adoption plan does.

What does not work: making it a performance objective, mocking them in front of juniors, or going round them to their staff.

The junior who has already raced ahead

You will have one of these too: competent, recently qualified, months of private practice with an assistant, genuinely faster at some things.

Your instinct will be to slow them down. Half-resist it. What you need first is honesty about what they have already put into which tool, because you cannot manage exposure you do not know about. Ask directly, once, with the no-consequences framing of week 1. Then give them a real job: writing prompts and run sheets, and teaching Level 2. That turns private practice into a documented team asset.

The line you hold is review. The risk is not that they use the tool; it is that they produce more finished-looking work than the reviewer's normal capacity. Say this to the reviewer out loud: the drafts now look better than they used to, and the appearance of the work no longer tells you much about the care behind it. Check the numbers and the sources, not the prose.

If they have used an unapproved tool on client data, treat it as an exposure to assess and report per your own obligations, not as a character question. Whether it must be reported anywhere is a matter for your professional body's rules and your local law. Ask; do not assume.

The ninety days on one page

Week	Phase	What happens	Who leads
1	Contain and choose	Ask everyone what they already use; issue the one rule	Owner
2		Draft the two-page policy; check tool terms	Owner + data/IT
3		Choose two use cases against the four tests	Owner + pilot group
4		Record the "before", labelled as estimates	Pilot group
5	Pilot	Levels 1–2 training; pilot starts; every run logged	Owner
6		"Show me the worst output"; list failure types	Owner + reviewer
7		Level 3; tighten prompts; drop what is not working	Pilot group
8		Compare against baseline; widen, extend or stop	Owner + sponsor
9	Widen what held	Write one-page run sheets	Pilot group
10		Teach Levels 1–2 to the next cohort	Pilot group
11		Stress test a real period-end; owner hands off	New cohort
12		Update the policy against what happened	Owner + data/IT
13		Write up; set the standing review; pick the next two	Owner

Part two: a four-level training curriculum

This is designed to be taught by a manager one step ahead of the room. Each level stands alone, and most of your team will get most of the benefit from Levels 1 to 3. You need a room with a screen, a working licence, and real examples from your own work with client-identifiable details removed — invented examples do not land with this audience.

Teach in order. The common mistake is starting at Level 2, because prompting looks like the practical part. Without Level 1, people write better prompts and trust the answers more, which is the worst combination available.

Level 1 — What the thing actually is, and why it is confident when wrong

Who it is for. Everyone, including people who will never use it, and especially people who review others' work.

Length. Forty-five minutes, one session. Do not extend it; the demonstration does the work.

Learning outcomes. Each person can say in their own words that the tool produces plausible text rather than looking up facts; that fluency is not evidence of accuracy; and that it will invent a reference or a figure in the same tone it uses when right, because confidence is a property of how it writes, not a signal of how sure it is.

Session outline.

1. *Ten minutes — what it is doing.* It produces the next piece of text

that fits, from patterns in an enormous amount of writing. It is very good at the shape of a document: a memo reads like a memo. The shape being right tells you nothing about the contents.

1. *Fifteen minutes — the exercise below.* This is the session.

2. *Ten minutes — where it is strong.* Rewriting, summarising a document

you supply, first drafts, spotting what is missing from a list. You supply the substance, it supplies the form.

1. *Five minutes — where it is weak.* Anything needing a fact you did not

give it. Arithmetic you have not checked. Anything where nearly right is the same as wrong.

1. *Five minutes — the team's rule,* plus the data rules from your policy.

The live exercise: the confident wrong answer. On screen, unrehearsed — the point is that you do not know what it will say either.

Summarise the disclosure requirements for [TOPIC] under [ACCOUNTING FRAMEWORK] for an entity of type [ENTITY TYPE]. Quote the exact paragraph reference for each requirement and state which pronouncement it comes from.

Take whatever it produces to your source material and read out what is really there. Sometimes it is right. Sometimes it produces a reference in the correct format that does not exist, or attaches a real reference to the wrong requirement. Either way the tone is identical, which is the lesson. Then paste in fifteen anonymised trial balance rows, ask for the movement and percentage change on each, and have someone check every line on a calculator while the room watches. If the arithmetic comes out right, say plainly that one correct run is not a guarantee.

How you tell whether it landed. At the end, ask everyone to write one sentence: *why did it produce a confident wrong answer?* Collect the slips. If most say the data must be out of date, it did not land — that model leads people to trust it on anything historic, so re-run the exercise with a different topic. The real test comes a fortnight later, in the question people ask about a colleague's draft. "Is this right?" means it did not land. "What did you check this against?" means it did.

Level 2 — Writing a usable prompt

Who it is for. Everyone who will actually use the tool for work.

Length. Ninety minutes, in two forty-five minute sessions a week apart, with homework between. The gap matters: people need to try it badly once on their own first.

Learning outcomes. Each person can write a prompt stating role, context, task, constraints and output format; can supply source material instead of expecting the tool to know things; and can tell a bad prompt from a bad tool.

Session outline.

Session A. Five minutes on the failure they have all seen: a one-line request that returns something generic. Fifteen minutes on the five parts — role, context, task, constraints, format. The part finance people skip is context: the entity, the framework, the period, materiality, who reads it. Twenty minutes on the exercise below. Five minutes setting homework: one real task, the prompt written for it, and the output, good or bad.

Session B. Twenty-five minutes workshopping the homework anonymously on screen and fixing prompts as a group — the most valuable half hour in the curriculum. Ten minutes on iteration: the second and third messages do more than the first ("too long", "you have assumed [X], I meant [Y]"). Ten minutes building the shared prompt list, so nobody rewrites the same prompt next month.

The live exercise: two prompts, one task. Run the thin version first:

Write a variance analysis commentary.

It will be plausible and useless. Then:

You are drafting internal management commentary for the finance team of [ENTITY], a [INDUSTRY] business reporting under [ACCOUNTING FRAMEWORK], for the period [PERIOD]. The audience is [AUDIENCE], who are not accountants.

Below is the actual variance data. Using only the figures given:

1. Write one paragraph per variance above [THRESHOLD].
2. For each, state the movement in both absolute and percentage terms.
3. Where the data does not explain a cause, say that the cause is not determinable from the data provided. Do not speculate about causes.
4. List separately any figure you could not reconcile.

Constraints: maximum [NUMBER] words. Plain English. No adjectives about performance. Do not describe any result as strong, weak, disappointing or encouraging.

Data:

[PASTE VARIANCE TABLE]

Compare the two on screen. Point out what instruction 3 prevented: invented explanations for variances are among the most damaging failure modes in finance drafting, because they read exactly like knowledge.

How you tell whether it landed. Two weeks after session B, look at the shared prompt list. Entries from more than half the attendees means it landed. Three entries, all from the owner, means it did not — usually because people are still using it privately for things they do not count as "real work".

Level 3 — Verification, and working with your team's own documents

Who it is for. Anyone whose output is reviewed, and everyone who reviews. In practice, most of the team.

Length. Two sessions of sixty minutes, a week apart.

Learning outcomes. Each person can name what must be checked in a given output and check it rather than redo the work; can ground a request in supplied documents and require quoted source lines; and can say out loud when verification costs more than the drafting saved.

Session outline.

Session A — verification. Ten minutes on the reviewer's problem: output now arrives looking finished, and the old signals — clumsy phrasing, inconsistent formatting — no longer show how much care went in. Fifteen minutes on the four checks. Numbers: recalculate, do not read. Sources: does the quoted line exist, in those words? Omissions: what is missing, which no output will tell you. Claims: any sentence asserting a fact you did not supply is unverified until you verify it. Twenty-five minutes on the exercise below. Ten minutes on how your team evidences that a review happened, using your existing file conventions.

Session B — your own documents. Fifteen minutes on grounding: the difference between asking it about a policy and giving it the policy. Twenty minutes on the quoting technique below. Fifteen minutes on limits: long documents get summarised unevenly, so split them and ask about one section at a time. Ten minutes on which documents may be uploaded at all under your policy — the most-asked question in the curriculum, so have the answer written down.

The live exercise: find the planted errors. Beforehand, take a real piece of assisted output — a summary of an internal policy, or commentary on a set of figures — and plant exactly three errors: one arithmetic, one a quoted phrase that does not appear in the source, one material omission. Do not say how many there are, or that they were planted, until afterwards. Hand out the output and the source documents. Fifteen minutes, in pairs, then compare findings.

Almost every group finds the arithmetic error and misses the omission. That is the lesson, and it is better felt than told: what is missing does not announce itself. Then teach the fix, on the same source material:

Using only the document below, answer this question: [QUESTION].

For every statement you make, quote the exact sentence from the document that supports it, in quotation marks, immediately after the statement. If the document does not address part of my question, list those parts under a heading "Not covered by this document". Do not use any knowledge from outside the document.

Document:
[PASTE DOCUMENT]

Show that the "Not covered" list is the part worth reading first, and check two quotes against the source in front of the room. Paraphrase offered as quotation is a failure to watch for.

How you tell whether it landed. Look at the pilot log. If the "rework needed?" column gets *more* entries after Level 3, the training worked: people are catching what they used to accept. A pilot whose error rate appears to fall straight after verification training has usually stopped looking properly. Say that in advance, so finding errors reads as doing the job well rather than as a bad week.

Level 4 — Designing a repeatable workflow, and knowing when to stop

Who it is for. Managers, process owners, and whoever maintains this after the ninety days.

Length. Three sessions of sixty minutes, plus a follow-up four weeks later to review what was built.

Learning outcomes. Each person can convert an ad hoc use into a documented workflow with a fixed prompt, defined inputs, a checking step and a named sign-off; can judge whether it is worth it once review time is counted; and can retire a workflow that has stopped earning its place.

Session outline.

Session A — from ad hoc to repeatable. What makes a task suitable: frequency, stable inputs, an existing check, a tolerable failure mode. The anatomy of a run sheet — trigger, inputs, prompt, steps, checks, sign-off, storage. Each attendee picks a candidate from their own work.

Session B — build one. Attendees write their run sheet in the room, then swap with a neighbour, who follows it exactly as written and marks every point where they had to guess. A run sheet somebody else cannot follow is not finished.

Session C — knowing when to stop. Four stop signals:

1. Verification costs more than the drafting saved. If review minutes exceed the time saved, stop: the workflow is moving effort, not reducing it.

1. The failure mode is intolerable regardless of frequency. If a wrong output goes straight out with no human in between, the error rate does not need to be high to be unacceptable.

1. The task is the training. Some tasks are how juniors learn the business, and automating the first draft of a reconciliation note may cost you something you will not notice for two years. Make that call deliberately rather than by default.

1. Nobody can explain the output. If the person signing cannot say how the result was arrived at, their name is still on it.

Follow-up, four weeks later. Each attendee reports on their workflow using the log, and the group retires at least one thing. A curriculum that never retires anything is accumulating obligations.

The live exercise: build and run the month-end commentary workflow. Do it as a group, on the real task, in the session. Write the run sheet on screen, then run it end to end with anonymised data while somebody times each step, including review, and writes the numbers where everyone can see them. Then ask the only question that matters: knowing these numbers, would you do this next month? Sometimes the answer is no. Let it be no — a documented decision not to automate something does more for your credibility with the sceptics than three sessions that end in enthusiasm. Use this as the workflow's fixed step:

You are producing the [PERIOD] management commentary for [ENTITY], following the same structure as the attached previous version.

Inputs: the current period figures below, and last period's commentary for structure and tone only – do not carry forward any figure or explanation from it.

Produce:

1. A commentary section for each of [SECTIONS], in the same order as the previous version.
2. Movements stated in both absolute and percentage terms.
3. A separate list headed "Requires input from the business" for anything the figures alone do not explain.
4. A separate list headed "Checks for the reviewer" naming every figure you calculated rather than copied.

Do not state causes that are not evidenced in the data provided. Do not characterise performance.

Current period figures:
[PASTE]

Previous commentary:
[PASTE]

Instruction 4 earns its place: it tells the reviewer where to point a calculator.

How you tell whether it landed. Four weeks later, look for two things: at least one workflow running without the trainer, and at least one workflow retired or never started because the numbers did not support it. If everything survives, nobody applied the stop signals, and you have taught enthusiasm rather than judgement.

When it stalls — and it usually does

Adoption commonly stalls between weeks six and ten. The novelty has worn off, month-end has arrived, and the people who were logging diligently have gone quiet. Plan for it rather than treating it as a verdict.

A stall is information. Read it before reacting; the cure differs entirely by cause.

- **It stalled at month-end.** Nothing is wrong. People fell back to the

known method under deadline pressure, which is a reasonable thing for a professional to do. Restart the following week, and expect a repeat each month until the workflow is faster than the old one under pressure, not just in a quiet week.

- **It stalled because the output was not good enough.** The most useful

outcome available to you. Write it down specifically — not "it isn't accurate" but "it could not reconcile the intercompany balances without more instruction than the task saved".

- **It stalled because verification cost too much.** Legitimate. Narrow

the task so the check is cheaper, or stop. Do not solve it by checking less.

- **It stalled because one person was carrying it.** The most common cause

and the least discussed. Fix the handover, not the enthusiasm.

- **It stalled because nobody was ever asked to.** There was a policy and a

licence and no instruction. If you want a task done a particular way, ask for it the way you ask for anything else.

What not to do is relaunch. A second launch with more energy and a new name, aimed at the people who watched the first one fade, costs you credibility. Fix the specific cause quietly, and let the second attempt be smaller than the first.

What works is dull. Return to one use case, with one group, with the log. Ask the reviewer what they are still correcting. Keep the review date even when there is nothing to report — "nothing happened this month" is the finding that most often needs saying out loud.

And be willing to conclude that some of this does not suit your team, your work or your obligations at the moment. That is a defensible professional judgement, it is cheaper than a slow expensive drift, and it leaves you free to look again in six months with a log full of things you know rather than things you were told.