

A WORKING PACK

AI for Finance Teams

The Working Pack



Tang Shiuan Jenn

Chartered Accountant

No invented statistics. No fabricated case studies. No promises about your income.
Everything in this pack is either an artefact you can use or a claim you can check.

Contents

1. How to Use This Pack
2. Choosing What to Automate First
3. The Prompt Library
4. Verification, Red Lines and Choosing a Tool
5. Your Policy Documents
6. The 90-Day Plan and the Training Curriculum

How to Use This Pack

What this is

This is a working pack, not a book. Nothing in it needs reading cover to cover. It is a set of artefacts you can put in front of a team on Monday: prompts to paste, a policy to adapt, a plan to follow, a scoring sheet to argue over, and a list of the places where the honest answer is to close the laptop.

It was written by a chartered accountant, for people whose name goes on the work.

What it deliberately does not contain

No statistics about AI adoption. No productivity percentages. No case studies about firms that saved a fortune. No promises about your income or your job.

Everything in that list is standard in this market and almost none of it is checkable. The moment a pack like this quotes a figure nobody can verify, every other claim in it becomes a guess too. So there are no numbers here except the ones you generate yourself, in your own team, measuring your own work.

What that costs you: this pack cannot tell you what results to expect. What it buys you: everything it does tell you can be checked.

How to use it

If you have not started at all. Read *Choosing What to Automate First*, run the 45-minute meeting at the end of it, and pick two use cases. Two, not ten. Then take the policy from *Your Policy Documents*, fill it in, and circulate it before anyone else starts experimenting. That order matters — a policy written after the tools have spread is a document about a problem you already have.

If your team is already using AI without rules. Start with *Your Policy Documents* today. Then *Verification*, *Red Lines* and *Choosing a Tool*. The prompts can wait a fortnight; the red lines cannot.

If you are the one who has to teach it. Go to the four-level curriculum in *The 90-Day Adoption Plan and the Training Curriculum*. Each level is written so you can teach it without outside help, with the live exercise already specified.

If you just want better output tomorrow morning. *The Prompt Library*. Find the workflow that matches what is on your desk, adapt the prompt rather than pasting it blind, and read the check line underneath it before you use anything it gives you.

The one rule underneath all of it

Whoever signs the work owns the work.

A tool is never the reason something went out wrong. Every prompt in this pack has a line under it telling you what to check, because a fluent answer and a correct answer look identical until you look. That is the whole discipline. The rest is technique.

A note on what is general and what is yours

Everything here that touches law, tax, professional standards or regulatory duty is general information. Rules differ between countries and professions, and they change. No document in this pack is legal advice, and adopting any of it is not a statement that you comply with anything. Have the policy read by whoever carries risk in your organisation, and check it against your own regulator's guidance and your client engagement terms before you use it.

Choosing What to Automate First

Ask a finance team where they want to start with AI and you will usually hear one of three answers: variance commentary, board paper summarising, or "the audit file". These are the three most interesting problems in the room. They are also, on any honest assessment, three of the worst places to begin.

The reason is not that the work is too hard. It is that in all three cases you cannot cheaply tell whether the output is right. A draft that sounds like a competent colleague wrote it, and is wrong in one place you did not look, costs you more than doing the job yourself. You pay twice: once for the draft, once for the review that has to be as careful as the original work.

What follows is a way of choosing that does not rely on anyone's enthusiasm: five factors, scored one to five, applied to the tasks your team actually does. It takes forty-five minutes with the right people in a room, and it produces two starter use cases you can defend to a sceptic, plus a written reason for everything you left out.

Part One — The five-factor score

Score every candidate task on five factors, one to five. Total out of twenty-five. Two of the factors are inverted, which matters — read the definitions before you score anything.

Factor 1 — Frequency

How often does this task recur?

Score	What it looks like
1	Once a year or less. A one-off.
2	A few times a year. Quarterly.
3	Monthly.
4	Weekly, or several times a month.
5	Daily, or many times a day.

Frequency is where the money is. A task that takes an hour and happens once a year is worth an hour a year. The same task done weekly is worth fifty. Teams routinely pick the annual monster over the weekly irritation, because the monster is memorable and the irritation is invisible.

One exception, and it is a real one. Some low-frequency tasks produce an artefact that then gets used constantly — a close checklist, a set of KPI definitions, a process note. Score the *task* honestly at 1 and note the durable output separately. A 1 with a permanent artefact behind it is a legitimate reason to promote a task up the list. A 1 with a disposable output is not.

Factor 2 — Time

How long does one instance take now, done by the person who normally does it?

Score	What it looks like
1	Under ten minutes.
2	Ten to thirty minutes.
3	Half an hour to two hours.
4	Half a day to a full day.
5	More than a day.

Score the time it takes *now*, not the time you imagine it should take. If nobody knows, time three instances before the meeting. Guessing here corrupts everything downstream, because time and frequency together are the only two factors that describe the prize. The other three describe the cost of collecting it.

Do not score the time you think the tool will save. You do not know that yet, and the point of the exercise is to stop deciding on things you do not know.

Factor 3 — Verifiability

How easily can you tell whether the output is right?

Score	What it looks like
1	You cannot tell without redoing the work from scratch, or correctness is a judgement no one available can make quickly.
2	Checking takes nearly as long as doing. Errors are plausible and quiet.
3	Checkable, but only by someone who knows the subject, reading the whole thing.
4	A knowledgeable person spots an error quickly — minutes, not hours.
5	Correctness is mechanical and obvious. It ties, it foots, it matches the ledger, or the sheet returns the answer you already knew.

This is the factor teams overscore, every time, because verification feels like an afterthought rather than part of the work. It is not an afterthought. It is the entire economics of the thing.

Write the verifiability score last, and write it as a sentence first: "I would know this was wrong because ____." If you cannot finish that sentence in one breath, the score is 1 or 2. If the sentence is "because I would read it carefully", that is a 2 or a 3, not a 5 — careful reading is exactly the expensive thing you were trying to avoid.

Factor 4 — Consequence of error — inverted

If a mistake gets through, what happens? **High score means cheap mistake.** This factor runs backwards on purpose, so that every factor points the same way and the total means something.

Score	What it looks like
1	The error reaches a client, a regulator, a lender, a court or a board, or a professional duty attaches to the work. Expensive and slow to unwind.
2	The error reaches someone outside your team and has to be corrected formally. Embarrassing.
3	The error flows into something else — a return, a report, a system — and gets found later, out of context.
4	The error is caught in normal internal review. Costs minutes and some grumbling.
5	The error is caught instantly by the person who asked, and nobody outside the room ever knows.

Score the consequence of the error *escaping*, not of it being made. Everything gets reviewed in theory. Ask what happens on a Friday afternoon in a busy week.

Factor 5 — Data sensitivity — inverted

What has to go into the tool for the task to work? **High score means nothing confidential.** Also inverted.

Score	What it looks like
1	Impossible without named personal data, client confidential material, unpublished price-sensitive information, or credentials.
2	Real internal records — ledger detail, supplier and customer names, transaction data.
3	Internal material that identifies your organisation and its people, but no client data and nothing price-sensitive.
4	Internal but generic. Procedures, structures, definitions.
5	Nothing confidential at all. Public, generic, or fully anonymised without effort.

Two practical notes. First, "we would anonymise it" only counts if anonymising is quick and reliable. Twenty minutes of careful redaction per instance belongs in the Time score, and the task is less attractive than it looked.

Second, this is general information and not legal advice. What you may put into which tool depends on your contracts with clients, your own confidentiality obligations, the rules of your professional body and the data protection law where you and your clients sit. Those differ by jurisdiction and they change. Get your own position checked rather than relying on a score in a table.

Reading the total

Total	Reading
20–25	Start here. Rare. If you have one, do it this month.
16–19	The working middle. Most of a sensible programme lives in this band.
11–15	Conditional. Worth doing only if you can redesign the task to lift the weak factor.
10 or below	Leave it. Revisit in a year if something changes.

Then apply three vetoes, which override the total completely:

1. **Verifiability of 1 or 2 — vetoed.** Whatever the total.
2. **Consequence of 1 — vetoed.** Whatever the total.

3. Data sensitivity of 1 — vetoed. Whatever the total.

A data sensitivity of 2 is not a veto but a constraint: that task is only available to you inside a tool where the material genuinely stays under your control, on terms someone has actually read. If you do not have such a tool, every 2 behaves like a 1 and drops out.

The vetoes exist because a total is an average, and averages hide the one thing that will hurt you. A task scoring 4, 5, 2, 1, 1 — frequent, slow, unverifiable, catastrophic if wrong, full of client data — totals 13 and sits comfortably mid-range. Thirteen out of twenty-five is not a comfortable task. It is a disaster with good attendance.

Why the "time" trap catches everyone

Here is the failure mode this whole method exists to prevent.

Someone proposes a task that scores 4 or 5 on time. It is the task that eats their Thursday. It is genuinely painful and everybody agrees it is painful, which makes it feel like an obvious candidate. But it scores 2 on verifiability, because the only way to know the output is right is to do the underlying work.

The arithmetic is brutal. The task takes four hours; the tool drafts it in ten minutes. You have not saved three hours and fifty minutes. You have saved nothing until the draft is checked, and checking an unverifiable four-hour output takes close to four hours, because the check *is* the work. You have added ten minutes and a false sense of progress.

Worse, the draft is persuasive. It is fluent, well-structured and confident. A checker reading a fluent draft looks for errors; a person doing the work from scratch encounters the facts. Those are different mental activities and they catch different things. A fluent wrong draft is harder to correct than a blank page, because it tells you where to look and the place it tells you to look is not where the error is.

So: time tells you the size of the prize; verifiability tells you whether you can collect it. A high time score with a low verifiability score is not a good use case with a snag. It is a bad use case that looks good, which is worse than one that looks bad, because someone will champion it.

The corollary is cheering. Tasks that score modestly on time — twenty minutes, done daily, checked in ten seconds — are the ones that quietly pay. They are unglamorous, nobody will present them at a conference, and they will save your team more hours this year than the Thursday monster ever will.

Part Two — Thirty tasks, scored

Below are thirty tasks common to finance functions and professional firms, scored against the five factors. The scores are mine and they are arguable — deliberately so. The point of a published scale is that you can look at a row, disagree with one number, and see exactly which definition we read differently. Change the score, change the total, carry on. What you must not do is change a score because you want a different answer.

Assume a general-purpose assistant of the kind anyone can reach from a browser. Where a row is constrained by data sensitivity, the constraint means an enterprise-tier arrangement your organisation controls, on terms someone has read.

Key: **F** frequency, **T** time, **V** verifiability, **C** consequence (inverted), **D** data sensitivity (inverted).

#	Task	F	T	V	C	D	Total	Verdict
1	Explaining an inherited spreadsheet formula	4	2	5	5	5	21	Best possible starter — paste the formula, not the data; the sheet marks the answer.
2	Writing a formula or query from a plain description	4	2	5	5	5	21	Same family. Test on a sample where you know the result before trusting it.
3	Drafting an internal training exercise for the team	2	3	4	5	5	19	Low risk, genuinely useful. Check the model answers yourself.
4	Reconciliation prep — formatting and first-pass matching	5	3	5	4	2	19	High frequency and the ledger marks the work. Constrained only by where the data may go.
5	Meeting notes into circulated minutes	5	2	4	4	3	18	Strong starter. You were in the room, so you can mark the output.
6	Routine client emails — updates, chasers, acknowledgements	5	1	5	4	3	18	Small saving each time, large in aggregate. Never send unread.
7	Process documentation for a routine task	2	4	4	4	4	18	High value; the person who does the task is the reviewer. Backlog makes it better than it scores.
8	Month-end close checklist — build or refresh	1	3	5	4	5	18	One afternoon, permanent artefact. An excellent pilot.
9	Internal finance policy — first draft	1	4	4	4	5	18	Good use. The draft is a skeleton; the control detail must be yours.
10	Job description for a finance vacancy	2	3	4	4	5	18	Fine filler task. Do not let it invent qualifications, or say anything about pay you have not decided.
11	Audit request list / information list for a client	2	2	4	4	5	17	Quick win before every audit. Your auditor will amend it anyway.
12	Interview questions and scoring sheet for a finance hire	1	3	4	4	5	17	Solid and low-stakes. Keep the scoring criteria yours.
13	User acceptance test cases for a system or control	1	4	4	4	4	17	Good. The tool supplies breadth; you supply the nasty cases.
14	KPI definition dictionary	1	3	5	4	4	17	Unglamorous, durable, week-one material. Your team recognises a wrong definition instantly.
15	Plain-language rewrite of a technical memo	3	2	4	4	3	16	Common and useful. The check is fast because you own the source.
16	Forecast scenario — building the structure	2	4	4	3	3	16	Use it for scaffolding, never for assumptions. Test the mechanics with known inputs.
17	Cleaning and standardising a messy data extract	4	3	3	3	2	15	Only with row-count and total controls in place. Silent row loss is the risk, not visible errors.
18	Invoice coding suggestions against the chart of accounts	5	1	4	3	2	15	Volume is the attraction. Treat every suggestion as a first pass a human approves.
19	Fixed asset register tidy-up — descriptions and categories	1	4	3	3	3	14	Middle of the pack. Check categories yourself; they drive depreciation and tax.

#	Task	F	T	V	C	D	Total	Verdict
20	Long email chain into a decision log	4	2	3	3	2	14	Useful for your own recall. Weak as a record anyone else relies on.
21	Variance commentary — first draft from figures you supply	3	3	3	3	2	14	The classic trap. Quick to draft, slow to trust — the tool cannot know <i>why</i> the number moved.
22	Contract summarising into a key-terms schedule	3	4	3	2	1	13	Vetoed on data. The check is re-reading the contract, which was the work.
23	Board paper summarising	2	3	3	2	1	11	Vetoed on data. The omission you cannot see is the entire risk.
24	Audit working paper — documenting work performed	4	3	2	1	1	11	Vetoed three ways. The record must describe what you actually did.
25	Tax computation review	2	4	2	1	1	10	Vetoed. Use it for the covering memo if anything, never the computation.
26	Payroll queries involving named individuals	3	2	3	1	1	10	Vetoed. Named pay data, and the error lands on a person.
27	Credit assessment of a named customer	3	3	2	1	1	10	Vetoed. A decision about a named party on data you cannot check.
28	Researching the treatment of a novel transaction	2	4	1	1	2	10	Vetoed. Where confident-sounding wrong answers do the most damage.
29	A difficult client email — bad news, fee dispute, complaint	2	2	3	1	1	9	Vetoed. In a live dispute your words become evidence. Write it yourself.
30	Drafting a response to a regulator or professional body	1	3	2	1	1	8	Vetoed. The lowest score here, and rightly.

What the table is telling you

Read the shape of it rather than the individual rows.

The top is boring. The two highest scores are explaining and writing spreadsheet formulas — tasks so small that nobody would propose them in a strategy meeting. They win because the sheet tells you immediately whether the answer works, and because you paste a formula rather than data. Frequency and verifiability, not ambition.

The bottom is everything anyone was excited about. Board papers, tax computations, working paper documentation, regulator correspondence. These are the tasks that hurt most, which is precisely why they involve the most confidential material and the heaviest consequences. Pain and suitability are not correlated. If anything they run against each other.

The middle is where the argument is. Rows 17 to 21 all score 14 or 15 — data cleaning, invoice coding, variance commentary. These are the rows worth an hour of your team's disagreement, because the score depends on facts about your organisation that I do not have. If your ledger data can sit in a tool you control, invoice coding moves from 15 to 18 and becomes a serious candidate. If it cannot, it drops out entirely. Same task, different organisation, different answer.

Two rows deserve a word each.

Variance commentary (row 21) is the single most requested use case in finance, and it scores 14. The figures are verifiable; the explanation is not. Say distribution costs are up on last month. A tool can see the movement. It cannot know that you changed courier in week three, and on its own it will not say "I don't know" — it will produce a plausible reason. Plausible reasons in management commentary are worse than blank space, because someone will act on one. If you use a tool here at all, use it to draft the *structure* and to check that every movement above your threshold has been addressed. Supply the causes yourself.

Audit working paper documentation (row 24) scores 11 and is vetoed three times over. It deserves the emphasis. A generated narrative describing work performed is indistinguishable, on the page, from a narrative describing work performed — whether or not it was. That is not a technology problem you can solve with a better prompt. It is the reason the record exists. Where professional duties attach to your file, the documentation is the evidence of the work, and it has to be written by the person who did it. Check what your own professional body and your firm's methodology require; this varies and it changes.

Disagreeing properly

Someone in your team will say a score is wrong. Good. Handle it like this:

1. Name the factor. Not "this should be higher" — "verifiability should be 4, not 3".

1. Quote the definition. Read the row from the scale out loud.
2. Say what you would do differently. "I would know it was wrong because the control total does not match" is an argument for 5. "I would read it properly" is not.

1. Change the number and let the total fall where it falls.

The one rule: nobody argues about a total. Totals are arithmetic. If you find yourself defending a total, you have decided the answer and are reverse-engineering the scores, and the method has stopped working.

Part Three — Running this in your own team in forty-five minutes

This is a single meeting. It does not need a workshop, a facilitator or a wall of sticky notes, and it produces a decision rather than a discussion.

Who is in the room

Four to eight people. You need three roles covered, and one person can cover two:

- **People who do the work.** They score frequency and time. Nobody else is entitled to.
- **People who review the work.** They score verifiability and consequence. The reviewer's verifiability score wins any tie, because they are the one who will pay for a bad one.
- **Whoever owns confidentiality and data protection.** They score data sensitivity, and they hold a veto that needs no explanation.

Do not invite anyone whose only contribution is enthusiasm.

Before the meeting

Circulate the five factor scales — just the five tables from Part One, one page. Ask everyone to bring **three tasks they personally do**, written on separate lines, described concretely. "Reporting" is not a task. "Writing the commentary on the departmental cost report" is.

If nobody knows how long a candidate task takes, ask the person who does it to time three instances that week. Ten minutes of measurement beats an hour of estimating, and the number is defensible afterwards.

The agenda

Time	Item	Output
0:00–0:05	Purpose and rules. Read out the three vetoes.	Everyone knows a high total can still lose.
0:05–0:12	Silent listing. Each person writes their three tasks. No discussion.	12–24 raw candidates.
0:12–0:20	Read out and merge. Combine duplicates, discard anything nobody actually does.	A clean list of 15–25.
0:20–0:35	Score. Thirty seconds a task, by role, in order F-T-V-C-D.	Every task has five numbers.
0:35–0:42	Apply vetoes. Strike the failures. Sort by total.	A ranked, filtered list.
0:42–0:45	Pick two. Name an owner each. Set the review date.	A decision.

Fifteen minutes for scoring sounds tight and is not. Thirty seconds a task is enough when the scales are on the table in front of people, and the clock is doing useful work: it stops the room relitigating a task everyone already understands.

Rules for the room

- **Silent listing first.** If you open with discussion, the most senior or most talkative person sets the agenda and you will end up scoring their three tasks.
- **Score by role, not by consensus.** The doer's time score stands. The reviewer's verifiability score stands. Consensus scoring drifts towards whatever the room already wanted.
- **No scoring your own task alone.** If you do the task and review it, someone else scores verifiability.
- ****Write the "I would know it was wrong because _____" sentence out loud**** for anything scoring 4 or 5 on verifiability. Half of them collapse to a 3 the moment they are spoken.
- **Vetoes are silent.** No debate. The veto is not a judgement about the

task or the person who proposed it.

What you leave with

Two use cases. Not five, and not one.

One gives you no comparison: if it goes badly you cannot tell whether the method or the task was at fault. Five gives you five half-started experiments and no evidence, because no small team learns five new habits at once.

For each of the two, write down before you leave the room:

1. The owner. One named person, not a team.
2. The baseline: how long the task takes now, measured over three

instances, and how often it happens.

1. What "working" would look like in four weeks, in time or rework — not in enthusiasm.

1. The verification step, in one sentence, written before anyone starts.
2. The date you will look at it again. Four weeks is about right.

After four weeks

Come back for fifteen minutes. Three possible answers:

- **It worked.** Write down what verification actually cost, because that

is the number you will need when someone proposes the next task. Then score the next two candidates from your ranked list.

- **It half-worked.** Usually this means the verification step was invented afterwards rather than designed first. Fix that and run another four weeks. Do not widen.

- **It did not work.** Say so plainly, record why against the five factors, and move to the next candidate. A use case that fails against a written scale is information. A use case that fails against nothing is just a bad mood.

If nothing scores well

This happens, particularly in small practices where nearly everything touches client confidential material. If your list has no task above 15 and every high scorer is vetoed on data, that is a genuine finding, and it is worth more than a forced pilot.

It usually points at one of two things. Either your tooling is the constraint — the work would qualify if the material stayed under your control, which is a procurement question rather than a use case question. Or your suitable tasks are all small and you have been looking for something big. In the second case, go back to the top of the table: formula explanation, checklists, definitions, internal drafting. Small, frequent, instantly checkable, nothing confidential. That is not a consolation prize. On the arithmetic, it is the answer.

Re-run it

Score again every quarter, or whenever one of three things changes: the tools available to you, your policy on what data may go where, or the shape of the work itself. Keep the old scoresheets. Being able to show why a task was rejected in March, and what changed by September, is worth as much internally as the decision itself — particularly to the person who has to answer for it.

The Prompt Library

A prompt is a briefing, not a spell. None of the forty prompts below will do much for you if you paste one, take whatever comes back and put your name on it. They are written on the assumption that you already know something the model does not: what your entity does, which numbers matter this period, what your reviewer will ask, and what the consequences are if the answer is wrong. The prompt's job is to get that knowledge out of your head and in front of the model in an orderly way. Your job is everything after.

Three things are worth changing in almost every prompt here. First, the role line — replace the generic description with the specific one that fits your situation, because "you are assisting a finance team in a group with several overseas subsidiaries and a manual consolidation" produces different work from "you are assisting a finance team". Second, the constraints — add your house rules, your materiality thresholds, your naming conventions, anything a new joiner would get wrong. Third, the output format — make it match the document the output actually has to feed, so you are pasting into a template rather than reformatting by hand.

Two things are worth leaving alone. The instruction to mark uncertainty and refuse to guess: take it out and you get confident-sounding filler, which is the single most expensive failure mode in this work. And the instruction not to invent figures. Before any of that, though, settle a separate question: whether you are allowed to put this particular information into this particular tool. Client data, staff data, unpublished results and anything under a confidentiality obligation are decisions for whoever owns that policy in your organisation, not for you at your desk on a Tuesday. If nobody has decided, assume the answer is no until they do.

Change this	Keep this
The role line, so it matches your entity and sector	The uncertainty rule and the refusal to guess
The constraints, so they carry your house rules	The ban on inventing figures
The output format, so it fits your template	The instruction to list what it could not determine
The placeholders in square brackets	The request to show its reasoning where judgement was used

1. Month-end close and reconciliations

P1.1 — Reconciliation break triage

Use when: you have a list of unexplained differences and limited time before close.

Paste this:

You are assisting a finance team during month-end close. I will give you a list of reconciling items from [ACCOUNT NAME] for [PERIOD].

Task: sort them into groups by the most likely underlying cause, and for each group tell me what evidence would confirm or rule out that cause.

Constraints: do not estimate, extrapolate or invent any amount, date or counterparty. Use only what is in the list I paste. Where an item could belong to more than one group, say so rather than picking one.

Format: a table with columns Item reference, Amount as given, Likely cause, Evidence that would confirm it, Who to ask. Below the table, a short list headed "Cannot classify from this information" with the reason for each.

If you are unsure about an item, put it in the cannot-classify list. Do not guess.

Here is the list:
[PASTE LIST]

Check before you use the output: the groupings are guesses dressed as categories — confirm each cause against the source document before you act on it, and watch for items it quietly assigned to a cause because they resembled the item above.

P1.2 — Break explanation, written up

Use when: you have worked out why a difference exists and now have to write it down for the file.

Paste this:

You are helping a qualified accountant write a file note. I will describe in rough terms why a reconciling difference on [ACCOUNT NAME] for [PERIOD] arose, and what I did about it.

Task: turn my notes into a clear file note that another accountant could follow in a year's time without asking me questions.

Constraints: use only the facts in my notes. Do not add causes, amounts, dates, control references or conclusions I have not given you. Do not soften or strengthen my conclusion. British English, plain sentences, no adjectives of reassurance.

Format: four short paragraphs – what the difference was, why it arose, what was done, what remains open. Under 300 words.

If my notes leave a gap that a reader would trip over, do not fill it. List the gaps at the end under "Questions I could not answer from your notes".

My notes:
[PASTE NOTES]

Check before you use the output: the "what remains open" paragraph — models tend to tidy loose ends into resolutions. If you left something unresolved, make sure the note still says so.

P1.3 — Close checklist builder

Use when: your close runs on memory and one person's spreadsheet.

Paste this:

You are helping a finance team document its month-end close.

Context: I will describe the tasks we perform, roughly in order, in my own words. Our team is [NUMBER] people. The close runs from working day [X] to working day [Y].

Task: turn this into a close checklist that could be handed to someone who has not run our close before.

Constraints: include only tasks I have described. Do not add standard tasks you assume we perform – instead list them separately as suggestions I can accept or reject. Do not assign deadlines I have not given you.

Format: a table with columns Working day, Task, Owner, Depends on, Evidence retained. Then a separate list headed "Tasks you did not mention that teams often include – confirm whether these apply to you".

Where my description is ambiguous about sequence or ownership, mark the cell "UNCLEAR" rather than choosing.

My description:
[PASTE DESCRIPTION]

Check before you use the output: the dependency column. A checklist that puts a task before the thing it depends on is worse than no checklist, and this is the field most often filled in by pattern rather than logic.

P1.4 — Intercompany mismatch investigation plan

Use when: two entities disagree and you need a plan before the calls start.

Paste this:

You are assisting a group finance team. Two entities, [ENTITY A] and [ENTITY B], disagree on the intercompany balance for [PERIOD]. I will describe what each side reports and what I already know.

Task: produce an investigation plan – the order in which to test possible causes, cheapest and most likely first.

Constraints: work only from what I tell you. Do not assume our systems, our currencies, our cut-off conventions or our transfer pricing arrangements. Do not calculate a difference I have not given you.

Format: numbered steps. Each step states what to check, what result would confirm the cause, and what to do next if it does not.

At the end, list under "Assumptions I had to make" anything you took as given that I did not actually state. If you cannot form a plan without information I have not provided, say what you need instead.

Situation:

[PASTE SITUATION]

Check before you use the output: the assumptions list, and then the steps that rely on those assumptions. Currency and cut-off conventions are the two it will invent most readily.

P1.5 — Accrual and prepayment review questions

Use when: you are reviewing someone else's schedule and want a second pair of eyes on what to ask.

Paste this:

You are acting as a reviewer of a schedule of accruals and prepayments for [ENTITY], [PERIOD]. I will paste the schedule as text.

Task: produce the questions a careful reviewer would ask about this schedule. Focus on what would be wrong if the preparer had been rushing.

Constraints: do not recalculate anything and do not tell me any figure is wrong – you cannot see the supporting evidence. Ask questions instead. Do not reference any accounting standard by number or name; if a point turns on a recognition question, describe it in plain words and mark it "judgement – confirm with your own reference sources".

Format: questions grouped under Completeness, Cut-off, Basis of estimate, Reversal and follow-through. Maximum 20 questions, ordered by how much money is likely to turn on them.

Where the schedule does not give you enough to form a question, say so rather than asking a generic one.

Schedule:
[PASTE SCHEDULE]

Check before you use the output: whether the questions are actually specific to your schedule. Generic review questions are a sign it did not have enough context, and they will waste your preparer's afternoon.

2. Variance analysis and management commentary

P2.1 — First-pass variance question list

Use when: the variance report has landed and you have not yet spoken to the business.

Paste this:

You are assisting a financial controller preparing to investigate variances for [ENTITY], [PERIOD]. I will paste a variance table showing actual, budget or prior period, and the difference.

Task: for each line I flag, list the specific questions I should put to the budget holder – the ones that would distinguish between a timing difference, a volume difference, a price or rate difference, a misposting, and a genuine change in the business.

Constraints: do not explain the variances yourself. You do not know our business. Do not calculate percentages or restate my figures. Do not speculate about causes in the wording of the questions.

Format: for each flagged line, the account name and then between three and six questions as bullets.

If a line is too small or too ambiguous to be worth a conversation, say so and move on.

Variance table:
[PASTE TABLE]

Check before you use the output: that it has not smuggled a presumed cause into a question ("why did headcount costs rise due to the new hires?"). Leading questions get you the answer you suggested, not the one that is true.

P2.2 — Commentary drafted from your explanations

Use when: you know the reasons and need them written up for the pack.

Paste this:

You are drafting management commentary for [ENTITY], [PERIOD], for an audience of [AUDIENCE].

I will give you, for each variance: the account, the figure, and the explanation I have obtained from the business. That explanation is the only permitted source of cause.

Task: write the commentary.

Constraints: every stated cause must come from my notes. Do not add context, market conditions, industry trends or forward-looking statements. Do not add any figure that is not in my notes. Do not use the words "significant", "material", "strong" or "challenging". Do not tell the reader how to feel about the result. British English.

Format: one short paragraph per variance, ordered by size of difference, each beginning with the account name in bold. Maximum [WORD LIMIT] words in total.

If an explanation I gave you does not actually account for the movement described, flag it at the end under "Explanations that do not fully account for the movement" rather than writing around it.

My notes:

[PASTE NOTES]

Check before you use the output: the flagged list at the end, then every causal word in the body — "because", "driven by", "reflecting". Each one has to trace back to something a named person actually told you.

P2.3 — The walk from last period to this

Use when: you need a bridge narrative that a reader can follow without the chart.

Paste this:

You are helping a finance team explain a movement in [MEASURE] from [PRIOR PERIOD] to [CURRENT PERIOD] for [ENTITY].

I will give you the opening figure, the closing figure, and the components of the movement with their amounts and my explanation of each.

Task: write the bridge as prose, in the order that makes the story clearest, not necessarily the order I listed.

Constraints: the components must sum to the movement as I have given them – if they do not, stop and tell me the difference rather than adjusting anything or inserting a balancing item. Add no components of your own. No causes beyond mine.

Format: a short lead sentence stating the movement, then one sentence per component, then a closing sentence on what it means for the next period only if I have told you what that is.

State plainly if you cannot reconcile my components to my movement.

Inputs:

[PASTE INPUTS]

Check before you use the output: whether it reported a failure to reconcile. If it produced a smooth narrative and your components do not add up, it has hidden your error inside good prose.

P2.4 — De-waffling a draft

Use when: the commentary is written but reads like a press release.

Paste this:

You are an unsentimental editor of finance writing. I will paste a draft of management commentary.

Task: rewrite it so that every sentence carries information. Cut hedging, cut throat-clearing, cut adjectives that assert rather than describe.

Constraints: do not change any figure, any cause or any conclusion. Do not add facts. If cutting a sentence would remove a fact, keep the fact and cut the decoration around it. British English. Preserve any wording I have marked [DO NOT CHANGE].

Format: the rewritten text first. Then a short list of sentences you could not tighten because they contained a claim you could not verify from the draft itself – these are the ones I need to check.

Do not invent a stronger version of a weak claim. If a sentence asserts something without support, leave it and flag it.

Draft:

[PASTE DRAFT]

Check before you use the output: the flagged list, and a line-by-line comparison of figures. Editors of all kinds have a habit of rounding numbers to make a sentence sit better.

P2.5 — Attack my explanation

Use when: you are about to present a reason and want it tested first.

Paste this:

You are playing the role of a sceptical audit committee member with a finance background. I will give you my explanation for a movement in [MEASURE] at [ENTITY] for [PERIOD].

Task: attack it. Find the questions that would expose it if it were wrong, incomplete, or an artefact of how the numbers were prepared.

Constraints: do not propose alternative explanations as though they were facts, and do not invent figures to support a challenge. Frame every challenge as a question I can go and answer. Assume I am competent and sceptical of you too.

Format: challenges in three groups – Does the cause explain the size of the movement, Does it explain the timing, and Is there a simpler explanation. Maximum ten challenges, hardest first.

If my explanation is internally consistent and you cannot fault it on the information given, say so plainly instead of manufacturing objections.

My explanation:
[PASTE EXPLANATION]

Check before you use the output: whether the challenges are answerable. A challenge you cannot resolve with evidence you can obtain is not useful before a meeting; park it and deal with the ones you can.

3. Budgeting and forecasting

P3.1 — Assumption map for a budget line

Use when: a budget line has become a single number nobody can defend.

Paste this:

You are helping a finance team make its budget assumptions explicit for [BUDGET LINE] at [ENTITY] for [PERIOD].

Task: break the line down into the drivers that would have to be estimated to build it up from scratch, and for each driver state what source inside a business would normally hold the answer.

Constraints: give no values, ranges, percentages or benchmarks. Not one number. You do not know our business, our sector or our size. Where a driver depends on a decision rather than a fact, mark it "decision required" and say whose decision it is.

Format: a table with columns Driver, Unit, Type (fact / estimate / decision), Where the answer lives, Question to ask.

If the line could be built up in more than one sensible way, show the two most common structures and say what would make one preferable.

Check before you use the output: that the drivers actually multiply or add up to your line. A tidy driver list that does not reconstruct the number is decoration.

P3.2 — Budget holder briefing

Use when: you are sending the budget pack to non-finance managers.

Paste this:

You are helping a finance team brief non-financial managers who must submit budget figures for [PERIOD].

Context: their cost lines are [LIST LINES]. The submission deadline is [DATE]. Common problems last cycle were [LIST PROBLEMS].

Task: write the briefing note that goes out with the template.

Constraints: no finance jargon without a plain-English gloss on first use. No threats and no cheerleading. Do not state any figure, limit or target – I will insert those. Do not promise what finance will do with the numbers beyond what I have told you.

Format: a short opening on what we need and when, a numbered list of what to submit, a numbered list of what not to do, and a closing line on who to contact. Under 500 words.

Leave a clearly marked [] placeholder wherever a figure, limit or name is required rather than supplying one.

Check before you use the output: every placeholder is still empty. If it has helpfully filled one in with a plausible-looking number, that number will end up in a manager's spreadsheet.

P3.3 — Forecast assumption challenge

Use when: the forecast is built and you want the weak joints found.

Paste this:

You are reviewing the assumption set behind a forecast for [ENTITY], covering [PERIOD]. I will paste the assumptions as a list, with values.

Task: identify which assumptions the outcome is most sensitive to, which are internally inconsistent with each other, and which are stated so vaguely that they cannot be tested later.

Constraints: do not tell me an assumption is too high or too low – you have no basis for that and I will not accept it. Restrict yourself to sensitivity, internal consistency and testability. Do not compute revised outcomes.

Format: three sections – Most sensitive, Mutually inconsistent, Not testable as written. In each, name the assumption and say precisely what the problem is.

Where two assumptions appear inconsistent but might be reconciled by something I have not told you, phrase it as a question.

Assumptions:
[PASTE ASSUMPTIONS]

Check before you use the output: the sensitivity claims. Without your model it is reasoning from structure alone, so treat "most sensitive" as a list to test, not a finding.

P3.4 — Scenario definitions, without numbers

Use when: you have been asked for scenarios and want them defined before anyone models them.

Paste this:

You are helping a finance team define scenarios for [ENTITY] over [PERIOD]. I will describe the business and the uncertainties that actually worry management.

Task: define [NUMBER] scenarios. Each is a coherent story about the world, not a percentage applied to a base case.

Constraints: no numbers, no probabilities, no likelihoods. Each scenario must state which of my named uncertainties resolve which way, and the consequences must be internally consistent – if demand falls in a scenario, do not also assume prices hold unless you say why. Use only the uncertainties I give you.

Format: for each scenario – a name, a one-paragraph description, a bulleted list of which assumptions move and in which direction, and one line on the earliest observable sign that this scenario is happening.

If two of my uncertainties cannot both be varied independently, say so.

The business and the uncertainties:
[PASTE DESCRIPTION]

Check before you use the output: the early-warning signs. They are the part of a scenario set that gets used, and the part most likely to be written as a platitude.

P3.5 — Budget submission review

Use when: submissions are in and you must review many quickly.

Paste this:

You are performing a first-pass review of a budget submission from [DEPARTMENT] for [PERIOD]. I will paste the submission and, separately, our submission rules.

Task: check the submission against the rules only. List every point of non-compliance, and every place where a required explanation is missing or is too vague to be an explanation.

Constraints: do not assess whether the figures are reasonable. That is my job and you lack the context. Do not compare to prior year unless I have given you prior year. Quote the exact wording from the submission when you flag it.

Format: a table with columns Rule, What the submission says, Problem, What to request. Then a one-line overall statement – ready for review, or return to sender.

If a rule is ambiguous as written, flag the rule itself rather than guessing which reading applies.

Rules:

[PASTE RULES]

Submission:

[PASTE SUBMISSION]

Check before you use the output: the quoted wording. If a quotation does not appear verbatim in the submission, treat the whole review as unreliable and run it again.

4. Audit and assurance support

P4.1 — Making sense of a request list

Use when: the audit request list has arrived and you need to allocate work.

Paste this:

You are helping a finance team plan its response to an auditor's information request list for [ENTITY], [PERIOD]. I will paste the list.

Task: for each item, state in plain words what the auditor is likely to be trying to establish, what document or extract would satisfy it, and which of my team roles would normally hold it.

Constraints: this is general information about how such requests are usually framed, not advice on what any standard or regulator requires. Say that plainly at the top of your output. Do not cite any standard, regulation or professional pronouncement. Do not tell me an item is unnecessary or can be refused.

Format: a table with columns Request, Likely purpose, Document to provide, Owner, Effort (low / medium / high).

Where you cannot tell what is being asked for, write "ask the audit team to clarify" rather than guessing.

Request list:
[PASTE LIST]

Check before you use the output: the "likely purpose" column. It is inference, and if you send documents based on a wrong inference you will do the work twice.

P4.2 — Drafting a response to an audit query

Use when: you have the facts and need the reply written cleanly.

Paste this:

You are helping a qualified accountant respond to an audit query on [TOPIC] for [ENTITY], [PERIOD]. I will give you the query and the facts and documents we hold.

Task: draft the response.

Constraints: use only the facts I supply. Do not assert that a treatment is correct, compliant or in accordance with anything – state what was done and why we did it, and let the auditor form their view. Do not name any standard, regulation or authority. Do not commit us to anything I have not said we will do. Neutral tone, no defensiveness.

Format: a short statement of our position, then the facts in order, then a list of the documents attached, then any point we accept remains a matter of judgement.

If the facts I have given you do not fully answer the query, say what is missing at the end instead of writing around the gap.

Query and facts:
[PASTE]

Check before you use the output: any sentence that asserts compliance. Delete it and replace it with what you actually did. Your position is what happened, not a claim about the rules.

P4.3 — Control description write-up

Use when: a control exists in practice but has never been written down.

Paste this:

You are documenting an internal control for [ENTITY]. I will describe what actually happens, in my own words, including the parts that are informal.

Task: write the control description.

Constraints: describe only what I have told you happens, including the weaknesses. Do not write the control as it ought to be. Do not add frequency, thresholds, approvers or evidence that I have not stated. Do not reference any control framework by name. Do not use the word "ensures" – say what the control does and what it would catch.

Format: What the control is, Who performs it, When and how often, What evidence it leaves, What it would catch, What it would not catch.

The last section is the important one. If my description does not let you complete it honestly, say so and ask me the questions that would.

What happens:
[PASTE DESCRIPTION]

Check before you use the output: the "what it would not catch" section. If it is short or reassuring, it is wrong — that section is the reason for writing the description at all.

P4.4 — Walkthrough interview questions

Use when: you are about to sit with a process owner and trace a transaction.

Paste this:

You are preparing an interview plan for a walkthrough of the [PROCESS] process at [ENTITY]. I will describe what we believe the process is.

Task: write the questions to ask the person who actually performs it, designed to reveal where the documented process and the real one differ.

Constraints: open questions only – nothing answerable with yes. Do not ask questions that presume the process is being followed. Do not assess anything; you are writing questions, not conclusions. No framework or standard references.

Format: questions in the order the transaction flows, each with a one-line note on what a concerning answer would sound like.

Add a final group headed "Questions to ask if they seem to be describing the manual rather than their day", with the questions that usually surface workarounds.

Believed process:
[PASTE DESCRIPTION]

Check before you use the output: that the questions do not read as an interrogation. You want a description of reality, and people stop giving you one the moment they feel tested.

P4.5 — Evidence gaps in a testing plan

Use when: you have drafted a testing approach and want the holes found.

Paste this:

You are reviewing a draft testing approach for [AREA] at [ENTITY], [PERIOD]. I will paste the approach, including what we intend to test, how, and over what population.

Task: identify where the evidence obtained would not actually support the conclusion drawn – gaps between what the test proves and what the approach claims it proves.

Constraints: this is a logical review, not a professional or regulatory one. State that at the top. Do not tell me whether the approach meets any standard, and do not name one. Do not propose sample sizes or statistical parameters. Do not assume facts about our population that I have not given you.

Format: for each gap – What the test does, What conclusion is drawn, Why the first does not support the second, What additional evidence would close it.

If the approach is sound as far as you can tell from the information given, say so and say what information you would need to be more confident.

Approach:
[PASTE APPROACH]

Check before you use the output: whether it stayed inside logic. The moment it starts telling you what is required, it has left the territory it can be trusted in.

5. Tax and compliance preparation

Everything in this group is preparation. None of it is advice, and none of the outputs should reach a filing without a qualified person in your jurisdiction reviewing it. Tax rules differ by country, differ within countries, and change — often with effect from a date that has already passed by the time you read about it. Use these prompts to get organised, not to reach conclusions.

P5.1 — Information gathering checklist

Use when: a filing or return is coming and the documents are scattered.

Paste this:

You are helping a finance team organise itself before a [FILING TYPE] for [ENTITY] covering [PERIOD]. I will list the activities the entity carried out during the period.

Task: produce a checklist of the documents and records the team should gather and have ready for whoever prepares the filing.

Constraints: this is a document-gathering aid and general information only, not tax advice. Rules differ by jurisdiction and change; I will have this reviewed by a qualified adviser in my jurisdiction. Do not state what is deductible, taxable, claimable or reportable. Do not reference any tax law, section, regulator or filing form. Do not assume my jurisdiction.

Format: a table with columns Activity, Records to gather, Where they usually live, Owner, Status.

Where an activity's treatment plainly depends on local rules, write "jurisdiction-dependent – for your adviser" in place of any conclusion.

Activities:
[PASTE LIST]

Check before you use the output: that nothing in it has drifted from "gather this" to "this is treated as". Strike any line that states a treatment.

P5.2 — Categorisation queries

Use when: you are cleaning a transaction listing before it goes to an adviser.

Paste this:

You are helping prepare a transaction listing for [ENTITY], [PERIOD], before it is sent to our tax adviser. I will paste the listing with descriptions and amounts.

Task: identify the transactions whose description is too vague, too generic or too inconsistent for someone else to categorise, and write the question we need to put to the person who incurred them.

Constraints: do not categorise anything yourself and do not state any tax treatment. This is preparation, not advice. Do not change or recalculate any amount. Do not infer a purpose from a supplier name.

Format: a table with columns Reference, Description as given, Why it is unclear, Question to ask, Who to ask.

Where a description is adequate, leave it out. A short list is a good outcome.

Listing:
[PASTE LISTING]

Check before you use the output: whether it has inferred purpose from vendor names. "Software supplier, therefore software cost" is exactly the kind of assumption that later needs unwinding.

P5.3 — Understanding a letter you have received

Use when: correspondence has arrived and you need to know what is being asked before you panic.

Paste this:

You are helping a finance professional understand a piece of correspondence. I will paste the text of a letter received by [ENTITY].

Task: tell me, in plain words, what the letter appears to ask for, what it appears to require by when, and what it does not say.

Constraints: this is a reading aid and general information only. It is not legal or tax advice, rules and procedures differ by jurisdiction and change, and I must confirm my position with a qualified adviser. Do not tell me how to respond, whether the position stated is correct, or what my rights or obligations are. Do not name any law, regulator or authority. Work only from the words in the letter.

Format: three sections – What it asks for, Dates and deadlines it states, What is unclear or unstated.

If a passage is ambiguous, quote it and say so rather than resolving it.

Letter:
[PASTE TEXT]

Check before you use the output: the deadlines, against the letter itself, word for word. A misread date is the one error here with an irreversible consequence.

P5.4 — Questions for your adviser

Use when: you have adviser time booked and want to use it well.

Paste this:

You are helping a finance team prepare for a meeting with its tax adviser about [TOPIC] at [ENTITY]. I will describe our situation and what we already know.

Task: write the questions we should ask, ordered so that the answers build on each other, and note for each what documents we should have to hand when we ask it.

Constraints: do not answer any of the questions. Do not indicate what the likely answer is, and do not state any rule or treatment. This is meeting preparation only, not advice; rules differ by jurisdiction and change.

Format: numbered questions, each with a sub-line "Bring:" listing the documents.

Add a final section, "Things we should tell them without being asked", for facts in my description that an adviser would want to know early.

Our situation:
[PASTE DESCRIPTION]

Check before you use the output: the final section. It is the most valuable part and the easiest to get wrong, because it depends on facts you may not have mentioned.

P5.5 — Obligations calendar

Use when: deadlines live in several people's heads.

Paste this:

You are helping a finance team build an internal compliance calendar for [ENTITY]. I will list the obligations we know we have, with the deadlines as we understand them.

Task: turn my list into a working calendar, with the internal preparation milestones that have to happen before each external deadline.

Constraints: use only the obligations and dates I give you. Do not add obligations you assume apply to us, and do not correct my dates – you do not know my jurisdiction and requirements change. Where you think I may have omitted something common, put it in a separate list of questions, phrased as questions.

Format: a table with columns External deadline, Obligation, Internal milestone, Days before, Owner. Sorted by date.

Add a line at the top stating that the dates come from me and must be confirmed against our own sources each year.

Obligations:
[PASTE LIST]

Check before you use the output: that no deadline has been silently adjusted, and that no obligation has appeared that you did not list.

6. Contracts, agreements and policy documents

Reading a contract for its finance consequences is not the same as advising on it. These prompts extract and organise. They do not interpret, and nothing here replaces your legal adviser.

P6.1 — Finance summary of an agreement

Use when: a signed agreement lands on your desk and you need to know what it does to your numbers.

Paste this:

You are helping a finance team understand the financial mechanics of an agreement. I will paste the text of an agreement involving [ENTITY].

Task: summarise only the provisions that affect money moving, timing, or what has to be recorded and tracked.

Constraints: this is a reading aid, not legal advice; I will have legal questions reviewed by a qualified adviser, and law differs by jurisdiction. Quote the clause reference and the operative wording for every point you make. Do not interpret ambiguous drafting – flag it. Do not state the accounting treatment. Do not summarise clauses you were not given.

Format: sections for Amounts payable or receivable, Timing and triggers, Variation and escalation mechanisms, Termination and its financial consequences, Anything creating an ongoing obligation to track.

End with "Clauses I could not interpret with confidence", quoting each.

Agreement:
[PASTE TEXT]

Check before you use the output: every quoted clause reference against the document. Quotation is the one thing you can verify quickly, and if the references drift, nothing else in the summary is safe.

P6.2 — Payment terms extraction

Use when: you are loading terms into a system or comparing across a portfolio.

Paste this:

You are extracting data from a contract for [ENTITY]. I will paste the text.

Task: extract the payment terms into a structured table for entry into our systems.

Constraints: extract only. Do not interpret, do not normalise into categories of your own, and do not fill an empty field with what is usual. Every value must appear in the text; give the clause reference for each. If a field is not addressed in the contract, write "not stated" – never a default.

Format: a table with columns Field, Value, Clause reference. Fields: currency, amount basis, invoicing frequency, payment period, late payment consequence, price adjustment mechanism, discount or rebate terms, expenses treatment, disputed amounts.

Below the table, list any provision that affects payment but did not fit these fields.

Contract:

[PASTE TEXT]

Check before you use the output: every "not stated" and every populated field with an odd clause reference.
Extraction failures cluster in documents with unusual numbering.

P6.3 — Obligations and dates register

Use when: the contract creates things somebody must remember to do.

Paste this:

You are building an obligations register from an agreement involving [ENTITY]. I will paste the text.

Task: list every obligation that falls on us, every deadline, notice period and renewal or expiry date, and every condition that triggers one.

Constraints: quote the clause reference and the wording for each entry. Do not calculate a date unless the contract states the reference point and the period, and where you do calculate one, show the working. Do not list obligations that fall on the other party unless they trigger one of ours. Do not interpret conditional wording – quote it.

Format: a table with columns Obligation, Owner function, Trigger or date, Notice period, Clause, Wording quoted.

Flag separately any obligation whose trigger you could not determine from the text.

Agreement:
[PASTE TEXT]

Check before you use the output: any calculated date. Notice periods interact with definitions of business days and month-ends, and that is where the arithmetic quietly goes wrong.

P6.4 — Questions for legal review

Use when: you have limited legal budget and want it spent on the right clauses.

Paste this:

You are helping a finance team prepare for legal review of a draft agreement involving [ENTITY]. I will paste the draft and describe the commercial deal we believe we have agreed.

Task: identify where the drafting and my description of the deal do not obviously match, and write the questions to put to our legal adviser.

Constraints: do not give legal advice, do not state what any clause means as a matter of law, and do not say whether a term is standard or reasonable. Law differs by jurisdiction and changes. Confine yourself to: this is what you said the deal is, this is what the words say, here is the question.

Format: a table with columns Commercial expectation, Clause and wording, The mismatch, Question for counsel. Ordered by financial exposure as I have described it.

If the drafting matches my description, say so for that point rather than inventing a concern.

The deal, then the draft:
[PASTE]

Check before you use the output: that the "commercial expectation" column reflects what you actually said, not a tidier version of it.

P6.5 — Internal policy drafting

Use when: you need a policy document that people will actually follow.

Paste this:

You are drafting an internal [POLICY TYPE] policy for [ENTITY]. I will tell you the rules we have decided, the limits we have set, and the behaviour we are trying to change.

Task: write the policy.

Constraints: include only rules I have given you. Do not add limits, thresholds, approval levels or exceptions of your own – if you think one is missing, list it as a question at the end. Do not reference any law or regulation; where a rule exists because of an external requirement, I will add that myself. Write for someone who will read it once, in a hurry. British English. No sentence longer than 25 words.

Format: Purpose, Who it applies to, The rules as numbered clauses, What to do if the rules do not cover your situation, Who owns this policy.

End with "Decisions you have not made yet" listing the gaps.

Our decisions:
[PASTE]

Check before you use the output: the gaps list, before circulating anything. A policy with an unmade decision inside it will be tested by the first person who finds the gap.

7. Board papers, reports and written communication

P7.1 — Board paper structure

Use when: you know what you need to say and not yet how to arrange it.

Paste this:

You are helping a finance director structure a board paper on [TOPIC] for [ENTITY]. I will tell you the decision or noting item required, what the board already knows, and the facts I hold.

Task: propose the structure – sections, what belongs in each, and roughly how long each should be.

Constraints: structure only. Do not draft the content, do not supply facts, and do not suggest a recommendation. Assume the board has limited time and will read the first page properly and the rest selectively.

Format: a numbered outline. For each section: purpose in one line, what goes in it, target length, and whether it belongs in the paper or an appendix.

At the end, list under "Facts you will need that you have not mentioned" the gaps a board would notice – as questions, not assumptions.

Background:
[PASTE]

Check before you use the output: the gaps list against your own knowledge. It is the closest thing here to a rehearsal for the meeting.

P7.2 — One page from twenty

Use when: the paper is written and the summary is due.

Paste this:

You are writing the executive summary of a paper I will paste. The audience is [AUDIENCE] and the decision required is [DECISION].

Task: write a summary of no more than one page.

Constraints: every statement must appear in the paper. Add nothing. If the paper contains a qualification on a key point, the qualification must survive into the summary – do not simplify a hedged conclusion into a firm one. Do not change any figure. Do not add a recommendation the paper does not make.

Format: the decision required in one sentence, three to five bullets of what the reader must know to make it, the main risk or uncertainty, and what happens next.

List separately any point where the paper is unclear about its own conclusion.

Paper:

[PASTE]

Check before you use the output: the hedges. Summarisation flattens qualified statements into confident ones, and the summary is the part that gets quoted back at you.

P7.3 — The difficult email

Use when: you have to tell someone something they will not like.

Paste this:

You are helping a finance professional write a difficult message. I will describe the situation, the recipient, the relationship, and the outcome I want.

Task: draft the message.

Constraints: state the substance in the first three sentences – no warm-up. Do not apologise for facts. Do not soften a deadline or a consequence I have given you, and do not add one I have not. No flattery, no false warmth, no exclamation marks. Do not characterise the recipient's likely feelings. British English.

Format: subject line, then the message. Under 200 words.

Then give me one alternative opening paragraph with a different balance of directness, and one line on when each is the better choice.

The situation:
[PASTE]

Check before you use the output: whether any commitment has crept in. Drafts of awkward messages tend to acquire an offer to help that you did not intend to make.

P7.4 — Anticipating the questions

Use when: you present tomorrow.

Paste this:

You are preparing a finance presenter for questions on [TOPIC] at [ENTITY]. I will paste my paper or notes and describe the audience.

Task: predict the questions, hardest first, including the ones that come from someone who has read only the summary and the ones that come from someone who has read the appendix closely.

Constraints: do not answer them and do not draft my answers – I need to know what I do not know. Do not invent facts about our business in the framing of a question. Where a question would be prompted by something absent from my paper, say what is absent.

Format: three groups – Questions the paper answers, Questions the paper raises but does not answer, Questions the paper does not anticipate at all. Within each, hardest first.

Say plainly if a likely question has no answer available from what I have given you.

My paper and audience:
[PASTE]

Check before you use the output: the middle group. Those are the questions you will be asked and are least prepared for.

P7.5 — Plain-English pass

Use when: a document has to be read by people outside finance.

Paste this:

You are rewriting a finance document for readers who are competent professionals but not accountants. I will paste the text.

Task: rewrite it so it can be read once and understood.

Constraints: keep every figure exactly as written. Keep every qualification. Where a technical term is unavoidable, keep it and add a short gloss in brackets on first use – do not replace it with a loose synonym. Do not simplify a conditional statement into an unconditional one. British English. Do not add examples, analogies or context of your own.

Format: the rewritten document, then a list of the terms you glossed and the gloss you used, so I can check each one.

If a passage cannot be made plainer without changing its meaning, leave it and list it under "Kept as written".

Document:
[PASTE]

Check before you use the output: every gloss. A slightly wrong definition of a technical term, in a document going to non-specialists, is harder to correct later than the jargon was.

8. Process documentation and team training

P8.1 — Procedure from your own narration

Use when: the process exists only in the hands of the person who does it.

Paste this:

You are turning a spoken description into a written procedure. I will paste a transcript or notes of someone describing how they perform [PROCESS] at [ENTITY].

Task: write the procedure as numbered steps that a competent newcomer could follow.

Constraints: include only steps described in the source. Do not add steps that seem obvious, do not reorder into a more logical sequence without flagging it, and do not add system names, field names or approvals not mentioned. Where the speaker says "usually" or "it depends", preserve the conditionality – do not turn it into a rule.

Format: numbered steps, each starting with a verb. Sub-steps where the source gives detail. A "Before you start" list of what you need to hand.

End with "Points the description left open", listing every place the speaker was vague, and a note where you changed the order.

Source:
[PASTE]

Check before you use the output: the conditionals. Turning "usually we check with the manager" into "check with the manager" changes the process, and turning it into nothing changes it more.

P8.2 — Desk instructions for a new joiner

Use when: somebody starts on Monday.

Paste this:

You are writing desk instructions for a new joiner in a finance team at [ENTITY], covering [TASK]. I will paste our existing procedure and describe what the new joiner does and does not already know.

Task: rewrite the procedure for someone doing it for the first time, without our internal shorthand.

Constraints: change no step and no sequence. Expand abbreviations only where I have told you what they mean – otherwise mark them [EXPAND THIS]. Do not invent screen names, menu paths or system behaviour. Add no reassurance and no encouragement.

Format: the instructions, with a "what this step is for" line under each step, and a final section "When to stop and ask" listing the situations where a newcomer should escalate rather than proceed.

List anything in our procedure that you could not make sense of.

Procedure and context:
[PASTE]

Check before you use the output: the "when to stop and ask" section, with the person who currently does the job. It is the part that prevents expensive mistakes and the part only they can write.

P8.3 — Training exercises

Use when: you are teaching a technique rather than a task.

Paste this:

You are designing practice exercises for a finance team learning [TECHNIQUE OR TASK] at [ENTITY]. I will describe the level of the team and the mistakes they currently make.

Task: design [NUMBER] exercises that would surface exactly those mistakes, in increasing difficulty.

Constraints: exercises must use fictional data that is clearly labelled as fictional, and no exercise may present invented figures as real market, industry or benchmark data. Do not reference any real organisation. Each exercise must have a definite right answer, or, if it is a judgement exercise, a definite list of what a good answer must address.

Format: for each – the scenario, what the learner must produce, the mistake it is designed to surface, and the marking notes.

If a mistake I described cannot be surfaced by an exercise, say so and suggest what would address it instead.

Team and mistakes:
[PASTE]

Check before you use the output: the marking notes. If they are vague, the exercise will produce an argument rather than learning.

P8.4 — Handover note

Use when: somebody is leaving, or going on leave, or moving on.

Paste this:

You are building a handover note for [ROLE] at [ENTITY]. I will describe what the role does, in whatever order it comes to me.

Task: organise it into a handover document, and identify what is missing.

Constraints: use only what I give you. Do not assume responsibilities based on the job title. Do not assign deadlines or contacts I have not named. Where I have described something vaguely, keep my words and mark it "needs detail".

Format: Recurring tasks by frequency, One-off items in progress with status, Relationships and who to contact for what, Access and permissions required, Known problems and workarounds, Things only this person knows.

The last section matters most. Prompt me with questions to fill it – the undocumented knowledge that walks out of the door.

What the role does:
[PASTE]

Check before you use the output: the questions it asks you at the end, and answer them before the person leaves. After they leave, the note is whatever it is.

P8.5 — Explaining a spreadsheet to its future maintainer

Use when: a model has grown and only one person understands it.

Paste this:

You are documenting a spreadsheet formula so that a future maintainer can understand it. I will paste the formula and describe the sheet layout and what the calculation is meant to achieve.

Task: explain what the formula actually does, step by step, and whether that matches my description of the intent.

Constraints: work only from the formula text and the layout I describe. Do not assume ranges, named ranges or sheet contents I have not given you. Do not rewrite the formula unless I ask. If the formula's behaviour in an edge case cannot be determined from what I have provided, say so.

Format: what it does in one sentence; then a numbered breakdown from the innermost operation outwards; then Edge cases, covering blanks, errors, text where a number is expected, and ranges that grow.

End with "Where this may not match your stated intent", or state plainly that it appears to match.

Formula, layout and intent:
[PASTE]

Check before you use the output: the edge cases, by testing them in a copy of the file. An explanation of a formula is a hypothesis until the spreadsheet agrees with it.

Building your own

Every prompt above is the same six parts in the same order, and once you see the frame you can write your own in about two minutes.

Part	What it does	Typical failure if you omit it
Role	Sets who is speaking and to whom	Generic writing pitched at nobody
Task	States the single job	Several half-done jobs
Context	Supplies the facts only you have	Plausible invention filling the gap
Constraints	Rules out what you do not want	Added figures, added causes, added cheer
Format	Names the shape of the output	Reformatting by hand afterwards
Uncertainty rule	Says what to do when it does not know	Confident, fluent, wrong

The last one is the part people leave out, and it is the part that makes the difference between a tool and a liability. "If you are unsure, say so" is a weak version. Better ones name the behaviour precisely: put unclassifiable items in a separate list; quote the clause you could not interpret; state what information you would need; say plainly that you cannot reconcile the components to the total. Give it somewhere to put its ignorance and it will use it. Give it nowhere, and it will smooth the gap over with something that reads well.

Two habits are worth forming. Write the check before you write the prompt — decide what would be wrong in the output and how you would spot it, and if you cannot answer that, you are not ready to use the output for anything that matters. And keep your prompts in a file where your team can see and improve them, with a line under each recording what went wrong the last time someone used it. A prompt library that nobody edits is a library of your first attempts.

Verification, Red Lines and Choosing a Tool

Everything in this section is general information for professional people thinking about how they work. It is not legal, regulatory, tax or professional advice. Rules differ by jurisdiction, firms differ in what they permit, and both change. Check your own position with your regulator, your professional body and, where it matters, your own advisers.

Part One — The Three-Check Rule

You already have a review habit, tuned over years to the mistakes humans make: transposed digits, a stale opening balance, a paragraph copied from last year and not updated. AI output fails differently. It fails smoothly. The wrong number is formatted correctly, sits in the right column, and is surrounded by three that are right. The invented source has a plausible title. The reasoning reads like reasoning.

So the habit needs a small extension, not a replacement. Three checks, in order, every time output leaves your hands: **the figures, the sources, the reasoning.**

The order matters. Figures first, because a wrong figure makes everything downstream irrelevant. Sources second, because that is where fabrication hides. Reasoning last, because it takes the most thought and you should not spend that thought on a document that has already failed.

The whole thing should take a few minutes on short work and under half an hour on something substantial. If it would take longer, the work was the wrong shape to give to a tool.

Check one — the figures

What you do. Take every number in the output and trace it back to where it came from. Not "does it look right" — trace it. Open the source file. Find the figure. Match it. Then check that the arithmetic between figures holds: the components add to the total, the variance equals the difference, the percentages are of the base you think they are of, the periods compared are the same length.

Numbers you did not supply get harder treatment. If a figure appears in the output that was not in your input, it came from somewhere you cannot see. Delete it or verify it independently. There is no third option.

What it catches. Numbers altered in transit — rounded when they should not have been, pulled from the wrong row, carried from one entity to another, restated to a different currency without saying so. Totals that do not foot. Percentages calculated the wrong way round. Prior-period figures that quietly became current-period. Numbers invented to fill a sentence that needed one.

What it misses. A number that is faithfully copied from a source that is itself wrong. The check verifies transmission, not truth. It also misses the number that should be there and is not — an omitted provision, a missing intercompany line, a segment left out of a total you did not think to foot. Absence is much harder to see than error.

How long. Aim for a minute or two per page of numbers. If you cannot trace a figure inside about thirty seconds, that is not a slow check, that is a finding — write it down and go back to whoever produced it.

A prompt that makes this check faster, because it forces the output to show its own working:

Before the answer, produce a table of every number you are about to use. Columns: the number, the exact label or cell reference it came from in the material I gave you, and one of these two words – "supplied" if it came from my material, "derived" if you calculated it. For anything derived, show the calculation. Do not include any number you cannot place in one of those two categories. If a figure I need is missing from my material, say so and leave it blank rather than estimating it.

Material: [PASTE OR ATTACH]

Task: [WHAT YOU WANT PRODUCED]

You still trace the figures, but now against a list the tool has committed to, which is faster than hunting through prose. The "derived" rows tell you which calculations to re-perform.

Check two — the sources

What you do. For every factual claim that is not a figure — a rule, a treatment, a deadline, a definition, a contractual term, "the standard requires", "market practice is" — find the source and read it yourself. The primary document, not a summary of it. If the output names a source, confirm it exists, confirm it says what the output says it says, and confirm it is current.

If the output does not name a source, that claim has no source. Treat it as a hypothesis you now have to prove or drop.

What it catches. Confident references to things that do not exist, or that exist but say something else. Rules described as they were several years ago. Guidance from one jurisdiction presented as though it applies in yours — this one is common and easy to miss, because the language of finance is shared across borders while the rules are not. Thresholds and dates that have moved. Requirements that apply to a class of entity yours is not in.

What it misses. A source that is real, current, correctly quoted, and beside the point. Relevance is a judgement no check can automate. It also misses the source that is right for your jurisdiction and wrong for your client's, and the well-known position that has an exception your client happens to sit inside.

How long. Budget two to five minutes per claim you actually rely on. That sounds heavy until you notice how few claims in a document carry weight. Mark the ones that do — the ones a reader would act on, or that would appear in a complaint — and check those properly. Skimming twelve claims is worth less than checking three.

A useful move here is to separate the drafting from the sourcing:

List every factual assertion in the text below that is not a number, one per line. For each, mark it "verifiable" if a reader could check it against a document, or "judgement" if it is an opinion or an inference. For each verifiable assertion, state what kind of document would settle it. Do not tell me the answer, do not cite anything, and do not add assertions that are not already in the text.

Text: [PASTE THE DRAFT]

That gives you a checking list rather than a citation list. It is more honest, because it does not invite the tool to produce references, which is exactly where fabrication lives.

Check three — the reasoning

What you do. Close the output. In your own words, out loud or on paper, state the argument: what the conclusion is, what it depends on, and what would have to be true for it to hold. Then reopen the output and compare. You are looking for three things — a step that is asserted rather than argued, an assumption that was never stated, and a conclusion that is wider than the evidence under it.

Then run the inversion. Ask yourself what the strongest argument against the conclusion is. If you cannot construct one, you do not yet understand the position well enough to send it.

What it catches. Confident conclusions built on one unrepresentative period. Correlation dressed as cause — "margin fell because of the pricing change", when three things changed that quarter. Recommendations that quietly assume the reader's circumstances. Analysis that answers a slightly different question from the one asked — the most common failure, and the hardest to see, because the answer is good.

What it misses. Anything that depends on knowing the client, the history, the personalities and the thing the last auditor flagged. It misses the reasoning that is sound in general and wrong here. And it misses your own blind spots — you are checking with the same head that approved the approach.

How long. Five to ten minutes on anything carrying a recommendation; less on descriptive work. It is the check people skip when busy, and the one that catches the errors that embarrass people: figure errors get found by the next reviewer, reasoning errors get found by the client.

The prompt that helps here is the one that argues with you:

Argue against the conclusion in the text below. Give the three strongest objections a well-informed sceptic would raise, the assumption each objection attacks, and what evidence would settle it. Do not soften them and do not offer a balanced view. Do not rewrite the text.

Text: [PASTE THE DRAFT]

Read the objections as prompts for your own thinking, not as findings. Some will be weak. The point is to make the shape of the argument visible so you can test it yourself.

The verification card

Print this. Pin it where the work gets reviewed. It is deliberately short, because a checklist nobody completes is worse than no checklist — it produces the paperwork of assurance without the assurance.

THE THREE-CHECK CARD

Before this leaves my hands.

1. FIGURES — trace, do not glance

- Every number traced to a named source file, cell or document
- Totals footed; variances re-performed; percentages checked against the right base
- Periods, entities and currencies confirmed to be the ones stated

- Any number I did not supply: verified independently, or deleted

2. SOURCES — *read the original, not the summary*

- Every claim I am relying on has a source I have opened myself
- The source is current, and applies in this jurisdiction and to this entity
- Unsourced claims: proved or removed
- Nothing quoted that I have not read in full

3. REASONING — *state it without looking*

- I can state the conclusion and what it depends on, in my own words
- Every assumption is on the page, not in my head
- The conclusion is no wider than the evidence under it
- I can make the strongest case against it — and it still stands

SIGN-OFF

- This answers the question that was actually asked
- I can defend every line of it, unaided, to someone who challenges it
- Where I am uncertain, the document says so
- If this is wrong, I am the one who is wrong. That is correct.

Checked by: [NAME] Date: [DATE] Time taken: [MINUTES]

Log the time taken. Not for anyone else — for you. After a fortnight you will know which work is genuinely faster with a tool and which is only faster until the checking starts. It is the only productivity figure in this pack, and the only one worth having, because it is yours and you measured it.

What the three checks do not catch

They do not catch everything, and a pack that claimed otherwise would be the sort of thing this pack exists to be an alternative to.

Here is what remains after all three checks are done properly.

The plausible omission. The checks are all directed at what is on the page. None of them looks for what should be there and is not. A tool that summarises fifty documents and silently drops the one that matters produces output that passes every check.

The correct-but-inapplicable. Right rule, right numbers, wrong client. Nothing in the checks tests fit.

Compounding. Verified output becomes input to the next piece of work, and the check does not repeat. Three steps down the chain, the assumption is load-bearing and nobody remembers it was an assumption.

Drift. The tenth document of the day gets a thinner check than the first, and reviewing fluent output is not the same mental act as writing it. Read enough finished-looking work and the reading gets lighter without you deciding to make it so.

Four things to do about the residual risk, none of which eliminates it:

1. **Write down what you assumed.** A short assumptions block at the top

of anything you pass on. It costs a minute, and it is the most useful thing to have when someone queries the work six months later.

1. **Keep the input.** Store the material you supplied alongside the output. If you cannot reconstruct what the tool was given, you cannot later work out how it went wrong.

1. **Vary who checks.** The person who framed the task is the worst person to check it, and the second-worst is anyone who agrees with them. Get fresh eyes on anything that carries real consequence.

1. **Set a ceiling.** Decide, in advance and in writing, the value or sensitivity of work above which the tool is not used at all. Deciding in the moment, under deadline, is deciding badly. Part Two is a starting draft of that ceiling.

Part Two — Ten Places Not to Use It

There is a way of writing about this that treats every limitation as a temporary inconvenience soon to be engineered away. This is not that. Some of what follows may change; most will not, because the reason is not capability. It is that a duty sits on a person, and a duty cannot be delegated to something that cannot be held to it.

For each: the situation, why the tool fails there in particular, and what to do instead.

1. Signing, certifying or expressing an opinion

The situation. Anything where your name, signature or firm goes on a document as the person vouching for it — accounts, certificates, confirmations, formal opinions, representations to a third party who will act on them.

Why it fails. Your signature is not a statement that the words are good. It says *you* did the work and *you* stand behind it. The signature is about accountability, not text, and a tool cannot supply accountability. Sign something you did not verify and the signature is false, whether or not the content turns out to be right.

Instead. Use a tool, if at all, on drafting and checking well before signature — never on the substance of what is certified. If a tool touched the draft, the person signing must have re-performed the underlying work themselves, not reviewed the prose.

2. Anything you cannot verify

The situation. The output is in a subject, language, jurisdiction or technical area where you have no way to check whether it is right.

Why it fails. Every safeguard in Part One assumes a checker who can tell good from bad. Remove that and the tool is not an assistant, it is an oracle. Worse, fluency scales with unfamiliarity: output in an area you do not know reads *more* convincing than output in an area you do, because you cannot see the joins.

Instead. Either learn enough to check it, find someone who can, or do not send it. A tool used to orient yourself in an unfamiliar area gives you a list of things to look up, never the answer. The tell: if you could not answer a follow-up question without going back to the tool, you do not own the work.

3. Live disputes, investigations and privileged matters

The situation. Litigation, arbitration, a regulatory investigation, a disciplinary process, a dispute with a former employee, anything where a lawyer is involved or is about to be.

Why it fails. Two reasons, both serious. In a contested matter the working papers may become visible to the other side, and a draft produced by a tool — including the discarded versions and the framing of your questions — is a record of how you thought about the case. And protections attaching to communications with your lawyers can be fragile; putting the same material through a third-party system may affect them, depending entirely on your jurisdiction.

Instead. Take instructions from the lawyers on the matter before anything goes near a tool, and get the answer in writing. Assume, until told otherwise, that nothing about a live dispute goes into any external system. This is general information, not legal advice; the rules here are jurisdiction-specific and consequential.

4. Decisions about a named individual

The situation. Pay, promotion, redundancy, hiring, performance ratings, discipline, creditworthiness, a reference, a decision to end someone's contract.

Why it fails. The decision has to be explicable to the person affected and possibly to a tribunal or a regulator, and "the system suggested it" is not an explanation. The tool has no access to the context that matters — the year they had, the thing their manager did, the reason the numbers look like that. And several jurisdictions place obligations around automated processing of personal data and decisions that significantly affect people; those vary and change, and you need to know your own.

Instead. Make the decision yourself, on evidence you can point to, and write your reasons before you draft any communication. A tool may help you make an already-made decision clearer or kinder in the telling. It must not help you reach it, and it must not receive the person's personal data to do so.

5. Anything where the client's data cannot leave your control

The situation. Data covered by a confidentiality undertaking, a data processing agreement that names permitted sub-processors, a client policy prohibiting external processing, or a regulatory requirement to keep information within a jurisdiction. Also, in practice: anything containing identifiable personal data, bank details, board material before it is public, or price-sensitive information.

Why it fails. Not because the tool is careless. Because pasting the data into it is a disclosure to a third party, and you either had permission for that or you did not. The tool's own security is not the question; whether you were allowed to send it is, and that is settled by contracts you signed before the tool existed.

Instead. Check the engagement terms and data agreements before the work starts, not during it. Where the answer is no, either work without the tool or use one deployed inside your own environment under terms you have read. When in doubt, redact: much analytical work can be done on structure and anonymised figures without names, account numbers or identifiers leaving your systems.

6. Matters where a professional duty attaches to you personally

The situation. Independence and conflict assessments. Ethical judgements. Deciding whether to accept or resign from an engagement. Deciding whether something must be reported, and to whom. Anything your professional body would describe as a matter for your own judgement.

Why it fails. These duties sit on you as a person precisely because they exist to be exercised when it is uncomfortable — when the client is large, the fee is significant and the easy answer is available. A tool produces the conventional answer, usually the comfortable one, and it has nothing at stake. That is the problem: the duty exists because you have something at stake.

Instead. Decide it yourself, document the reasoning at the time in your own words, and consult a human — a partner, an ethics helpline, your professional body — where you are unsure. The written record of your reasoning is the point, and it has to be yours.

7. Suspicions of fraud, error or dishonesty

The situation. Something looks wrong and a person may have done it deliberately.

Why it fails. These situations carry reporting obligations that vary by jurisdiction and can be time-critical, and some carry prohibitions on telling anyone what you suspect. A drafting tool cannot know your obligations or who you are permitted to speak to. There is also a practical trap: writing up a suspicion in a system you do not control creates a record you may not be able to retrieve, delete or explain.

Instead. Follow your firm's escalation procedure immediately, and if you do not know it, find out today rather than the day you need it. Take advice on your reporting obligations from someone qualified to give it. Keep your own contemporaneous handwritten or locally stored note. Nothing external, nothing informal.

8. Advice that turns on the specific rules of a place

The situation. Tax positions, filing obligations, statutory deadlines, entity-specific reporting requirements, licensing, anything where the answer changes when you cross a border or a threshold.

Why it fails. These rules are numerous, local and revised often. Material describing them ages badly and is not evenly available across jurisdictions, so the answer tends to reflect the places most written about rather than the place you are in. That is dangerous because the *structure* of the answer is usually right — the right considerations, in the right order — while the threshold, rate or date inside it is wrong or foreign.

Instead. Use the primary source published by the relevant authority, or a subscription service that states its currency date. A tool may help you frame the question or draft the covering letter. The rule itself comes from the source every time, and you note the date you checked.

9. The record of what happened

The situation. Minutes, file notes, records of a conversation, attendance notes, evidence of what was said and decided.

Why it fails. A record is not a summary. Its value is that it is a first-hand account by someone who was present, which is exactly what a generated version is not. A tool asked to tidy up rough notes smooths them, fills gaps with what usually happens, and turns the awkward half-sentence someone said into a clean sentence they did not. In a dispute that difference is everything, and the smoothing is invisible once done.

Instead. Write records yourself, close to the event, in plain language, including the parts that are unresolved or unflattering. If you use a transcription tool, keep the transcript as the record and mark it as such — and check what your policy and the law in your jurisdiction say about recording people, which is a genuine question and not a formality.

10. Anything you would not be able to defend line by line

The situation. The general case. You are about to send work you could not walk through, unaided, if the recipient rang and asked why.

Why it fails. Not a limitation of the tool at all. A limitation of the arrangement. Work you cannot defend is work you have not done, and whether it was a tool, a junior or a template that produced it changes nothing about your position.

Instead. Either do the work until you can defend it, or say plainly that you have not done it. "I have not verified this — it is a starting point" is a perfectly professional sentence, and it is the honest one. Most of the trouble people get into with these tools starts with an unwillingness to say it.

Part Three — Assessing a Tool Before It Touches Client Data

Marketing pages are written to be reassuring. Contracts are written to be enforceable. They are different documents and they routinely say different things. The purpose of this section is to get answers in the second kind of document.

Work through the table below with the vendor, or with your own IT team if the tool is deployed internally. Two rules make it worth doing. Get the answers in writing, with the clause reference: "yes, of course" is not an answer, "clause 8.3 of the data processing addendum" is. And if a question comes back vague after you have asked twice, record the vagueness as the answer. It is information.

Fill in the third column with what they said and the fourth with your own decision. The fourth column is a judgement your firm makes, not one a checklist makes for you.

What happens to our input

#	Question to ask	Their answer	Acceptable?
1	Is anything we submit used to train, tune or improve your models, or anyone else's?		
2	If it is, how do we switch that off, is the switch on our plan, and does it apply to every feature or only some?		
3	Is our input processed by any third party — an underlying model provider, a hosting provider, a subcontractor? Name them all.		
4	Do you keep our input to review it for safety, abuse or quality? Who reviews it and for how long is it kept?		
5	Does any of this change for uploaded files, images or connected data sources, as against typed text?		

Where it is stored, and for how long

#	Question to ask	Their answer	Acceptable?
6	In which countries is our data stored, processed and backed up? List every one.		
7	Can we require it to stay in a named region, and is that a contractual commitment or a setting you can change?		
8	How long is our input retained by default, and can we shorten it?		
9	How long does it persist in backups, logs and caches after deletion?		
10	Is it encrypted at rest and in transit, and who holds the keys?		

Who at your end can see it

#	Question to ask	Their answer	Acceptable?
11	Which of your staff can access our content, under what circumstances, and with what approval?		
12	Is that access logged, and can we see the log?		
13	Do your staff sign confidentiality undertakings that would cover our clients' information?		
14	Can government or law-enforcement requests reach our data, and do you commit to telling us when one does, where you are permitted to?		
15	Have you had a security incident affecting customer content, and what is your contractual notification period if you do?		

What happens when we leave

#	Question to ask	Their answer	Acceptable?
16	On termination, what is deleted, when, and how is deletion evidenced?		
17	Can we export everything — content, history, configuration — in a usable format, and is that a right in the contract?		
18	What survives termination in your favour: licences to our content, aggregated data, derived material?		
19	If you suspend or terminate our account, do we still get our data out, and for how long?		
20	If the service is discontinued, what notice do we get?		

The contract versus the marketing page

#	Question to ask	Their answer	Acceptable?
21	Which document actually governs: the terms of service, the data processing addendum, the security page, or the sales deck? Point to the order of precedence.		
22	Which of the claims on your website appear as obligations in the contract? Name the clauses.		
23	Can you change the terms unilaterally, with what notice, and can we exit without penalty if you do?		
24	What is your liability cap, and what is excluded from it?		
25	Do you accept the role of processor, on our instructions, for the data we submit — and is that written down?		
26	Do you have current independent security certification or an audit report, and may we see the actual report rather than the badge?		

What changes if you are acquired

#	Question to ask	Their answer	Acceptable?
27	On a change of control, do our terms transfer unchanged, and for how long?		
28	Is our data an asset that could be sold or transferred in insolvency or acquisition?		
29	Do we get notice and a right to terminate on a change of control?		
30	If your funding or ownership changes, what happens to the region and retention commitments in questions 6 to 9?		

Two practical notes on using this. Question 22 is the one that produces the most silence, and the silence is the finding. And a sensible way to handle the paperwork is to have a tool extract, not interpret:

Read the attached terms and produce a table with three columns: the question below, the exact quoted wording from the document that addresses it, and the clause number. If the document does not address a question, write "not addressed" – do not infer, summarise or fill the gap. Quote verbatim. Do not tell me whether the terms are good.

Questions:

[PASTE THE QUESTIONS FROM THE TABLE ABOVE]

Document: [ATTACH THE TERMS]

That gives you a map of where to read. You still read the clauses yourself, and anything that matters commercially or legally goes to someone qualified to advise on it. Extraction is a task a tool does well. Deciding whether a limitation of liability is acceptable is not.

The assessment has a shelf life

An assessment describes a supplier on the day you did it, and that is all it describes. Terms are revised. Sub-processors are added. Features arrive that handle data differently from the feature you assessed, and they arrive

switched on. Free and paid tiers can differ on exactly the points you cared about, so an upgrade can change your position without anyone deciding to change it. Companies are bought. Retention settings get reset by a migration.

None of this is bad faith. It is what happens to software companies. The consequence is that a one-off assessment filed away is a record of a decision, not evidence of a current position.

So: repeat it. Once a year for anything touching client data, and immediately on any of these triggers — you are told the terms have changed, the ownership changes, you start using a materially new feature, your plan changes, a client asks you a direct question about where their data goes, or you take on work with a confidentiality regime stricter than your usual one. Keep the completed tables with the date and the name of the person who filled them in, so that the next review is a comparison rather than a fresh start.

And keep one thing in your own hands regardless of what any assessment says: know which of your clients' information must never be submitted to an external system at all, write that list down, and make sure everyone who could paste something has read it. Vendor assurances describe what someone else promises to do. That list describes what you control.

Your Policy Documents

What is in this section

Three documents, in ascending order of length:

1. **The one-page team note** — for a team of two to fifteen who need a rule by Friday.
1. **The firm policy** — for a practice with clients, engagement terms and a regulator.
1. **The client-facing paragraph** — the wording you put in an engagement letter or a proposal when a client asks how you use these tools.

Take whichever fits, put your own name in the brackets, delete what does not apply, and have someone who carries risk read it before you circulate it. All three are starting drafts and general information, not legal advice. Rules differ by jurisdiction and profession and they change.

There is also a free web version of the one-page note that assembles it from nine questions. It is on the same site this pack came from, under Free tools. It runs entirely in your browser and collects nothing.

1. The one-page team note

USING AI TOOLS — [TEAM OR FIRM NAME]

Applies to: everyone in the team, including contractors. **Owner:** [name and role]. **Next review:** [date].

Why this exists. We use these tools because they save real time. They also produce confident, well-written output that is sometimes wrong. This note exists so the time saved is not paid back later with interest, and so nobody has to guess what is allowed.

Approved tools. [List them.] Anything not on this list is not approved yet — that is not a refusal, ask and the list can grow. Personal accounts and free consumer tiers are not approved for any work involving client material.

What may go in. [Choose one: no client or personal data at all / client data only in the approved tools whose terms state our input is not used for training / anonymised material only.] Identification numbers, bank details and payroll figures for named individuals stay out regardless of the tool. If you are unsure whether something counts, it counts. Ask first.

What you check before it goes out. Every figure, every quotation, every reference, every claim about a rule — against its source. Recalculate anything numerical independently. Never cite a standard, a case or a section you have not opened yourself.

What we tell clients. [Choose one: internal aid, not routinely disclosed / we answer honestly and specifically when asked / we state it in engagement terms.] We never deny using these tools.

Where they are not to be used. Signing, approving or certifying anything. Any final decision about a person's pay, employment or standing. Anything under privilege, live dispute or investigation without written approval. Anything you would not be able to check.

Records. Where a tool materially helped produce a deliverable, note it in the working papers: which tool, what for, who checked it. One line is enough.

If something goes wrong. Tell [owner] the same day. Say plainly what happened. Nobody is disciplined for reporting quickly. Concealment is the only thing treated seriously.

2. The firm policy

Longer, for a practice that has clients and a regulator. The clauses below are drafting blocks — keep the ones that apply.

2.1 Purpose and scope

This policy governs the use of artificial intelligence tools in the work of [FIRM]. It applies to all partners, employees, contractors and temporary staff, on any device, whether the tool is provided by the firm or not. It covers general-purpose assistants, AI features embedded in software the firm already licenses, transcription and note-taking tools, and any tool that sends firm or client material to a third party for processing.

It does not change any existing duty. Where this policy and a professional obligation appear to conflict, the professional obligation governs and you raise it with [ROLE] immediately.

2.2 Responsibility

The person who signs, reviews or issues a piece of work is fully responsible for it. That responsibility is not reduced, shared or transferred by the use of any tool. "The system produced it" is not an explanation the firm will offer to a client, a regulator or a court, and it is not one you may offer internally.

2.3 Approved tools and how the list changes

Only tools on the approved register may be used for client work. The register is maintained by [ROLE] and records, for each tool: the provider, the licence tier, whether firm input may be used to train the provider's models, where data is processed, the retention period, and the date the terms were last reviewed.

Requesting an addition is normal and encouraged. Using an unapproved tool on client material is a disciplinary matter, not because the tool is suspect but because the firm cannot answer for what it cannot see.

2.4 Client and personal information

[Choose the stance that matches the firm's risk appetite and delete the others.]

Option A — prohibition. No client information, personal data or confidential material may be entered into any AI tool. This includes names, initials, entity names, account numbers, identification numbers, addresses and figures attributable to a named person or client, and it includes combinations of details that would identify someone together even if each is innocuous alone.

Option B — approved tools only. Client information may be entered only into tools on the approved register whose terms state that firm input is not used to train the provider's models and is not retained beyond the period recorded on the register. It may never be entered into a personal account or a consumer tier.

Option C — anonymised only. Only material from which identifying details have been removed may be entered. Where anonymising would take longer than performing the task unaided, perform the task unaided.

In every case: identification numbers, bank account details, payroll figures for named individuals, health information and anything received under an explicit confidentiality undertaking stay out.

2.5 Verification

Before AI-assisted work leaves the person who produced it:

(a) every figure is recalculated or agreed to source; (b) every quotation, reference, standard and authority is opened and read in the original; (c) the reasoning is read as though a reviewer were trying to break it, with particular attention to anything stated confidently and without qualification; (d) [for higher-risk work] a second qualified person reviews the deliverable and records that they have done so.

Fluent output is not verified output. These tools are at their most persuasive on the things they have invented.

2.6 Disclosure

[Choose one.] The firm does not deny the use of these tools in any circumstances. Where a client asks, the answer given is honest and specific: what the tool was used for, what it was not used for, and who checked the result.

2.7 Prohibited uses

AI tools are not to be used to: sign, certify, approve or give an opinion; make a final decision about any individual's employment, remuneration, credit or professional standing; process any matter under legal privilege, live dispute or regulatory investigation without the written approval of [ROLE] for that specific matter; produce anything the user would be unable to verify; or generate material presented as the independent work of a person who did not review it.

2.8 Records

Where an AI tool materially contributed to a deliverable, the working papers record the tool, the purpose, and the reviewer. This is not bureaucracy. It is what allows the firm to answer a question about the work two years after the engagement closed.

2.9 Incidents

Any suspected error in issued work, or any client material entered somewhere it should not have been, is reported to [ROLE] on the day it is noticed. Prompt reporting attracts no sanction. Concealment does.

2.10 Review

This policy is reviewed every [3/6/12] months by [ROLE], and sooner if a tool on the register materially changes its terms, if the firm's professional body or regulator issues guidance, or if an incident suggests a clause is not working. Superseded versions are retained.

3. The client-facing paragraph

For engagement letters and proposals. Short on purpose — clients want to know you have thought about it, not to read your internal policy.

Use of technology. We use software tools, including tools that use artificial intelligence, as an aid in performing our work. They are used under a written internal policy which restricts what information may be processed, requires that outputs be verified against source before use, and prohibits their use for judgements and approvals. The individuals named in this engagement remain fully responsible for the work and for every conclusion in it. We are happy to describe our approach in more detail on request.

If your engagement terms already contain a confidentiality or subcontracting clause, check that this paragraph sits consistently with it rather than beside it.

Before you adopt any of this

These are drafts assembled from general practice, not legal advice, and not a statement of compliance with any law, standard or professional requirement. Requirements differ between jurisdictions and professions and they change. Have whoever carries risk in your organisation read the version you intend to use, and check it against your regulator's current guidance and your own client engagement terms.

The 90-Day Plan and the Training Curriculum

Ninety days is long enough to change how a finance team works and short enough that one person can hold the whole thing in their head. What follows is a plan for that period, then a curriculum you can teach yourself, in your own meeting room, with no outside help and no budget beyond people's time.

The plan assumes a finance function of somewhere between five and fifty people. Smaller than that and you compress the phases; larger and you run this in one department first and treat the rest of the organisation as a later problem. It also assumes you already have access to some tool — a general-purpose assistant, an enterprise-tier version of one, or an AI feature inside software you already pay for. Choosing the tool is a separate exercise. This section is about what happens after the licence exists.

Two warnings before you start. First, none of the numbers in this section are benchmarks, because there are none worth quoting; the only numbers that will matter to you are the ones your own team records. Second, this is general guidance on running a change in a team. Anything touching confidentiality, client consent, professional standards, data protection or your engagement terms differs by jurisdiction and changes over time. Check your own position with your professional body, your regulator and whoever handles your data obligations before you commit to anything here.

Part one: the ninety-day plan

What you settle before day one

Four decisions, made once, in writing. They take an hour. Skipping them is what causes most of the mess later.

Who owns this. One named person, not a committee. They do not have to be the most senior person in the room and they should not be your busiest month-end resource. They need enough standing to say "not yet" to a partner.

What you will measure. Two things only: the time a task takes end to end, and whether the output needed rework. Not how many people logged in. Not how many prompts were run. Enthusiasm is not a measure; it is a mood, and it evaporates around the twentieth of the month.

The one rule that applies from day one, before any policy is written: nothing produced with an assistant leaves the team — not to a client, not to a board, not to a regulator, not to a colleague outside finance — without a named human who has checked it and is willing to have their name on it. That sentence can go in an email this afternoon. The full policy comes later.

What is not going in. Decide, now, what data does not get pasted into anything until the policy exists. As a starting position: client-identifiable information, personal data, anything under an NDA, anything price-sensitive, anything you would not send to an external printer. Whether that starting position is right for you depends on your tool's terms, your contracts and your local rules. Read the terms; do not assume.

Role	Who it usually is	What they own	Realistic time cost
Sponsor	FD, controller, managing partner	Says yes to the pilot, protects the time, hears the results	1 hour a month
Owner	A manager with credibility	The plan, the log, the policy draft, the sessions	Half a day a week for 90 days
Pilot group	3–5 people across levels	Doing the two use cases and recording honestly	2–3 hours a week
Reviewer	A senior who is already the quality backstop	Signing off pilot output, flagging what they had to fix	Their normal review time
Data/IT contact	Whoever handles systems and contracts	Tool terms, access, what data may go where	2 hours in month one

Phase one — Days 1 to 30: contain it, then choose

The goal. Two things by day 30: a written policy people have actually read, and two chosen use cases with a baseline measurement. Nothing else. You are not trying to produce results this month. You are trying to stop the uncontrolled version of this happening while you think.

Week 1 — find out what is already happening. The owner asks every person in the team, individually and without consequence, what they have already used and what for. Do this in person or in a short call, not in a form. Say plainly that nobody is in trouble and that you need the truth to write a sensible policy. You will usually find more than you expected, including someone who has been pasting things they should not have. React to that with a fix, not a disciplinary — this week is the only chance you get to hear the truth cheaply.

Who: owner. *Done when:* you have a written list of current use, tool by tool, person by person, and you have told the whole team the one rule above.

Week 2 — draft the policy. Short. Two pages at most, because nobody reads four. It needs to cover: which tools are approved and which are not; what data may never be entered and what may; that a named human reviews everything before it leaves; how you record that a review happened; what to do when someone thinks the tool got something wrong; and who to ask when the policy does not cover the situation. Have your data/IT contact check it against your tool's terms and your obligations. If you have client confidentiality duties, get the consent question answered properly rather than assumed — that answer differs by jurisdiction and by engagement letter.

Who: owner drafts, sponsor approves, data/IT contact checks. *Done when:* the policy exists, the sponsor has signed it, and it has been sent to everyone with a request to reply "read".

Week 3 — choose two use cases. Two. Not ten. The temptation is to list every candidate and start them all, which produces ten half-finished experiments and no evidence about any of them. Pick using four tests: the task happens at least monthly; it currently eats time you can name; a wrong answer is caught by someone before it does damage; and you can tell within a week whether the output was good. Good early candidates tend to be drafting and summarising work with a strong review step already around them. Poor early candidates are anything where an error goes straight out of the door.

Who: owner and pilot group, sponsor decides ties. *Done when:* two use cases are written down in one sentence each, with the name of the person who reviews the output.

Week 4 — measure the "before". For each of the two use cases, record how long it takes now and how much rework it currently needs, using the log in the table below. Be honest that this baseline is an estimate — people are poor at recalling how long things took, and the act of measuring changes the number. Record it anyway, and label it as an estimate wherever you quote it later. A soft baseline you admit is soft beats a precise-sounding figure you invented.

Date	Task	Who	Minutes, unassisted (est.)	Minutes, assisted	Review minutes	Rework needed?	What went wrong
------	------	-----	----------------------------	-------------------	----------------	----------------	-----------------

Who: pilot group. *Done when:* at least three baseline observations per use case exist, and everyone in the pilot group has run Level 1 of the curriculum.

How phase one most commonly fails. Someone senior sees a demonstration, gets interested, and the scope grows in a fortnight from two use cases to "a strategy". You end up with a slide deck, no baseline and no policy, while the team quietly carries on doing whatever it was already doing. The defence is the written pair of use cases: everything else goes on a list marked "after day 90", and the list is shown, not argued with.

The second failure is a policy written to be safe rather than to be followed — a document that bans everything, which people comply with by not telling you what they do. If your policy would make the honest answers you got in week 1 into sackable offences, you have written the wrong policy.

Phase two — Days 31 to 60: run a real pilot

The goal. Evidence. By day 60 you want thirty or so logged observations across your two use cases and a clear, unexcited sense of whether either one is better, worse or the same.

Week 5 — train and start. Pilot group completes Levels 1 and 2 of the curriculum. They begin using the assistant on the two chosen use cases only. They log every single run, including the ones they abandoned halfway. Abandoned runs are data; quietly dropping them is how a pilot lies to you.

Who: owner trains, pilot group runs. *Done when:* every pilot member has logged at least two runs.

Week 6 — first honest look. Thirty minutes with the pilot group. Not a status update — a specific question to each person: "Show me the worst output you got this week." Then: "What did you have to fix?" The reviewer reports what they corrected before sign-off. Write down the failure types. You are building a list of what this tool gets wrong on your work specifically, which is more useful than any general warning anyone can give you.

Who: owner chairs, reviewer contributes. *Done when:* a written list of observed failure types exists and has been circulated to the pilot group.

Week 7 — tighten, do not widen. Take what you learned in week 6 and improve the prompts, not the scope. This is the week the pilot group does Level 3. If a use case is clearly not working, say so out loud now and either fix the prompt once or drop it. Dropping one of two use cases at week 7 is a normal outcome, not a defeat, and it is much cheaper than dropping it at week 12.

Who: pilot group. *Done when:* each use case has one agreed prompt, saved somewhere everyone can reach, with the placeholders written out.

Week 8 — the honest read. Compare the log against the baseline. Ask three questions and answer them in writing:

1. Did the task get faster once review time is included? Review time

counts. A draft produced in two minutes that takes forty minutes to check has not saved you thirty-eight minutes.

1. Did rework go up, down, or stay flat?
2. Would the reviewer be comfortable if this output went out under their

name every month?

If the answers are good, you widen. If they are mixed, you extend the pilot for a month rather than widening — an extension is a legitimate decision and it will make the eventual answer stronger. If they are poor, you stop that use case and say why, in one paragraph, to the whole team.

Who: owner writes, sponsor reads. *Done when:* a one-page written finding exists, with the log attached.

How phase two most commonly fails. The pilot group logs the successes and forgets the failures, usually without meaning to. People remember the run that worked and skip logging the three that did not, and by week 8 the data says the pilot went beautifully while the reviewer is quietly rewriting everything. Two defences: ask for the worst output, not the best, every week; and treat the reviewer's correction time as the real measure, because it is the one number nobody can flatter.

The other failure is the reviewer being cut out. Once people find the output "usually fine", the review step starts getting skipped under deadline pressure, and the first time you learn about it is when something wrong has left the building. Make sign-off a recorded step, not a habit.

Phase three — Days 61 to 90: widen only what held

The goal. Move from a pilot that worked to a routine the team follows without you standing over it, and leave behind something that survives the owner going on leave.

Week 9 — write the run sheets. For each surviving use case, write a one-page run sheet: the exact prompt with placeholders, what to attach, what to check before accepting the output, who signs off, and where the result is saved. If it does not fit on one page, the workflow is too complicated to hand over and you should simplify it before you spread it.

Who: pilot group writes, owner edits. *Done when:* one page per use case, in the same place as the policy.

Week 10 — teach the next group. The pilot group teaches Levels 1 and 2 to the next cohort, with the owner present. The people who did it teaching the people who have not is more convincing than any external session, and it exposes gaps in your run sheets within about twenty minutes.

Who: pilot group teaches. *Done when:* the second cohort has run the workflow once each, under supervision.

Week 11 — a deliberate stress test. Run one period-end, or one client job, with the new workflow in place and the owner not intervening. Log it. The interesting output is the list of questions people asked that the run sheet did not answer.

Who: the new cohort. *Done when:* the run sheet has been amended with the answers to those questions.

Week 12 — review the policy against reality. You now know things you did not know on day 14. Update the policy: the approved tools list, the data rules that turned out to be unworkable or too loose, the sign-off record, the failure types you found. Get the sponsor to re-approve. Re-circulate. Two-page limit still applies.

Who: owner and data/IT contact. *Done when:* version two is signed and sent.

Week 13 — decide what happens next, and write it down. One page to the sponsor: what you tried, what the log says, what you have stopped doing, what the next two use cases are, and what you are deliberately not doing yet. Add the standing arrangement: who keeps the log, how often it is reviewed, and when the policy is next looked at.

Who: owner. *Done when:* the sponsor has the page and the next review date is in a calendar.

How phase three most commonly fails. Everything widens at once. The pilot held for two use cases with five people who had been trained and watched; it is then treated as a general permission, thirty people start using it for whatever they like, quality drops, and the whole thing gets banned in month five by someone reacting to a single bad incident. Widen the use cases you tested, with the training you tested, and leave the rest on the list.

The quieter failure is that nothing sticks because the owner was the only one who ever cared. If, by week 13, no one other than the owner can run the workflow from the run sheet, you have a hobby rather than a process.

The senior who will not engage

You will have one. Often more than one, and often your best technical reviewer, which is not a coincidence — they got good by not trusting things.

Do not put them in the pilot group. Do not sit them in a demonstration. Both read as pressure and produce a public position they then have to defend for a year.

Instead, give them the role they are already good at: reviewer. Ask them to be the person who checks the pilot output and tells you what is wrong with it. That is an honest job, it uses their scepticism as an asset, and it puts them in contact with the actual quality of the output rather than the claims about it. Several people who start out flatly opposed end up adopting parts of it in their own work once they have seen where it fails, because now they know the shape of the risk. Some do not, which is a legitimate outcome as long as they are not blocking the team.

Two things to hold firm on. They do not get to veto other people's use inside the policy. And they do get to say "not on my client work" for a period, if they carry the professional responsibility for it — that is not obstruction, it is someone exercising judgement they are accountable for. If your work is regulated, whose name is on the file matters more than your adoption plan does.

What does not work: making it a performance objective, mocking them in front of juniors, or going around them to their staff. All three convert a sceptic into an opponent, and opponents are much more expensive.

The junior who has already raced ahead

You will have one of these too, usually a competent qualified-recently type who has been using an assistant for six months, is genuinely faster at some things, and has been quietly building their own way of working with it.

Your instinct will be to slow them down. Half-resist it. What you need from them, immediately and explicitly, is honesty about what they have already put into which tool — you cannot manage exposure you do not know about. Ask directly, once, with the same no-consequences framing as week 1.

Then give them a real job: writing prompts and run sheets for the pilot, and teaching Level 2. This turns private practice into a documented team asset and it gives them status for doing the thing they were going to do anyway, which is far more effective than a warning.

The one line you hold is review. The specific risk with a fast, confident junior is not that they use the tool; it is that they produce more work than their reviewer's normal capacity, at a quality that looks finished. Polished, well-

formatted output invites a lighter read. Say this to the reviewer out loud: the drafts now look better than they used to, and the appearance of the work no longer tells you much about the care behind it. Check the numbers and the source references, not the prose.

If they have been using an unapproved tool on client data, deal with it as an exposure to be assessed and reported per your own obligations, not as a character question. Whether it needs to be reported anywhere is a matter for your professional body's rules and your local law, and you should ask rather than assume.

The ninety days on one page

Week	Phase	What happens	Who leads	Done when
1	Contain and choose	Ask everyone what they already use; issue the one rule	Owner	Written list of current use; one rule circulated
2		Draft the two-page policy; check tool terms and data obligations	Owner + data/IT	Sponsor signs; team confirms they have read it
3		Choose exactly two use cases against the four tests	Owner + pilot group	Two one-sentence use cases with named reviewers
4		Record the "before": time and rework, labelled as estimates	Pilot group	Three baseline observations per use case
5	Pilot	Levels 1–2 training; pilot begins; every run logged	Owner	Two logged runs per person
6		"Show me the worst output" review; list failure types	Owner + reviewer	Written failure list circulated
7		Level 3 training; tighten prompts; drop what is not working	Pilot group	One agreed saved prompt per use case
8		Compare against baseline; decide widen, extend or stop	Owner + sponsor	One-page written finding with log attached
9	Widen what held	Write one-page run sheets per use case	Pilot group	One page each, stored with the policy
10		Pilot group teaches Levels 1–2 to the next cohort	Pilot group	Second cohort has run it once, supervised
11		Stress test on a real period-end or job, owner hands off	New cohort	Run sheet amended with the gaps found
12		Update policy against what actually happened	Owner + data/IT	Version two signed and re-circulated
13		Write up; set the standing review; name the next two use cases	Owner	Sponsor has the page; next review date booked

Part two: a four-level training curriculum

This is designed to be taught by a manager who is one step ahead of the room, not by a specialist. Each level stands alone. Nobody needs Level 4 unless they are designing workflows, and most of your team will get most of the benefit from Levels 1 to 3.

What you need to run it: a room with a screen, a working licence for whatever tool you have approved, and real examples from your own work with client-identifiable details removed. Made-up examples do not land with this audience; they will spot that the numbers are not real and switch off.

Teach in the order given. The most common curriculum mistake is starting at Level 2 because prompting seems like the practical bit. Without Level 1, people write better prompts and trust the answers more, which is the worst combination available.

Level	For	Length	The point
1	Everyone	45 minutes, one session	What it is and why it is confident when wrong
2	Everyone who will use it	90 minutes, two sessions	Writing a prompt that produces usable work
3	Anyone whose output is reviewed or reviews others	2 × 60 minutes	Verification, and working with your own documents
4	Managers and process owners	3 × 60 minutes plus follow-up	Repeatable workflows, and when to stop

Level 1 — What the thing actually is, and why it is confident when wrong

Who it is for. Everyone in the finance function, including the people who will never use it. Especially the people who review other people's work.

Length. Forty-five minutes. One session. Do not extend it; the content is small and the demonstration does the work.

Learning outcomes. By the end, each person can say in their own words: that the tool predicts plausible text rather than looking up facts; that fluency is not evidence of accuracy; that it will invent a reference, a figure or a paragraph number with exactly the same confident tone it uses when it is right; that it does not know what it does not know; and that its confident tone is a property of how it writes, not a signal about how sure it is.

Session outline.

1. *Ten minutes — what it is doing.* One explanation, no jargon: it

produces the next piece of text that fits, based on patterns in an enormous amount of writing. It is very good at the shape of a document — a memo reads like a memo, a reconciliation note reads like a reconciliation note. The shape being right tells you nothing about whether the contents are.

1. *Fifteen minutes — the live exercise below.* This is the session.
2. *Ten minutes — where it is strong.* Rewriting, summarising a document

you provide, drafting a first version, explaining something in simpler words, spotting what is missing from a list. Note what these have in common: you supply the substance, it supplies the form.

1. *Five minutes — where it is weak.* Anything requiring it to know a fact

you did not give it. Arithmetic you have not checked. Anything current. Anything where being 95% right is the same as being wrong.

1. *Five minutes — the team's rule.* Read out the one rule about named human review, and the data rules from your policy.

The live exercise: the confident wrong answer. On screen, in front of everyone, run these two prompts in order. Do not rehearse them first — the point is that you do not know what it will say either.

Summarise the disclosure requirements for [TOPIC] under [ACCOUNTING FRAMEWORK] for an entity of type [ENTITY TYPE]. Quote the exact paragraph reference for each requirement and state which pronouncement it comes from.

Then, whatever it produces, put the reference into your actual source material and read out what is really there. Sometimes it will be right. Sometimes it will produce a reference in the correct format that does not exist, or attach a real reference to the wrong requirement. Either result teaches the lesson, because the tone is identical both times.

Then run:

Here is a trial balance extract. Calculate the movement in each account between [PERIOD 1] and [PERIOD 2], the percentage change, and tell me which three movements a reviewer should question first.

[PASTE 15-20 ROWS OF A REAL TRIAL BALANCE, ANONYMISED]

Have someone check every arithmetic line on a calculator while the room watches. Discuss what happened. If it got the arithmetic right, that is also a useful result — say plainly that one correct run is not a guarantee, and that this is why the checking step is not optional.

How you tell whether it landed. At the end, ask each person to write one sentence: *why did it produce a confident wrong answer?* Collect the slips. If most of them say something about the data being out of date, it did not land — that is the "it needs newer information" model, and it leads people to trust it on anything historic. Re-run the exercise with a different topic. What you want to see is people describing it producing text that fits rather than text that is true.

The real test comes a fortnight later. Listen for the question people ask about someone else's draft. "Is this right?" means Level 1 did not land. "What did you check this against?" means it did.

Level 2 — Writing a usable prompt

Who it is for. Everyone who will actually use the tool for work.

Length. Ninety minutes, split into two forty-five minute sessions a week apart, with homework in between. The gap matters; people need to try it badly once on their own before the second half is useful.

Learning outcomes. Each person can write a prompt that states the role, the context, the task, the constraints and the output format; can supply source material rather than expecting the tool to know things; can ask it to show its working or flag its assumptions; and can recognise when a bad output is caused by a bad prompt rather than by a bad tool.

Session outline.

Session A (45 minutes).

1. *Five minutes — the failure they have all seen.* Someone types "write a variance analysis", gets something generic, concludes it is useless.

1. *Fifteen minutes — the five parts.* Role, context, task, constraints,

format. Walk through each with a finance example. The part finance people skip is context: the entity, the framework, the period, the materiality, who reads the output.

1. *Twenty minutes — the live exercise below.*

2. *Five minutes — homework.* Everyone brings one real task and the prompt

they wrote for it, plus the output, good or bad.

Session B (45 minutes).

1. *Twenty-five minutes — workshop the homework.* Put prompts on screen

anonymously. Fix them as a group. This is the most valuable half hour in the whole curriculum.

1. *Ten minutes — iteration.* Show that the second and third messages in a

conversation do more than the first: "that is too long", "drop the introduction", "you have assumed [X], I meant [Y]".

1. *Ten minutes — build the shared prompt list.* Every prompt that worked

goes in one place, with placeholders, so nobody rewrites it next month.

The live exercise: two prompts, one task. Run the thin version first:

Write a variance analysis commentary.

Read the result out. It will be plausible and useless. Then run:

You are drafting internal management commentary for the finance team of [ENTITY], a [INDUSTRY] business reporting under [ACCOUNTING FRAMEWORK], for the period [PERIOD]. The audience is [AUDIENCE], who are not accountants.

Below is the actual variance data. Using only the figures given:

1. Write one paragraph per variance above [THRESHOLD].
2. For each, state the movement in both absolute and percentage terms.
3. Where the data does not explain a cause, say that the cause is not determinable from the data provided. Do not speculate about causes.
4. List separately any figure you could not reconcile.

Constraints: maximum [NUMBER] words. Plain English. No adjectives about performance. Do not describe any result as strong, weak, disappointing or encouraging.

Data:

[PASTE VARIANCE TABLE]

Compare the two on screen. Point out specifically what the instruction "do not speculate about causes" prevented — invented explanations for variances are one of the most common and most damaging failure modes in finance drafting, because they read exactly like knowledge.

How you tell whether it landed. Two weeks after session B, look at the shared prompt list. If it has entries from more than half the attendees, it landed. If it has three entries, all from the owner, it did not — and the usual reason is that people are still using it in a private chat window for things they do not consider "real work". Ask them what they used it for yesterday; the answer is usually more informative than the list.

Level 3 — Verification, and working with your team's own documents

Who it is for. Anyone whose output gets reviewed, and everyone who reviews. In practice this is most of the team.

Length. Two sessions of sixty minutes, ideally a week apart.

Learning outcomes. Each person can name what must be checked in a given output and check it efficiently rather than re-doing the work; can ground a request in supplied documents and require the tool to quote its source lines; can recognise the specific failure of a confident summary of a document that says something slightly different; and can decide when verification costs more than the assistance saved, and say so.

Session outline.

Session A — verification.

1. *Ten minutes — the reviewer's problem.* Output now arrives looking finished. The old signals — clumsy phrasing, inconsistent formatting — no longer tell you how much care went in.

1. *Fifteen minutes — the four checks.* Numbers: recalculate, do not read.

Sources: does the quoted line exist in the document, in those words? Omissions: what is missing, which no output will tell you. Claims: any sentence asserting a fact you did not supply is unverified until you verify it.

1. *Twenty-five minutes — the live exercise below.*

2. *Ten minutes — recording it.* How your team evidences that a review happened. Whatever your file or working paper conventions are, decide this now rather than inventing it per job.

Session B — your own documents.

1. *Fifteen minutes — grounding.* The difference between asking it about a policy and giving it the policy. Give it the document, every time.

1. *Twenty minutes — the quoting technique.* Requiring the tool to quote the exact line it relied on, so that verification takes seconds rather than a re-read.

1. *Fifteen minutes — the limits.* Long documents get summarised unevenly, and the middle is where things go missing. Split them. Ask about one section at a time.

1. *Ten minutes — the data rules again.* Which of your documents may be uploaded at all, under your policy and your obligations. This will be the most-asked question in the whole curriculum, so have the answer written down.

The live exercise: find the three planted errors. Before the session, take a real piece of assisted output — a summary of an internal policy, or a commentary on a set of figures — and plant exactly three errors: one arithmetic, one a quoted phrase that does not appear in the source, and one material omission. Do not tell the room how many errors there are, or that they were planted, until afterwards.

Give everyone the output and the source documents. Fifteen minutes, working in pairs. Then compare what each pair found.

Almost every group finds the arithmetic error and misses the omission. That is the lesson, and it is much better felt than told: what is missing does not announce itself, and no amount of reading the output will reveal it. Only knowing what should have been there does.

Then teach the fix. Run this prompt on the same source material:

Using only the document below, answer this question: [QUESTION].

For every statement you make, quote the exact sentence from the document that supports it, in quotation marks, immediately after the statement. If the document does not address part of my question, list those parts under a heading "Not covered by this document". Do not use any knowledge from outside the document.

Document:
[PASTE DOCUMENT]

Show that the "Not covered" list is the part worth reading first. Check two of the quotes against the source in front of the room, including one that turns out to be paraphrased rather than quoted — if you cannot produce one naturally, do not fake it; say that paraphrase-instead-of-quote is a failure to watch for and move on.

How you tell whether it landed. Look at the pilot log. If the "rework needed?" column starts getting more entries after Level 3, the training worked — people are now catching things they previously accepted. A pilot whose error rate appears to fall immediately after verification training is usually a pilot that has stopped looking properly. Tell the team that in advance, so that finding errors reads as doing the job well rather than as a bad week.

Level 4 — Designing a repeatable workflow, and knowing when to stop

Who it is for. Managers, process owners, and the one or two people who will maintain this after the ninety days.

Length. Three sessions of sixty minutes, plus a follow-up four weeks later to review what was built.

Learning outcomes. Each person can convert an ad hoc use into a documented workflow with a fixed prompt, defined inputs, a checking step and a named sign-off; can estimate whether the workflow is actually worth it once review time is counted; can recognise a task that should not be automated at all; and can retire a workflow that has stopped earning its place.

Session outline.

Session A — from ad hoc to repeatable. What makes a task suitable: frequency, stable inputs, an existing check, a tolerable failure mode. The anatomy of a run sheet: trigger, inputs, prompt with placeholders, processing steps, checks, sign-off, where the output is stored. Each attendee picks one candidate from their own work.

Session B — build one. Attendees write their run sheet in the room and swap with a neighbour, who must try to follow it exactly as written and mark every point where they had to guess. A run sheet somebody else cannot follow is not finished.

Session C — knowing when to stop. Four stop signals, in order of how often they apply:

1. Verification costs more than the drafting saved. Count the review

minutes honestly. If they exceed the time saved, stop; the workflow is moving effort, not reducing it.

1. The failure mode is intolerable regardless of frequency. If a wrong output goes straight out with no human between it and the reader, the error rate does not need to be high to be unacceptable.

1. The task is the training. Some tasks are how juniors learn the business. Automating the first draft of a reconciliation note may cost you something you will not notice for two years. This is a judgement call, and it is yours to make deliberately rather than by default.

1. Nobody can explain the output. If the person signing cannot say how the result was arrived at, they are signing something they do not understand, and their name is still on it.

Follow-up session, four weeks later. Each attendee reports on their workflow using the log. Retire at least one thing. A curriculum that never retires anything is accumulating obligations.

The live exercise: build and run the month-end commentary workflow. Do this as a group, on the real task, in the session. Write the run sheet on screen. Then run it end to end with real anonymised data while somebody times each step and writes the numbers where everyone can see them. Include the review time. Then ask the room the only question that matters: knowing these numbers, would you do this next month?

Sometimes the answer is no. Let it be no. A Level 4 session that ends in a documented decision not to automate something is a success, and it will do more for your credibility with the sceptics than three that end in enthusiasm.

Use this prompt as the workflow's fixed step, filled in and saved:

You are producing the [PERIOD] management commentary for [ENTITY], following the same structure as the attached previous version.

Inputs: the current period figures below, and last period's commentary for structure and tone only – do not carry forward any figure or explanation from it.

Produce:

1. A commentary section for each of [SECTIONS], in the same order as the previous version.
2. Movements stated in both absolute and percentage terms.
3. A separate list headed "Requires input from the business" for anything the figures alone do not explain.
4. A separate list headed "Checks for the reviewer" naming every figure you calculated rather than copied.

Do not state causes that are not evidenced in the data provided. Do not characterise performance.

Current period figures:
[PASTE]

Previous commentary:
[PASTE]

The fourth instruction is the one that earns its place: a list of what the tool calculated tells the reviewer exactly where to point a calculator.

How you tell whether it landed. Four weeks later, look for two things: at least one workflow running without the trainer, and at least one workflow retired or never started because the numbers did not support it. If everything survives, nobody applied the stop signals, and you have taught enthusiasm rather than judgement.

When it stalls — and it usually does

Adoption commonly stalls somewhere between weeks six and ten. The novelty has worn off, month-end has arrived, and the people who were logging diligently have gone quiet. This is normal enough that you should plan for it rather than treat it as a verdict.

A stall is information. Read what it is telling you before you react, because the cure differs entirely by cause:

- **It stalled at month-end.** Nothing is wrong. People fell back to the

known method under deadline pressure, which is a reasonable thing for a professional to do. Restart the following week and expect the same thing next month until the workflow is genuinely faster than the old one under pressure, not just in a quiet week.

- **It stalled because the output was not good enough.** The most useful

outcome available to you. Something you hoped would work does not, on your data, for your task. Write that down specifically — not "it isn't accurate" but "it could not reconcile the intercompany balances without more instruction than the task saved". That sentence is worth more than a successful pilot on a task you did not care about.

- **It stalled because verification cost too much.** Legitimate. Either

narrow the task so the check is cheaper, or stop. Do not solve it by checking less.

- **It stalled because one person was carrying it.** The most common cause

and the least discussed. The owner got busy, and nothing had been handed over. Fix the handover, not the enthusiasm.

- **It stalled because nobody was ever asked to.** Also common. There was a

policy and a licence and no instruction. If you want a task done a particular way, ask for it in the way you ask for anything else.

What not to do: relaunch. A second launch with more energy and a new name, aimed at the same people who watched the first one fade, costs you credibility you will want later. Fix the specific cause, quietly, and let the second attempt be smaller than the first.

What to do instead is dull and works. Return to one use case, with one group, with the log. Ask the reviewer what they are still correcting. Put the next review date in the calendar and keep it even when there is nothing to report — especially then, because "nothing happened this month" is the finding that most often needs saying out loud.

And be willing to conclude that some of it does not suit your team, your work or your obligations right now. That is a defensible professional judgement, it is a great deal cheaper than a slow expensive drift, and it leaves you free to look again in six months with a log full of things you actually know rather than things you were told.